

Chapter 1

April 19, 2025

1 Hands-On Data Preprocessing in Python

Learn how to effectively prepare data for successful data analytics

AUTHOR: Dr. Roy Jafari

1.0.1 Chapter 1: Review of the core modules NumPy, Pandas, and Matplotlib

Hello World! This is a Markdown chunk!

```
[2]: print('Hello World! This is Code Chunk!')
```

Hello World! This is Code Chunk!

```
[4]: import numpy as np
```

```
[5]: lst_nums = [2,5,7,11,13,17,23,31,37,41,43,47]
      np.mean(lst_nums)
```

```
[5]: np.float64(23.083333333333332)
```

```
[6]: lst_nums = [2,5,7,11,13,17,23,31,37,41,43,47]
      ary_nums = np.array(lst_nums)
      ary_nums.mean()
```

```
[6]: np.float64(23.083333333333332)
```

```
[7]: type(lst_nums)
```

```
[7]: list
```

```
[8]: type(ary_nums)
```

```
[8]: numpy.ndarray
```

```
[9]: np.arange(15)
```

```
[9]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14])
```

```
[10]: np.arange(5,15)
```

```
[10]: array([ 5,  6,  7,  8,  9, 10, 11, 12, 13, 14])
```

```
[11]: np.arange(-7.1,7)
```

```
[11]: array([-7.1, -6.1, -5.1, -4.1, -3.1, -2.1, -1.1, -0.1,  0.9,  1.9,  2.9,
           3.9,  4.9,  5.9,  6.9])
```

```
[12]: np.zeros([4,5])
```

```
[12]: array([[0., 0., 0., 0., 0.],
           [0., 0., 0., 0., 0.],
           [0., 0., 0., 0., 0.],
           [0., 0., 0., 0., 0.]])
```

```
[13]: np.ones(7)
```

```
[13]: array([1., 1., 1., 1., 1., 1., 1.])
```

Example #1

The following is the grade data of ten students. Create a code using NumPy that calculate and report their grade average.

```
Names = ['Jevon', 'Dawn', 'Kayleigh', 'Jadene', 'Kennedy', 'Kaydee', 'Ansh', 'Flynn', 'Kier',
Math_grades = [80, 50, 60, 70, 60, 100, 70, 70, 60, 70]
Science_grades = [90, 80, 50, 50, 60, 50, 90, 70, 80, 80]
History_grades = [60, 90, 50, 90, 100, 100, 100, 100, 90, 70]
```

```
[14]: Names = ['Jevon', 'Dawn', 'Kayleigh', 'Jadene', 'Kennedy', 'Kaydee',
           'Ansh', 'Flynn', 'Kier', 'Clarence']
Math_grades = [80, 50, 60, 70, 60, 100, 70, 70, 60, 70]
Science_grades = [90, 80, 50, 50, 60, 50, 90, 70, 80, 80]
History_grades = [60, 90, 50, 90, 100, 100, 100, 100, 90, 70]
```

```
[15]: Average_grades = np.zeros(10)
print(Average_grades)

for i, name in enumerate(Names):
    Average_grades[i] = np.mean([Math_grades[i], Science_grades[i],
                                History_grades[i]])

print(Average_grades)
```

```
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
[76.66666667 73.33333333 53.33333333 70.          73.33333333 83.33333333
 86.66666667 80.          76.66666667 73.33333333]
```

```
[16]: # better-looking report

for i, name in enumerate(Names):
```

```
print("Average for {} : {}".format(name,Average_grades[i]))
```

```
Average for Jevon : 76.66666666666667
Average for Dawn : 73.33333333333333
Average for Kayleigh : 53.33333333333336
Average for Jadene : 70.0
Average for Kennedy : 73.33333333333333
Average for Kaydee : 83.33333333333333
Average for Ansh : 86.66666666666667
Average for Flynn : 80.0
Average for Kier : 76.66666666666667
Average for Clarence : 73.33333333333333
```

```
[17]: np.linspace(0,1,21)
```

```
[17]: array([0. , 0.05, 0.1 , 0.15, 0.2 , 0.25, 0.3 , 0.35, 0.4 , 0.45, 0.5 ,
          0.55, 0.6 , 0.65, 0.7 , 0.75, 0.8 , 0.85, 0.9 , 0.95, 1.  ])
```

```
[18]: np.linspace(10,1000,100)
```

```
[18]: array([ 10.,  20.,  30.,  40.,  50.,  60.,  70.,  80.,  90.,
          100., 110., 120., 130., 140., 150., 160., 170., 180.,
          190., 200., 210., 220., 230., 240., 250., 260., 270.,
          280., 290., 300., 310., 320., 330., 340., 350., 360.,
          370., 380., 390., 400., 410., 420., 430., 440., 450.,
          460., 470., 480., 490., 500., 510., 520., 530., 540.,
          550., 560., 570., 580., 590., 600., 610., 620., 630.,
          640., 650., 660., 670., 680., 690., 700., 710., 720.,
          730., 740., 750., 760., 770., 780., 790., 800., 810.,
          820., 830., 840., 850., 860., 870., 880., 890., 900.,
          910., 920., 930., 940., 950., 960., 970., 980., 990.,
          1000.] )
```

Example 2

We are interested in finding the value(s) that holds the following mathematical statement.

$$x^2-5x+6=0$$

And imagine that we don't know that the statement can be simplified easily to find either -1 or +1 will hold the statement.

$$x^2-5x+6=(x-2)(x-3)$$

so we would like to use NumPy to try out any values between -100 and 100 and see what the answer is.

```
[20]: Candidates = np.linspace(-1000,1000,2001)
print(Candidates)

for candidate in Candidates:
```

```
if(candidate**2 - 5*candidate +6 ==0):
    print("Just found a possible answer: {}".format(candidate))
```

```
[-1000. -999. -998. ... 998. 999. 1000.]
```

```
Just found a possible answer: 2.0
```

```
Just found a possible answer: 3.0
```

2 Adult Dataset

“Census Income” dataset.

Number of Instances: 48842 Number of Attributes: 14 Date Donated: 1996-05-01 Missing Values?: Yes

2.0.1 Attributes:

Number of Attributes: 6 continuous, 8 nominal attributes

- age: continuous.
- workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
- fnlwgt: continuous.
- education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
- education-num: continuous.
- marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
- occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
- relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
- race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
- sex: Female, Male.
- capital-gain: continuous.
- capital-loss: continuous.
- hours-per-week: continuous.
- native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.
- class: >50K, <=50K

```
[22]: import pandas as pd

adult_df = pd.read_csv('adult.csv')
adult_df.head()
```

```
[22]: age          workclass  fnlwgt  education  education-num \
0    39          State-gov   77516   Bachelors           13
1    50  Self-emp-not-inc   83311   Bachelors           13
2    38          Private  215646    HS-grad            9
3    53          Private  234721      11th             7
4    28          Private  338409   Bachelors           13

      marital-status      occupation  relationship  race    sex \
0      Never-married      Adm-clerical  Not-in-family  White   Male
1  Married-civ-spouse  Exec-managerial      Husband  White   Male
2          Divorced  Handlers-cleaners  Not-in-family  White   Male
3  Married-civ-spouse  Handlers-cleaners      Husband  Black   Male
4  Married-civ-spouse  Prof-specialty      Wife    Black  Female

      capitalGain  capitalLoss  hoursPerWeek  nativeCountry  income
0          2174           0           40  United-States  <=50K
1           0           0           13  United-States  <=50K
2           0           0           40  United-States  <=50K
3           0           0           40  United-States  <=50K
4           0           0           40          Cuba  <=50K
```

```
[23]: type(adult_df.age)
```

```
[23]: pandas.core.series.Series
```

```
[24]: type(adult_df)
```

```
[24]: pandas.core.frame.DataFrame
```

```
[27]: adult_df.loc[0]
```

```
[27]: age          39
workclass      State-gov
fnlwgt         77516
education      Bachelors
education-num   13
marital-status  Never-married
occupation      Adm-clerical
relationship    Not-in-family
race            White
sex             Male
capitalGain     2174
capitalLoss     0
hoursPerWeek    40
nativeCountry   United-States
income          <=50K
Name: 0, dtype: object
```

```
[25]: adult_df.loc[0].index
```

```
[25]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',  
         'marital-status', 'occupation', 'relationship', 'race', 'sex',  
         'capitalGain', 'capitalLoss', 'hoursPerWeek', 'nativeCountry',  
         'income'],  
        dtype='object')
```

```
[26]: adult_df.age.index
```

```
[26]: RangeIndex(start=0, stop=32561, step=1)
```

```
[28]: adult_df.set_index(np.arange(10000,42561),inplace=True)
```

```
[29]: adult_df.set_index(np.arange(10000,42561))
```

```
[29]:
```

	age	workclass	fnlwgt	education	education-num	\
10000	39	State-gov	77516	Bachelors	13	
10001	50	Self-emp-not-inc	83311	Bachelors	13	
10002	38	Private	215646	HS-grad	9	
10003	53	Private	234721	11th	7	
10004	28	Private	338409	Bachelors	13	
...	
42556	27	Private	257302	Assoc-acdm	12	
42557	40	Private	154374	HS-grad	9	
42558	58	Private	151910	HS-grad	9	
42559	22	Private	201490	HS-grad	9	
42560	52	Self-emp-inc	287927	HS-grad	9	

	marital-status	occupation	relationship	race	sex	\
10000	Never-married	Adm-clerical	Not-in-family	White	Male	
10001	Married-civ-spouse	Exec-managerial	Husband	White	Male	
10002	Divorced	Handlers-cleaners	Not-in-family	White	Male	
10003	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
10004	Married-civ-spouse	Prof-specialty	Wife	Black	Female	
...	
42556	Married-civ-spouse	Tech-support	Wife	White	Female	
42557	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	
42558	Widowed	Adm-clerical	Unmarried	White	Female	
42559	Never-married	Adm-clerical	Own-child	White	Male	
42560	Married-civ-spouse	Exec-managerial	Wife	White	Female	

	capitalGain	capitalLoss	hoursPerWeek	nativeCountry	income
10000	2174	0	40	United-States	<=50K
10001	0	0	13	United-States	<=50K
10002	0	0	40	United-States	<=50K
10003	0	0	40	United-States	<=50K

10004	0	0	40	Cuba	<=50K
...
42556	0	0	38	United-States	<=50K
42557	0	0	40	United-States	>50K
42558	0	0	40	United-States	<=50K
42559	0	0	20	United-States	<=50K
42560	15024	0	40	United-States	>50K

[32561 rows x 15 columns]

```
[32]: # adult_df.education-num # Error
adult_df["education-num"]
```

```
[32]: 10000    13
      10001    13
      10002     9
      10003     7
      10004    13
      ..
      42556    12
      42557     9
      42558     9
      42559     9
      42560     9
      Name: education-num, Length: 32561, dtype: int64
```

```
[33]: adult_df.iloc[2].loc['education']
```

```
[33]: 'HS-grad'
```

```
[34]: adult_df.education.loc[10002]
```

```
[34]: 'HS-grad'
```

```
[35]: adult_df['education'].iloc[2]
```

```
[35]: 'HS-grad'
```

```
[36]: adult_df.at[10002, 'education']
```

```
[36]: 'HS-grad'
```

```
[37]: row_series = adult_df.loc[10002]
      print(row_series.loc['education'])
      print(row_series.iloc[3])
      print(row_series['education'])
      print(row_series.education)
```

HS-grad
HS-grad
HS-grad
HS-grad

```
[38]: columns_series = adult_df.education
      print(columns_series.loc[10002])
      print(columns_series.iloc[2])
      print(columns_series[10002])
      # print(row_series.10002) This will give syntax error!
```

HS-grad
HS-grad
HS-grad

3 Slicing

```
[39]: my_array = np.array([[2,3,5,7],[11,13,17,19],
                          [23,29,31,37,],[41,43,47,49]])
      my_array
```

```
[39]: array([[ 2,  3,  5,  7],
            [11, 13, 17, 19],
            [23, 29, 31, 37],
            [41, 43, 47, 49]])
```

```
[40]: my_array[1,1]
```

```
[40]: np.int64(13)
```

```
[41]: my_array[1,:]
```

```
[41]: array([11, 13, 17, 19])
```

```
[42]: my_array[:,1]
```

```
[42]: array([ 3, 13, 29, 43])
```

```
[43]: my_array
```

```
[43]: array([[ 2,  3,  5,  7],
            [11, 13, 17, 19],
            [23, 29, 31, 37],
            [41, 43, 47, 49]])
```

```
[44]: my_array[1:3,:]
```



```
[44]: array([[11, 13, 17, 19],
           [23, 29, 31, 37]])
```

```
[45]: my_array[1:3,0:2]
```

```
[45]: array([[11, 13],
           [23, 29]])
```

```
[46]: my_array[1:3,[0,2]]
```

```
[46]: array([[11, 17],
           [23, 31]])
```

```
[47]: adult_df.loc[:, 'education': 'occupation']
```

```
[47]:
```

	education	education-num	marital-status	occupation
10000	Bachelors	13	Never-married	Adm-clerical
10001	Bachelors	13	Married-civ-spouse	Exec-managerial
10002	HS-grad	9	Divorced	Handlers-cleaners
10003	11th	7	Married-civ-spouse	Handlers-cleaners
10004	Bachelors	13	Married-civ-spouse	Prof-specialty
...
42556	Assoc-acdm	12	Married-civ-spouse	Tech-support
42557	HS-grad	9	Married-civ-spouse	Machine-op-inspct
42558	HS-grad	9	Widowed	Adm-clerical
42559	HS-grad	9	Never-married	Adm-clerical
42560	HS-grad	9	Married-civ-spouse	Exec-managerial

```
[32561 rows x 4 columns]
```

```
[48]: adult_df.sort_values('education-num').reset_index().iloc[1:32561:3617]
```

```
[48]:
```

	index	age	workclass	fnlwgt	education	education-num	\
1	42432	36	Private	208068	Preschool	1	
3618	14676	17	Private	262511	11th	7	
7235	17294	53	Private	141388	HS-grad	9	
10852	11766	24	Private	228424	HS-grad	9	
14469	25218	30	Private	144064	HS-grad	9	
18086	31147	39	Private	348521	Some-college	10	
21703	25115	35	Private	257042	Some-college	10	
25320	25231	30	State-gov	193380	Bachelors	13	
28937	40173	24	Private	330571	Bachelors	13	
32554	18280	55	Local-gov	37869	Doctorate	16	

	marital-status	occupation	relationship	race	sex	\
1	Divorced	Other-service	Not-in-family	Other	Male	
3618	Never-married	Sales	Own-child	White	Male	
7235	Married-civ-spouse	Transport-moving	Husband	White	Male	

10852	Never-married	Handlers-cleaners	Other-relative	Black	Male
14469	Married-civ-spouse	Exec-managerial	Husband	White	Male
18086	Married-civ-spouse	Farming-fishing	Husband	White	Male
21703	Married-civ-spouse	Prof-specialty	Husband	White	Male
25320	Never-married	Prof-specialty	Other-relative	White	Male
28937	Married-civ-spouse	Prof-specialty	Wife	White	Female
32554	Never-married	Prof-specialty	Not-in-family	White	Female

	capitalGain	capitalLoss	hoursPerWeek	nativeCountry	income
1	0	0	72	Mexico	<=50K
3618	0	0	20	United-States	<=50K
7235	0	0	40	United-States	>50K
10852	0	0	40	United-States	<=50K
14469	0	0	62	United-States	<=50K
18086	0	2415	99	United-States	>50K
21703	0	0	50	United-States	<=50K
25320	0	0	35	United-States	<=50K
28937	0	0	40	United-States	<=50K
32554	0	0	60	United-States	<=50K

4 Boolean Masking

```
[49]: twopowers_sr = pd.Series([1,2,4,8,16,32,64,128,256,512,1024])
      BM = [False,False,False,True,False,False,False,True,True,True,True]
      twopowers_sr[BM]
```

```
[49]: 3      8
      7     128
      8     256
      9     512
     10    1024
      dtype: int64
```

```
[50]: twopowers_sr >=500
```

```
[50]: 0      False
      1      False
      2      False
      3      False
      4      False
      5      False
      6      False
      7      False
      8      False
      9       True
     10       True
```

dtype: bool

```
[51]: BM = twopowers_sr >=500  
      twopowers_sr[BM]
```

```
[51]: 9      512  
      10     1024  
      dtype: int64
```

```
[52]: twopowers_sr[twopowers_sr >=500]
```

```
[52]: 9      512  
      10     1024  
      dtype: int64
```

```
[53]: BM = adult_df.education == 'Preschool'  
      print('Mean: {}'.format(np.mean(adult_df[BM].age)))  
      print('Median: {}'.format(np.median(adult_df[BM].age)))
```

Mean: 42.76470588235294

Median: 41.0

```
[54]: BM1 = adult_df['education-num'] > 10  
      BM2 = adult_df['education-num'] < 10  
  
      print('More than 10 years of education - Capital Gain: {}'.  
            .format(np.mean(adult_df[BM1].capitalGain)))  
      print('Less than 10 years of education - Capital Gain: {}'.  
            .format(np.mean(adult_df[BM2].capitalGain)))
```

More than 10 years of education - Capital Gain: 2230.9397109166985

Less than 10 years of education - Capital Gain: 492.25532059102613

5 Get to know a dataset

```
[55]: adult_df.shape
```

```
[55]: (32561, 15)
```

```
[56]: adult_df.columns
```

```
[56]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',  
        'marital-status', 'occupation', 'relationship', 'race', 'sex',  
        'capitalGain', 'capitalLoss', 'hoursPerWeek', 'nativeCountry',  
        'income'],  
        dtype='object')
```

```
[57]: adult_df.columns = ['age', 'workclass', 'fnlwgt', 'education',
                          'education_num', 'marital_status', 'occupation',
                          'relationship', 'race', 'sex', 'capitalGain',
                          'capitalLoss', 'hoursPerWeek', 'nativeCountry',
                          'income']
```

```
[59]: adult_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 32561 entries, 10000 to 42560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             30725 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education_num         32561 non-null  int64
5   marital_status        32561 non-null  object
6   occupation            30718 non-null  object
7   relationship          32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
10  capitalGain           32561 non-null  int64
11  capitalLoss           32561 non-null  int64
12  hoursPerWeek          32561 non-null  int64
13  nativeCountry         31978 non-null  object
14  income                32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 5.0+ MB
```

```
[58]: adult_df.describe()
```

```
[58]:
```


	age	fnlwgt	education_num	capitalGain	capitalLoss	\
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	

	hoursPerWeek
count	32561.000000
mean	40.437456
std	12.347429
min	1.000000

```
25%      40.000000
50%      40.000000
75%      45.000000
max       99.000000
```


```
[63]: adult_df.age.plot.hist()
```

```
[63]: <Axes: ylabel='Frequency'>
```



```
[64]: adult_df.age.plot.box()
```

```
[64]: <Axes: >
```



```
[65]: adult_df.relationship.unique()
```

```
[65]: array(['Not-in-family', 'Husband', 'Wife', 'Own-child', 'Unmarried',
          'Other-relative'], dtype=object)
```

```
[66]: adult_df.relationship.value_counts()
```

```
[66]: relationship
      Husband      13193
Not-in-family    8305
    Own-child     5068
    Unmarried     3446
      Wife        1568
Other-relative     981
Name: count, dtype: int64
```


```
[67]: adult_df.relationship.value_counts(normalize=True)
```

```
[67]: relationship
      Husband      0.405178
Not-in-family    0.255060
    Own-child     0.155646
```

```
Unmarried      0.105832
Wife           0.048156
Other-relative  0.030128
Name: proportion, dtype: float64
```

```
[68]: adult_df.relationship.value_counts().plot.bar()
```

```
[68]: <Axes: xlabel='relationship'>
```



6 Apply a function

```
[71]: def MultiplyBy2(n):
      return n*2

adult_df.age.apply(MultiplyBy2)
```

```
[71]: 10000    78
      10001   100
      10002    76
      10003   106
      10004    56
      ...
      42556    54
      42557    80
      42558   116
      42559    44
      42560   104
      Name: age, Length: 32561, dtype: int64
```

6.0.1 Applying a Function - Analytic Example 1

Divide every value in column fnlwgt by the sum of all its values.

```
[72]: total_fnlwgt = adult_df.fnlwgt.sum()

def CalculatePercentage(v):
    return v/total_fnlwgt*100

adult_df.fnlwgt = adult_df.fnlwgt.apply(CalculatePercentage)
adult_df
```

```
[72]:      age      workclass      fnlwgt      education      education_num \
10000    39      State-gov    0.001254      Bachelors              13
10001    50  Self-emp-not-inc    0.001348      Bachelors              13
10002    38      Private    0.003490      HS-grad              9
10003    53      Private    0.003798      11th              7
10004    28      Private    0.005476      Bachelors              13
...      ...      ...      ...      ...      ...
42556    27      Private    0.004164      Assoc-acdm              12
42557    40      Private    0.002498      HS-grad              9
42558    58      Private    0.002458      HS-grad              9
42559    22      Private    0.003261      HS-grad              9
42560    52  Self-emp-inc    0.004659      HS-grad              9

      marital_status      occupation      relationship      race      sex \
10000      Never-married      Adm-clerical      Not-in-family      White      Male
10001      Married-civ-spouse      Exec-managerial      Husband      White      Male
10002      Divorced      Handlers-cleaners      Not-in-family      White      Male
10003      Married-civ-spouse      Handlers-cleaners      Husband      Black      Male
10004      Married-civ-spouse      Prof-specialty      Wife      Black      Female
...      ...      ...      ...      ...      ...
42556      Married-civ-spouse      Tech-support      Wife      White      Female
42557      Married-civ-spouse      Machine-op-inspct      Husband      White      Male
42558      Widowed      Adm-clerical      Unmarried      White      Female
```


42559	Never-married	Adm-clerical	Own-child	White	Male
42560	Married-civ-spouse	Exec-managerial	Wife	White	Female

	capitalGain	capitalLoss	hoursPerWeek	nativeCountry	income
10000	2174	0	40	United-States	<=50K
10001	0	0	13	United-States	<=50K
10002	0	0	40	United-States	<=50K
10003	0	0	40	United-States	<=50K
10004	0	0	40	Cuba	<=50K
...
42556	0	0	38	United-States	<=50K
42557	0	0	40	United-States	>50K
42558	0	0	40	United-States	<=50K
42559	0	0	20	United-States	<=50K
42560	15024	0	40	United-States	>50K

[32561 rows x 15 columns]

```
[73]: total_fnlwgt = adult_df.fnlwgt.sum()

adult_df.fnlwgt = adult_df.fnlwgt.apply(lambda v: v/total_fnlwgt*100)
adult_df
```

```
[73]:
```

	age	workclass	fnlwgt	education	education_num	\
10000	39	State-gov	0.001254	Bachelors	13	
10001	50	Self-emp-not-inc	0.001348	Bachelors	13	
10002	38	Private	0.003490	HS-grad	9	
10003	53	Private	0.003798	11th	7	
10004	28	Private	0.005476	Bachelors	13	
...	
42556	27	Private	0.004164	Assoc-acdm	12	
42557	40	Private	0.002498	HS-grad	9	
42558	58	Private	0.002458	HS-grad	9	
42559	22	Private	0.003261	HS-grad	9	
42560	52	Self-emp-inc	0.004659	HS-grad	9	

	marital_status	occupation	relationship	race	sex	\
10000	Never-married	Adm-clerical	Not-in-family	White	Male	
10001	Married-civ-spouse	Exec-managerial	Husband	White	Male	
10002	Divorced	Handlers-cleaners	Not-in-family	White	Male	
10003	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
10004	Married-civ-spouse	Prof-specialty	Wife	Black	Female	
...	
42556	Married-civ-spouse	Tech-support	Wife	White	Female	
42557	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	
42558	Widowed	Adm-clerical	Unmarried	White	Female	
42559	Never-married	Adm-clerical	Own-child	White	Male	

42560	Married-civ-spouse	Exec-managerial		Wife	White	Female
	capitalGain	capitalLoss	hoursPerWeek	nativeCountry	income	
10000	2174	0	40	United-States	<=50K	
10001	0	0	13	United-States	<=50K	
10002	0	0	40	United-States	<=50K	
10003	0	0	40	United-States	<=50K	
10004	0	0	40	Cuba	<=50K	
...	
42556	0	0	38	United-States	<=50K	
42557	0	0	40	United-States	>50K	
42558	0	0	40	United-States	<=50K	
42559	0	0	20	United-States	<=50K	
42560	15024	0	40	United-States	>50K	

[32561 rows x 15 columns]

```
[74]: def CalcLifeNoEd(row):
      return row.age - row.education_num

adult_df.apply(CalcLifeNoEd,axis=1)
```

```
[74]: 10000    26
      10001    37
      10002    29
      10003    46
      10004    15
      ..
      42556    15
      42557    31
      42558    49
      42559    13
      42560    43
      Length: 32561, dtype: int64
```

```
[75]: adult_df.apply(lambda r: r.age-r.education_num,axis=1)
```

```
[75]: 10000    26
      10001    37
      10002    29
      10003    46
      10004    15
      ..
      42556    15
      42557    31
      42558    49
      42559    13
```

```
42560    43
Length: 32561, dtype: int64
```

```
[76]: adult_df['lifeNoEd'] = adult_df.apply(
        lambda r: r.age-r.education_num,axis=1)

adult_df['capitalNet'] = adult_df.apply(
        lambda r: r.capitalGain - r.capitalLoss,axis=1)

adult_df[['education_num','lifeNoEd','capitalNet']].corr()
```

```
[76]:
```

	education_num	lifeNoEd	capitalNet
education_num	1.000000	-0.150452	0.117891
lifeNoEd	-0.150452	1.000000	0.051490
capitalNet	0.117891	0.051490	1.000000

7 Groupby

```
[77]: adult_df.groupby('marital_status').size()
```

```
[77]: marital_status
Divorced          4443
Married-AF-spouse    23
Married-civ-spouse 14976
Married-spouse-absent  418
Never-married     10683
Separated         1025
Widowed           993
dtype: int64
```

```
[81]: adult_df.groupby(['marital_status', 'sex']).size()
```

```
[81]: marital_status    sex
Divorced            Female    2672
                   Male      1771
Married-AF-spouse   Female     14
                   Male       9
Married-civ-spouse  Female   1657
                   Male  13319
Married-spouse-absent Female    205
                   Male    213
Never-married       Female   4767
                   Male   5916
Separated           Female    631
                   Male    394
Widowed             Female    825
                   Male    168
```

dtype: int64

```
[82]: adult_df.groupby(['marital_status', 'sex']).age.median()
```

```
[82]: marital_status    sex
Divorced              Female    43.0
                   Male      42.0
Married-AF-spouse     Female    31.0
                   Male      29.0
Married-civ-spouse    Female    38.0
                   Male      43.0
Married-spouse-absent Female    39.0
                   Male      41.0
Never-married         Female    25.0
                   Male      25.0
Separated             Female    39.0
                   Male      38.0
Widowed              Female    60.0
                   Male      62.5
Name: age, dtype: float64
```

```
[83]: adult_df.groupby(['race', 'sex']).capitalNet.mean()
```

```
[83]: race    sex
Amer-Indian-Eskimo Female    530.142857
                   Male      628.864583
Asian-Pac-Islander Female    727.583815
                   Male     1707.440115
Black          Female    471.142765
                   Male      627.268324
Other           Female    218.385321
                   Male     1314.438272
White           Female    508.219857
                   Male     1266.413112
Name: capitalNet, dtype: float64
```

```
[84]: adult_df.groupby(['race', 'sex', 'income']).fnlwgt.mean()
```

```
[84]: race    sex    income
Amer-Indian-Eskimo Female <=50K    0.001764
                   >50K    0.002395
                   Male  <=50K    0.002046
                   >50K    0.001954
Asian-Pac-Islander Female <=50K    0.002398
                   >50K    0.002305
                   Male  <=50K    0.002652
                   >50K    0.002762
```

Black	Female	<=50K	0.003454
		>50K	0.003331
	Male	<=50K	0.003922
		>50K	0.003971
Other	Female	<=50K	0.002803
		>50K	0.002593
	Male	<=50K	0.003478
		>50K	0.003310
White	Female	<=50K	0.002969
		>50K	0.002978
	Male	<=50K	0.003074
		>50K	0.003025

Name: fnlwgt, dtype: float64

```
[85]: grb_result =adult_df.groupby(['race','sex']).capitalNet.mean()

print(grb_result.index)
```

```
MultiIndex([('Amer-Indian-Eskimo', 'Female'),
            ('Amer-Indian-Eskimo', 'Male'),
            ('Asian-Pac-Islander', 'Female'),
            ('Asian-Pac-Islander', 'Male'),
            ('Black', 'Female'),
            ('Black', 'Male'),
            ('Other', 'Female'),
            ('Other', 'Male'),
            ('White', 'Female'),
            ('White', 'Male')],
            names=['race', 'sex'])
```

```
[86]: grb_result =adult_df.groupby(['race','sex']).capitalNet.mean()
grb_result
```

```
[86]: race      sex
Amer-Indian-Eskimo  Female    530.142857
                   Male      628.864583
Asian-Pac-Islander  Female    727.583815
                   Male     1707.440115
Black               Female    471.142765
                   Male      627.268324
Other               Female    218.385321
                   Male     1314.438272
White               Female    508.219857
                   Male     1266.413112
Name: capitalNet, dtype: float64
```

```
[87]: grb_result.unstack()
```

```
[87]: sex                Female      Male
      race
Amer-Indian-Eskimo  530.142857   628.864583
Asian-Pac-Islander  727.583815  1707.440115
Black               471.142765   627.268324
Other               218.385321  1314.438272
White               508.219857  1266.413112
```

```
[88]: mlt_seris =adult_df.groupby(['race','sex','income']).fnlwgt.mean()
      mlt_seris
```

```
[88]: race                sex      income
Amer-Indian-Eskimo  Female  <=50K      0.001764
                  >50K      0.002395
                  Male    <=50K      0.002046
                  >50K      0.001954
Asian-Pac-Islander  Female  <=50K      0.002398
                  >50K      0.002305
                  Male    <=50K      0.002652
                  >50K      0.002762
Black               Female  <=50K      0.003454
                  >50K      0.003331
                  Male    <=50K      0.003922
                  >50K      0.003971
Other               Female  <=50K      0.002803
                  >50K      0.002593
                  Male    <=50K      0.003478
                  >50K      0.003310
White               Female  <=50K      0.002969
                  >50K      0.002978
                  Male    <=50K      0.003074
                  >50K      0.003025
Name: fnlwgt, dtype: float64
```

```
[89]: mlt_seris.unstack()
```

```
[89]: income                <=50K      >50K
      race                sex
Amer-Indian-Eskimo  Female  0.001764  0.002395
                  Male    0.002046  0.001954
Asian-Pac-Islander  Female  0.002398  0.002305
                  Male    0.002652  0.002762
Black               Female  0.003454  0.003331
                  Male    0.003922  0.003971
Other               Female  0.002803  0.002593
                  Male    0.003478  0.003310
White               Female  0.002969  0.002978
```

	Male	0.003074	0.003025
--	------	----------	----------

```
[90]: mlt_seris.unstack().unstack()
```

```
[90]: income          <=50K          >50K
      sex             Female      Male      Female      Male
      race
Amer-Indian-Eskimo  0.001764  0.002046  0.002395  0.001954
Asian-Pac-Islander  0.002398  0.002652  0.002305  0.002762
Black               0.003454  0.003922  0.003331  0.003971
Other               0.002803  0.003478  0.002593  0.003310
White               0.002969  0.003074  0.002978  0.003025
```

```
[91]: mlt_df= mlt_seris.unstack().unstack()
      mlt_df.columns
```

```
[91]: MultiIndex([('<=50K', 'Female'),
                  ('<=50K', 'Male'),
                  ('>50K', 'Female'),
                  ('>50K', 'Male')],
              names=['income', 'sex'])
```

```
[92]: mlt_df.stack()
```

/tmp/ipykernel_2332935/2163858474.py:1: FutureWarning: The previous implementation of stack is deprecated and will be removed in a future version of pandas. See the What's New notes for pandas 2.1.0 for details. Specify future_stack=True to adopt the new implementation and silence this warning.

```
mlt_df.stack()
```

```
[92]: income          <=50K          >50K
      race          sex
Amer-Indian-Eskimo Female  0.001764  0.002395
                  Male   0.002046  0.001954
Asian-Pac-Islander Female  0.002398  0.002305
                  Male   0.002652  0.002762
Black              Female  0.003454  0.003331
                  Male   0.003922  0.003971
Other              Female  0.002803  0.002593
                  Male   0.003478  0.003310
White              Female  0.002969  0.002978
                  Male   0.003074  0.003025
```

```
[93]: mlt_df.stack().stack()
```

/tmp/ipykernel_2332935/2517862891.py:1: FutureWarning: The previous implementation of stack is deprecated and will be removed in a future version of pandas. See the What's New notes for pandas 2.1.0 for details. Specify

```
future_stack=True to adopt the new implementation and silence this warning.
mlt_df.stack().stack()
```

```
[93]: race      sex      income
Amer-Indian-Eskimo  Female  <=50K    0.001764
                                >50K    0.002395
                                Male    <=50K    0.002046
                                >50K    0.001954
Asian-Pac-Islander  Female  <=50K    0.002398
                                >50K    0.002305
                                Male    <=50K    0.002652
                                >50K    0.002762
Black               Female  <=50K    0.003454
                                >50K    0.003331
                                Male    <=50K    0.003922
                                >50K    0.003971
Other               Female  <=50K    0.002803
                                >50K    0.002593
                                Male    <=50K    0.003478
                                >50K    0.003310
White               Female  <=50K    0.002969
                                >50K    0.002978
                                Male    <=50K    0.003074
                                >50K    0.003025

dtype: float64
```

8 Pivot & Melt

```
[94]: wide_df = pd.read_csv('wide.csv')
wide_df
```

```
[94]:   ReadingDateTime  NO  NO2  NOX  PM10  PM2.5
0  01/01/2017 00:00  3.5  30.8  36.2  35.7  31.0
1  01/01/2017 01:00  3.6  31.5  37.0  28.5  31.0
2  01/01/2017 02:00  2.2  27.3  30.7  22.7  31.0
```

```
[95]: wide_df.melt(id_vars='ReadingDateTime',
                  value_vars=['NO', 'NO2', 'NOX', 'PM10', 'PM2.5'],
                  var_name='Species',
                  value_name='Value')
```

```
[95]:   ReadingDateTime Species  Value
0  01/01/2017 00:00      NO     3.5
1  01/01/2017 01:00      NO     3.6
2  01/01/2017 02:00      NO     2.2
3  01/01/2017 00:00     NO2    30.8
4  01/01/2017 01:00     NO2    31.5
```


5	01/01/2017 02:00	NO2	27.3
6	01/01/2017 00:00	NOX	36.2
7	01/01/2017 01:00	NOX	37.0
8	01/01/2017 02:00	NOX	30.7
9	01/01/2017 00:00	PM10	35.7
10	01/01/2017 01:00	PM10	28.5
11	01/01/2017 02:00	PM10	22.7
12	01/01/2017 00:00	PM2.5	31.0
13	01/01/2017 01:00	PM2.5	31.0
14	01/01/2017 02:00	PM2.5	31.0

```
[96]: long_df = pd.read_csv('long.csv')
long_df
```

```
[96]:
```

	ReadingDateTime	Species	Value
0	01/01/2017 00:00	NO	3.5
1	01/01/2017 01:00	NO	3.6
2	01/01/2017 02:00	NO	2.2
3	01/01/2017 00:00	NO2	30.8
4	01/01/2017 01:00	NO2	31.5
5	01/01/2017 02:00	NO2	27.3
6	01/01/2017 00:00	NOX	36.2
7	01/01/2017 01:00	NOX	37.0
8	01/01/2017 02:00	NOX	30.7
9	01/01/2017 00:00	PM10	35.7
10	01/01/2017 01:00	PM10	28.5
11	01/01/2017 02:00	PM10	22.7
12	01/01/2017 00:00	PM2.5	31.0
13	01/01/2017 01:00	PM2.5	31.0
14	01/01/2017 02:00	PM2.5	31.0

```
[97]: long_df.pivot(index='ReadingDateTime',
                      columns='Species',
                      values='Value')
```

```
[97]:
```

Species	NO	NO2	NOX	PM10	PM2.5
ReadingDateTime					
01/01/2017 00:00	3.5	30.8	36.2	35.7	31.0
01/01/2017 01:00	3.6	31.5	37.0	28.5	31.0
01/01/2017 02:00	2.2	27.3	30.7	22.7	31.0