

Excercises Ch1

April 19, 2025

1 Hands-On Data Preprocessing in Python

Learn how to effectively prepare data for successful data analytics

AUTHOR: Dr. Roy Jafari

1.0.1 Chapter 1: Review of the core modules NumPy, Pandas, and Matplotlib

Excercise 1 Use the adult.csv dataset and run the codes shown in the following Screenshots. Then answer the questions.

```
[1]: import pandas as pd
import numpy as np
adult_df = pd.read_csv('adult.csv')
adult_df.set_index(np.arange(10000,42561),inplace=True)

print(adult_df.iloc[5:7,0:2])

print(adult_df.loc['10005':'10007','age':'fnlwgt'])
```

```
      age workclass
10005   37   Private
10006   49   Private
      age      workclass fnlwgt
10005   37      Private 284582
10006   49      Private 160187
10007   52 Self-emp-not-inc 209642
```

- a) Use the output to answer what is the difference in the behavior of .loc and .iloc when it comes to indexing.
- b) Without running but by only looking at the data, what will be the output of adult_df.loc['10005':'10007', 'age']?
- c) Without running but by only looking at the data, what will be the output of adult_df.iloc[5:7, 0:2]?

a) iloc's lower bound is exclusive, loc's lower bound is inclusive

b) relationship race sex

10000 Not-in-family White Male 10001 Husband White Male 10002 Not-in-family White Male
10003 Husband Black Male

c) relationship race

10000 Not-in-family White 10001 Husband White 10002 Not-in-family White

Excercise 2 For `adult_df` use the `.groupby()` function to run the following code and create the multi-index Series `mlt_sr`.

```
[2]: import pandas as pd

adult_df = pd.read_csv('adult.csv')
mlt_seris =adult_df.groupby(['race','sex','income']).fnlwgt.mean()
mlt_seris
```

```
[2]: race          sex      income
Amer-Indian-Eskimo  Female  <=50K    109018.626168
                             >50K    148012.000000
                             Male    <=50K    126428.767857
                             >50K    120721.541667
Asian-Pac-Islander  Female  <=50K    148165.333333
                             >50K    142426.093023
                             Male    <=50K    163885.130435
                             >50K    170698.347639
Black               Female  <=50K    213411.580205
                             >50K    205806.033333
                             Male    <=50K    242343.990566
                             >50K    245390.356902
Other               Female  <=50K    173235.932039
                             >50K    160223.333333
                             Male    <=50K    214891.832168
                             >50K    204551.736842
White               Female  <=50K    183485.120961
                             >50K    184030.255837
                             Male    <=50K    189950.465418
                             >50K    186917.770077

Name: fnlwgt, dtype: float64
```

Now that you have created a multi-index Series, run the following codes, study the outputs, and answer the questions.

a) Run the code below first and then answer the following. When we use `.iloc[]` for a multi-index

```
[3]: print(mlt_seris.iloc[0])
      print(mlt_seris.iloc[1])
      print(mlt_seris.iloc[2])
```

```
109018.6261682243
148012.0
126428.76785714286
```

Answer:

`.iloc[]` selects rows and columns by their integer positions

b) Run the code below first and then answer the following. When we use `.loc[]` to access the data

```
[5]: mlt_seris.loc['Other']
```

```
[5]: sex      income
     Female  <=50K      173235.932039
           >50K      160223.333333
     Male    <=50K      214891.832168
           >50K      204551.736842
     Name: fnlwgt, dtype: float64
```

Answer:

.loc[] selects data based on index labels, not positions.

c) Run the codes below first and then answer the following. When we use .loc[] to access the

When you run either line of code below, you will get an error and that is the point of this qu

```
[6]: mlt_seris.loc['Female']
```

```
-----
KeyError                                Traceback (most recent call last)
File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/
  ↳ data_preprocessing/lib/python3.11/site-packages/pandas/core/indexes/base.py:
  ↳ 3805, in Index.get_loc(self, key)
    3804 try:
-> 3805     return self._engine.get_loc(casted_key)
    3806 except KeyError as err:
```

```
File index.pyx:167, in pandas._libs.index.IndexEngine.get_loc()
```

```
File index.pyx:196, in pandas._libs.index.IndexEngine.get_loc()
```

```
File pandas/_libs/hashtable_class_helper.pxi:7081, in pandas._libs.hashtable.
  ↳ PyObjectHashTable.get_item()
```

```
File pandas/_libs/hashtable_class_helper.pxi:7089, in pandas._libs.hashtable.
  ↳ PyObjectHashTable.get_item()
```

```
KeyError: 'Female'
```

The above exception was the direct cause of the following exception:

```
KeyError                                Traceback (most recent call last)
Cell In[6], line 1
----> 1 mlt_seris.loc[ ]
```

```
File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/
  ↳ data_preprocessing/lib/python3.11/site-packages/pandas/core/indexing.py:1191,
  ↳ in _iLocIndexer._getitem__(self, key)
```

```

1189 maybe_callable = com.apply_if_callable(key, self.obj)
1190 maybe_callable = self._check_deprecated_callable_usage(key,
↳ maybe_callable)
-> 1191 return self._getitem_axis(maybe_callable, axis=axis)

```

```

File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/
↳ data_preprocessing/lib/python3.11/site-packages/pandas/core/indexing.py:1431,
↳ in _LocIndexer._getitem_axis(self, key, axis)
1429 # fall thru to straight lookup
1430 self._validate_key(key, axis)
-> 1431 return self._get_label(key, axis=axis)

```

```

File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/
↳ data_preprocessing/lib/python3.11/site-packages/pandas/core/indexing.py:1381,
↳ in _LocIndexer._get_label(self, label, axis)
1379 def _get_label(self, label, axis: AxisInt):
1380     # GH#5567 this will fail if the label is not present in the axis.
-> 1381     return self.obj.xs(label, axis=axis)

```

```

File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/
↳ data_preprocessing/lib/python3.11/site-packages/pandas/core/generic.py:4293,
↳ in NDFrame.xs(self, key, axis, level, drop_level)
4290     index = self.index
4292 if isinstance(index, MultiIndex):
-> 4293     loc, new_index = index._get_loc_level(key, level=0)
4294     if not drop_level:
4295         if lib.is_integer(loc):
4296             # Slice index must be an integer or None

```

```

File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/
↳ data_preprocessing/lib/python3.11/site-packages/pandas/core/indexes/multi.py:
↳ 3290, in MultiIndex._get_loc_level(self, key, level)
3288     return indexer, maybe_mi_droplevels(indexer, ilevels)
3289 else:
-> 3290     indexer = self._get_level_indexer(key, level=level)
3291     if (
3292         isinstance(key, str)
3293         and self.levels[level]._supports_partial_string_indexing
3294     ):
3295         # check to see if we did an exact lookup vs sliced
3296         check = self.levels[level].get_loc(key)

```

```

File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/
↳ data_preprocessing/lib/python3.11/site-packages/pandas/core/indexes/multi.py:
↳ 3391, in MultiIndex._get_level_indexer(self, key, level, indexer)
3388     return slice(i, j, step)
3390 else:
-> 3391     idx = self._get_loc_single_level_index(level_index, key)
3393     if level > 0 or self._lexsort_depth == 0:

```

```

3394         # Desired level is not sorted
3395         if isinstance(idx, slice):
3396             # test_get_loc_partial_timestamp_multiindex

File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/
↳data_preprocessing/lib/python3.11/site-packages/pandas/core/indexes/multi.py:
↳2980, in MultiIndex._get_loc_single_level_index(self, level_index, key)
2978     return -1
2979 else:
-> 2980     return level_index.get_loc(key)

File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/
↳data_preprocessing/lib/python3.11/site-packages/pandas/core/indexes/base.py:
↳3812, in Index.get_loc(self, key)
3807     if isinstance(casted_key, slice) or (
3808         isinstance(casted_key, abc.Iterable)
3809         and any(isinstance(x, slice) for x in casted_key)
3810     ):
3811         raise InvalidIndexError(key)
-> 3812     raise KeyError(key) from err
3813 except TypeError:
3814     # If we have a listlike key, _check_indexing_error will raise
3815     # InvalidIndexError. Otherwise we fall through and re-raise
3816     # the TypeError.
3817     self._check_indexing_error(key)

KeyError: 'Female'

```

```
[7]: mlt_seris.loc['<=50K']
```

```

-----
KeyError                                Traceback (most recent call last)
File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/
↳data_preprocessing/lib/python3.11/site-packages/pandas/core/indexes/base.py:
↳3805, in Index.get_loc(self, key)
3804 try:
-> 3805     return self._engine.get_loc(casted_key)
3806 except KeyError as err:

File index.pyx:167, in pandas._libs.index.IndexEngine.get_loc()

File index.pyx:196, in pandas._libs.index.IndexEngine.get_loc()

File pandas/_libs/hashtable_class_helper.pxi:7081, in pandas._libs.hashtable.
↳PyObjectHashTable.get_item()

```

```
File pandas/_libs/hashtable_class_helper.pxi:7089, in pandas._libs.hashtable.  
↳PyObjectHashTable.get_item()
```

```
KeyError: '<=50K'
```

The above exception was the direct cause of the following exception:

```
KeyError                                Traceback (most recent call last)
```

```
Cell In[7], line 1
```

```
----> 1 mlt_seris.loc[      ]
```

```
File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/  
↳data_preprocessing/lib/python3.11/site-packages/pandas/core/indexing.py:1191,  
↳in _LocationIndexer.__getitem__(self, key)  
    1189 maybe_callable = com.apply_if_callable(key, self.obj)  
    1190 maybe_callable = self._check_deprecated_callable_usage(key,  
↳maybe_callable)  
-> 1191 return self._getitem_axis(maybe_callable, axis=axis)
```

```
File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/  
↳data_preprocessing/lib/python3.11/site-packages/pandas/core/indexing.py:1431,  
↳in _iLocIndexer._getitem_axis(self, key, axis)  
    1429 # fall thru to straight lookup  
    1430 self._validate_key(key, axis)  
-> 1431 return self._get_label(key, axis=axis)
```

```
File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/  
↳data_preprocessing/lib/python3.11/site-packages/pandas/core/indexing.py:1381,  
↳in _iLocIndexer._get_label(self, label, axis)  
    1379 def _get_label(self, label, axis: AxisInt):  
    1380     # GH#5567 this will fail if the label is not present in the axis.  
-> 1381     return self.obj.xs(label, axis=axis)
```

```
File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/  
↳data_preprocessing/lib/python3.11/site-packages/pandas/core/generic.py:4293,  
↳in NDFrame.xs(self, key, axis, level, drop_level)  
    4290     index = self.index  
    4292 if isinstance(index, MultiIndex):  
-> 4293     loc, new_index = index._get_loc_level(key, level=0)  
    4294     if not drop_level:  
    4295         if lib.is_integer(loc):  
    4296             # Slice index must be an integer or None
```

```
File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/  
↳data_preprocessing/lib/python3.11/site-packages/pandas/core/indexes/multi.py:  
↳3290, in MultiIndex._get_loc_level(self, key, level)  
    3288     return indexer, maybe_mi_droplevels(indexer, ilevels)  
    3289 else:  
-> 3290     indexer = self._get_level_indexer(key, level=level)
```

```

3291     if (
3292         isinstance(key, str)
3293         and self.levels[level]._supports_partial_string_indexing
3294     ):
3295         # check to see if we did an exact lookup vs sliced
3296         check = self.levels[level].get_loc(key)

File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/
↳data_preprocessing/lib/python3.11/site-packages/pandas/core/indexes/multi.py:
↳3391, in MultiIndex._get_level_indexer(self, key, level, indexer)
3388     return slice(i, j, step)
3390 else:
-> 3391     idx = self._get_loc_single_level_index(level_index, key)
3393     if level > 0 or self._lexsort_depth == 0:
3394         # Desired level is not sorted
3395         if isinstance(idx, slice):
3396             # test_get_loc_partial_timestamp_multiindex

File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/
↳data_preprocessing/lib/python3.11/site-packages/pandas/core/indexes/multi.py:
↳2980, in MultiIndex._get_loc_single_level_index(self, level_index, key)
2978     return -1
2979 else:
-> 2980     return level_index.get_loc(key)

File /run/media/weida/SSD1/Study Stuff/Hands-On Data Preprocessing in Python/
↳data_preprocessing/lib/python3.11/site-packages/pandas/core/indexes/base.py:
↳3812, in Index.get_loc(self, key)
3807     if isinstance(casted_key, slice) or (
3808         isinstance(casted_key, abc.Iterable)
3809         and any(isinstance(x, slice) for x in casted_key)
3810     ):
3811         raise InvalidIndexError(key)
-> 3812     raise KeyError(key) from err
3813 except TypeError:
3814     # If we have a listlike key, _check_indexing_error will raise
3815     # InvalidIndexError. Otherwise we fall through and re-raise
3816     # the TypeError.
3817     self._check_indexing_error(key)

KeyError: '<=50K'

```

Answer:

error

d) Run the codes below first and then answer the following. How is using `.loc[]` or `.iloc[]` di

```
[9]: print(mlt_seris.loc['Other']['Female']['<=50K'])
print(mlt_seris.iloc[12])
```

```
173235.93203883496
173235.93203883496
```

Answer:

Using `.loc[]` is label-based (selects by index names at any level), while `.iloc[]` is position-based (ignores index labels and selects by integer position).

Exercise 3 For this exercise you need to use a new dataset: `billboard.csv`. Visit <https://www.billboard.com/charts/hot-100> and see the latest song rankings of the day. This dataset presents information and ranking of 317 song tracks in 80 columns. The first four columns are artist, track, time, and `date_e`. The first columns are intuitive descriptions of song tracks. The column `date_e` shows the date that the songs entered the hot-100 list. The rest of 76 columns are songs ranking at the end of each weeks from 'w1' to 'w76'. Download and read this dataset using pandas and answer the following questions.

```
[12]: import pandas as pd

billboard_df = pd.read_csv("billboard.csv")
billboard_df.head()
```

```
[12]:
```

	artist	track	time	\
0	Destiny's Child	Independent Women Part I	3:38	
1	Santana	Maria, Maria	4:18	
2	Savage Garden	I Knew I Loved You	4:07	
3	Madonna	Music	3:45	
4	Aguilera, Christina	Come On Over Baby (All I Want Is You)	3:38	

	date_e	w1	w2	w3	w4	w5	w6	...	w67	w68	w69	w70	w71	\
0	2000-09-23	78	63.0	49.0	33.0	23.0	15.0	...	NaN	NaN	NaN	NaN	NaN	
1	2000-02-12	15	8.0	6.0	5.0	2.0	3.0	...	NaN	NaN	NaN	NaN	NaN	
2	1999-10-23	71	48.0	43.0	31.0	20.0	13.0	...	NaN	NaN	NaN	NaN	NaN	
3	2000-08-12	41	23.0	18.0	14.0	2.0	1.0	...	NaN	NaN	NaN	NaN	NaN	
4	2000-08-05	57	47.0	45.0	29.0	23.0	18.0	...	NaN	NaN	NaN	NaN	NaN	

	w72	w73	w74	w75	w76
0	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN

[5 rows x 80 columns]

a) Write one line of code that gives you a great idea of how many null values each column has


```
[17]: print(billboard_df.isnull().sum())
billboard_df = billboard_df.dropna(axis=1, how='all')
```

```
artist      0
track       0
time        0
date_e      0
w1          0
...
w72         317
w73         317
w74         317
w75         317
w76         317
Length: 80, dtype: int64
```

```
[19]: print(billboard_df.isnull().sum())
billboard_df
```

```
artist      0
track       0
time        0
date_e      0
w1          0
...
w61         315
w62         315
w63         315
w64         315
w65         316
Length: 69, dtype: int64
```

```
[19]:
```

	artist	track	time	\
0	Destiny's Child	Independent Women Part I	3:38	
1	Santana	Maria, Maria	4:18	
2	Savage Garden	I Knew I Loved You	4:07	
3	Madonna	Music	3:45	
4	Aguilera, Christina	Come On Over Baby (All I Want Is You)	3:38	
..	
312	Ghostface Killah	Cherchez LaGhost	3:04	
313	Smith, Will	Freakin' It	3:58	
314	Zombie Nation	Kernkraft 400	3:30	
315	Eastsidaz, The	Got Beef	3:58	
316	Fragma	Toca's Miracle	3:22	

	date_e	w1	w2	w3	w4	w5	w6	...	w56	w57	w58	w59	\
0	2000-09-23	78	63.0	49.0	33.0	23.0	15.0	...	NaN	NaN	NaN	NaN	
1	2000-02-12	15	8.0	6.0	5.0	2.0	3.0	...	NaN	NaN	NaN	NaN	

2	1999-10-23	71	48.0	43.0	31.0	20.0	13.0	...	NaN	NaN	NaN	NaN
3	2000-08-12	41	23.0	18.0	14.0	2.0	1.0	...	NaN	NaN	NaN	NaN
4	2000-08-05	57	47.0	45.0	29.0	23.0	18.0	...	NaN	NaN	NaN	NaN
..
312	2000-08-05	98	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN
313	2000-02-12	99	99.0	99.0	99.0	NaN	NaN	...	NaN	NaN	NaN	NaN
314	2000-09-02	99	99.0	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN
315	2000-07-01	99	99.0	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN
316	2000-10-28	99	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN

	w60	w61	w62	w63	w64	w65
0	NaN	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN	NaN
..
312	NaN	NaN	NaN	NaN	NaN	NaN
313	NaN	NaN	NaN	NaN	NaN	NaN
314	NaN	NaN	NaN	NaN	NaN	NaN
315	NaN	NaN	NaN	NaN	NaN	NaN
316	NaN	NaN	NaN	NaN	NaN	NaN


[317 rows x 69 columns]

b) With a for loop, draw and study the values in each of the remaining W columns.


```
[22]: import matplotlib.pyplot as plt

for col in billboard_df.columns[billboard_df.columns.str.startswith('w')]:
    print(f"\n--- Column: {col} ---")
    print(billboard_df[col].describe())
    billboard_df[col].plot.hist(title=col)
    plt.show()
```


```
--- Column: w1 ---
count      317.000000
mean        79.958991
std         14.686865
min         15.000000
25%         74.000000
50%         81.000000
75%         91.000000
max         100.000000
Name: w1, dtype: float64
```




```
--- Column: w2 ---
count    312.000000
mean      71.173077
std       18.200443
min        8.000000
25%       63.000000
50%       73.000000
75%       84.000000
max      100.000000
Name: w2, dtype: float64
```




```
--- Column: w3 ---  
count    307.000000  
mean      65.045603  
std       20.752302  
min        6.000000  
25%       53.000000  
50%       66.000000  
75%       79.000000  
max      100.000000  
Name: w3, dtype: float64
```




```
--- Column: w4 ---  
count    300.000000  
mean      59.763333  
std       22.324619  
min        5.000000  
25%       44.750000  
50%       61.000000  
75%       76.000000  
max      100.000000  
Name: w4, dtype: float64
```




```
--- Column: w5 ---  
count    292.000000  
mean      56.339041  
std       23.780022  
min        2.000000  
25%       38.750000  
50%       57.000000  
75%       73.250000  
max      100.000000  
Name: w5, dtype: float64
```




```
--- Column: w6 ---  
count    280.000000  
mean      52.360714  
std       24.473273  
min        1.000000  
25%       33.750000  
50%       51.500000  
75%       72.250000  
max       99.000000  
Name: w6, dtype: float64
```




```
--- Column: w7 ---  
count    269.000000  
mean      49.219331  
std       25.654279  
min        1.000000  
25%       30.000000  
50%       47.000000  
75%       67.000000  
max      100.000000  
Name: w7, dtype: float64
```

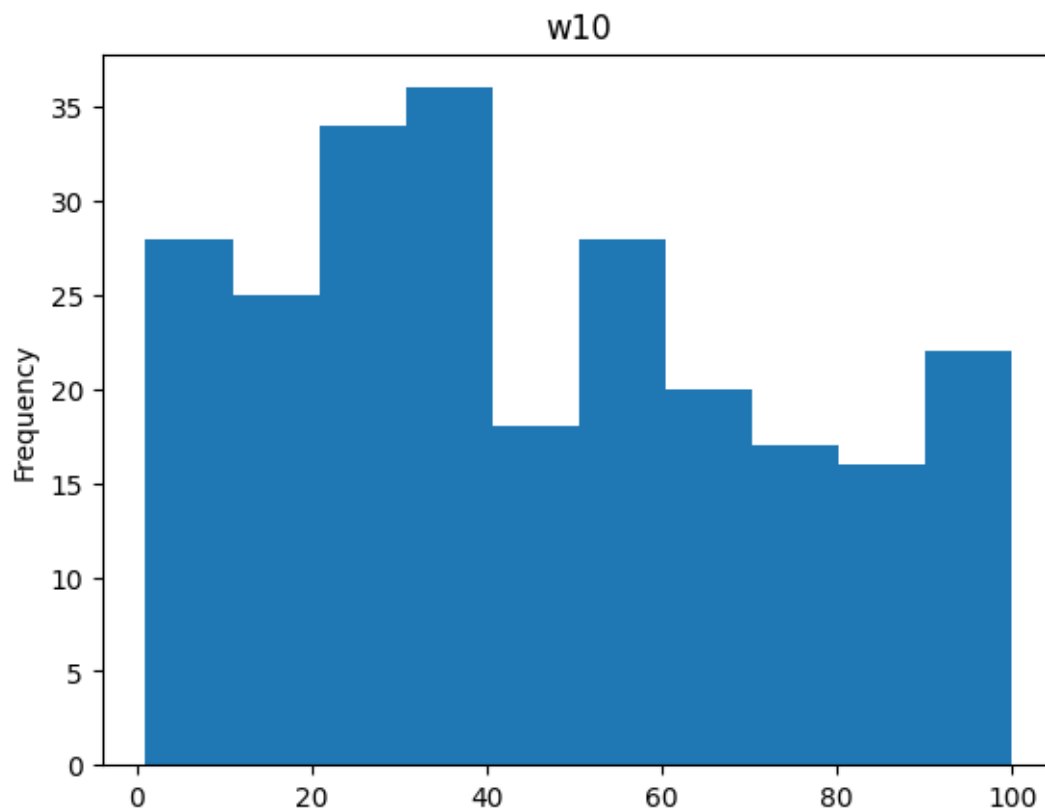
```
--- Column: w8 ---  
count    260.000000  
mean      47.119231  
std       26.370782  
min        1.000000  
25%       27.000000  
50%       45.500000  
75%       67.000000  
max       99.000000  
Name: w8, dtype: float64
```



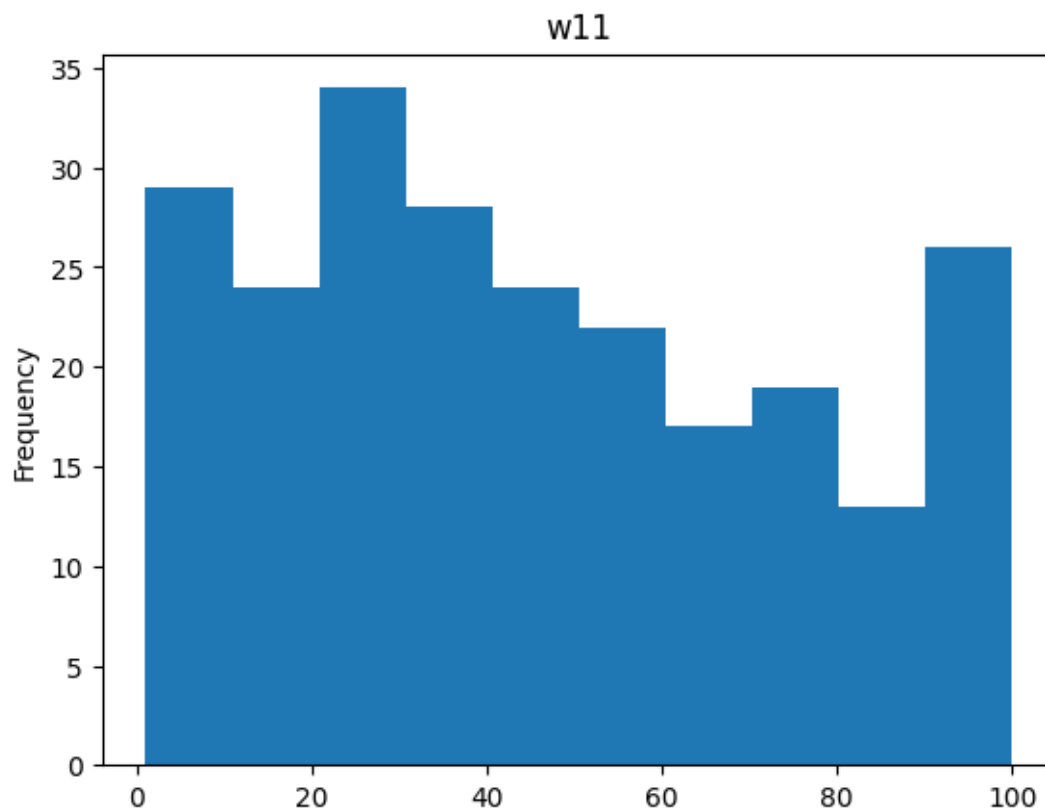
```
--- Column: w9 ---  
count    253.000000  
mean      46.343874  
std       27.136419  
min        1.000000  
25%       26.000000  
50%       42.000000  
75%       67.000000  
max      100.000000  
Name: w9, dtype: float64
```




```
--- Column: w10 ---
count    244.000000
mean      45.786885
std       28.152357
min        1.000000
25%       24.750000
50%       40.000000
75%       69.000000
max      100.000000
Name: w10, dtype: float64
```




```
--- Column: w11 ---  
count    236.000000  
mean      45.474576  
std       29.060527  
min        1.000000  
25%       22.000000  
50%       42.500000  
75%       69.250000  
max      100.000000  
Name: w11, dtype: float64
```




```
--- Column: w12 ---  
count    222.000000  
mean      44.198198  
std       28.893100  
min        1.000000  
25%       20.250000  
50%       42.000000  
75%       67.750000  
max      100.000000  
Name: w12, dtype: float64
```




```
--- Column: w13 ---
count    210.000000
mean      44.352381
std       29.481315
min        1.000000
25%       19.000000
50%       42.500000
75%       68.750000
max      100.000000
Name: w13, dtype: float64
```




```
--- Column: w14 ---  
count    204.000000  
mean      45.892157  
std       30.529936  
min        1.000000  
25%       20.000000  
50%       44.000000  
75%       72.500000  
max      100.000000  
Name: w14, dtype: float64
```




```
--- Column: w15 ---
count    197.000000
mean      47.898477
std       31.542486
min        1.000000
25%       18.000000
50%       46.000000
75%       76.000000
max      100.000000
Name: w15, dtype: float64
```


```
--- Column: w16 ---  
count    182.000000  
mean      48.241758  
std       31.899476  
min        1.000000  
25%       17.000000  
50%       47.000000  
75%       76.000000  
max      100.000000  
Name: w16, dtype: float64
```




```
--- Column: w17 ---  
count    177.000000  
mean      51.265537  
std       33.093233  
min        1.000000  
25%       19.000000  
50%       52.000000  
75%       85.000000  
max      100.000000  
Name: w17, dtype: float64
```




```
--- Column: w18 ---
count    166.000000
mean      52.108434
std       33.194199
min        1.000000
25%       19.000000
50%       56.500000
75%       84.000000
max      100.000000
Name: w18, dtype: float64
```




```
--- Column: w19 ---  
count    156.000000  
mean      53.166667  
std       33.022540  
min        1.000000  
25%       20.750000  
50%       55.500000  
75%       85.000000  
max       99.000000  
Name: w19, dtype: float64
```




```
--- Column: w20 ---  
count    146.000000  
mean      54.267123  
std       32.890475  
min        2.000000  
25%       22.250000  
50%       58.500000  
75%       87.000000  
max      100.000000  
Name: w20, dtype: float64
```




```
--- Column: w21 ---
count    65.000000
mean     26.184615
std      20.232162
min       3.000000
25%      10.000000
50%      24.000000
75%      35.000000
max      97.000000
Name: w21, dtype: float64
```




```
--- Column: w22 ---
count      55.000000
mean       25.636364
std        21.127798
min         3.000000
25%         9.000000
50%        21.000000
75%        36.500000
max        100.000000
Name: w22, dtype: float64
```




```
--- Column: w23 ---
count    48.00000
mean     23.81250
std      18.23564
min       3.00000
25%      10.00000
50%      20.50000
75%      36.00000
max      91.00000
Name: w23, dtype: float64
```

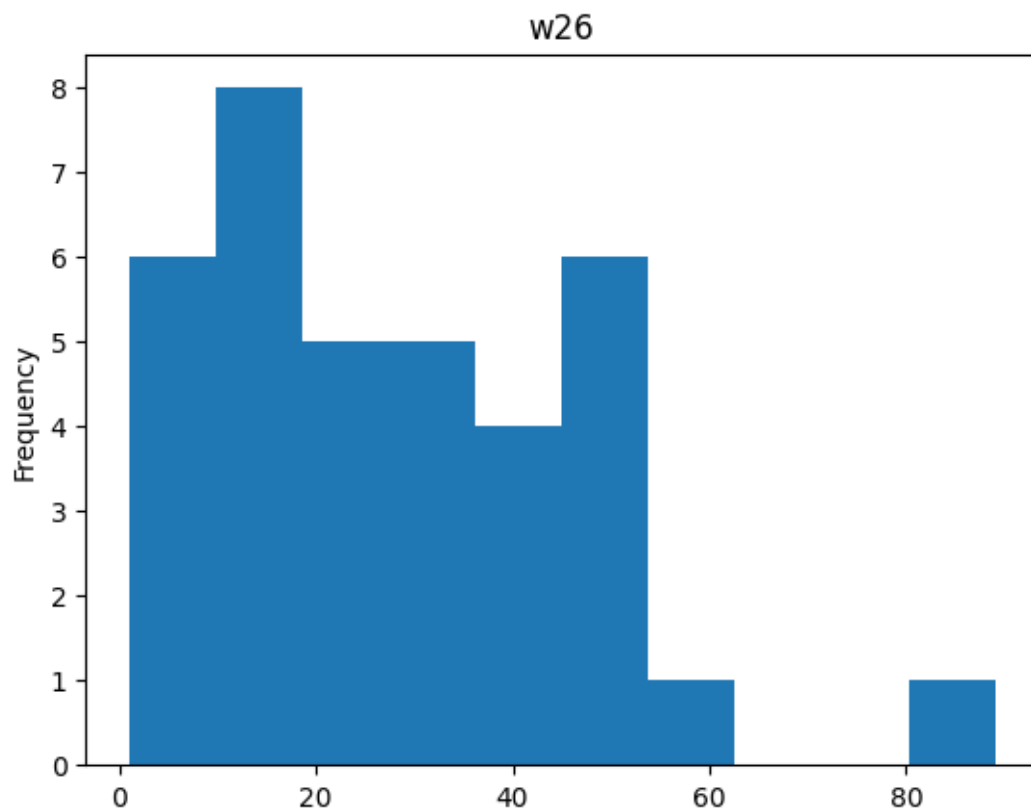
```
--- Column: w24 ---  
count    46.000000  
mean     26.782609  
std      18.556716  
min       3.000000  
25%      12.000000  
50%      21.500000  
75%      40.500000  
max      91.000000  
Name: w24, dtype: float64
```




```
--- Column: w25 ---  
count    38.000000  
mean     26.131579  
std      18.737404  
min       2.000000  
25%      12.250000  
50%      22.500000  
75%      38.250000  
max      90.000000  
Name: w25, dtype: float64
```




```
--- Column: w26 ---  
count    36.00000  
mean     28.00000  
std      19.14904  
min       1.00000  
25%      13.75000  
50%      26.00000  
75%      40.25000  
max       89.00000  
Name: w26, dtype: float64
```




```
--- Column: w27 ---
count    29.000000
mean     27.344828
std      19.736254
min       1.000000
25%      12.000000
50%      26.000000
75%      38.000000
max       86.000000
Name: w27, dtype: float64
```




```
--- Column: w28 ---
count    24.000000
mean     25.000000
std      16.229335
min       2.000000
25%      14.500000
50%      23.500000
75%      37.500000
max      58.000000
Name: w28, dtype: float64
```




```
--- Column: w29 ---  
count    20.000000  
mean     23.650000  
std      15.239405  
min       2.000000  
25%      13.750000  
50%      22.500000  
75%      36.000000  
max      49.000000  
Name: w29, dtype: float64
```




```
--- Column: w30 ---  
count    20.000000  
mean     26.400000  
std      15.104548  
min       2.000000  
25%      16.000000  
50%      31.500000  
75%      38.000000  
max      45.000000  
Name: w30, dtype: float64
```




```
--- Column: w31 ---  
count    19.000000  
mean     28.736842  
std      16.002924  
min       3.000000  
25%      16.000000  
50%      34.000000  
75%      41.000000  
max      49.000000  
Name: w31, dtype: float64
```


```
--- Column: w32 ---  
count    18.000000  
mean     30.277778  
std      17.032169  
min       3.000000  
25%      14.500000  
50%      38.000000  
75%      44.000000  
max      47.000000  
Name: w32, dtype: float64
```




```
--- Column: w33 ---
count    12.000000
mean     23.166667
std      17.071152
min       2.000000
25%       5.500000
50%      24.000000
75%      36.250000
max      50.000000
Name: w33, dtype: float64
```




```
--- Column: w34 ---  
count    10.000000  
mean     19.300000  
std      14.929836  
min       3.000000  
25%       5.750000  
50%      18.500000  
75%      28.250000  
max      49.000000  
Name: w34, dtype: float64
```




```
--- Column: w35 ---
count      9.000000
mean       16.888889
std        12.333333
min         4.000000
25%         5.000000
50%        14.000000
75%        27.000000
max        34.000000
Name: w35, dtype: float64
```




```
--- Column: w36 ---  
count      9.000000  
mean       18.555556  
std        13.893444  
min         5.000000  
25%         6.000000  
50%        16.000000  
75%        29.000000  
max        41.000000  
Name: w36, dtype: float64
```




```
--- Column: w37 ---
count      9.000000
mean       20.888889
std        15.799086
min         5.000000
25%         9.000000
50%        16.000000
75%        36.000000
max        48.000000
Name: w37, dtype: float64
```




```
--- Column: w38 ---  
count      8.000000  
mean       19.375000  
std        12.693502  
min         9.000000  
25%         9.000000  
50%        16.000000  
75%        23.000000  
max        40.000000  
Name: w38, dtype: float64
```




```
--- Column: w39 ---  
count      8.000000  
mean       20.250000  
std        18.289341  
min         3.000000  
25%         8.750000  
50%        13.000000  
75%        25.500000  
max        50.000000  
Name: w39, dtype: float64
```


```
--- Column: w40 ---  
count      7.000000  
mean       17.857143  
std        14.264508  
min         1.000000  
25%        13.000000  
50%        14.000000  
75%        18.500000  
max        47.000000  
Name: w40, dtype: float64
```




```
--- Column: w41 ---
count      7.000000
mean       18.571429
std        15.306395
min         1.000000
25%        13.000000
50%        14.000000
75%        19.500000
max         50.000000
Name: w41, dtype: float64
```




```
--- Column: w42 ---
count      6.000000
mean       14.166667
std        7.250287
min         2.000000
25%        11.500000
50%        15.000000
75%        19.250000
max        22.000000
Name: w42, dtype: float64
```




```
--- Column: w43 ---
count      6.000000
mean       17.666667
std        10.366613
min         3.000000
25%        11.750000
50%        18.000000
75%        23.500000
max        32.000000
Name: w43, dtype: float64
```




```
--- Column: w44 ---  
count      6.000000  
mean       21.666667  
std        15.081998  
min         4.000000  
25%        10.750000  
50%        21.000000  
75%        29.000000  
max        45.000000  
Name: w44, dtype: float64
```




```
--- Column: w45 ---
count      5.000000
mean       17.800000
std        11.691878
min         4.000000
25%         9.000000
50%        17.000000
75%        28.000000
max        31.000000
Name: w45, dtype: float64
```




```
--- Column: w46 ---  
count      5.000000  
mean       20.200000  
std        14.184499  
min         5.000000  
25%         7.000000  
50%        21.000000  
75%        31.000000  
max        37.000000  
Name: w46, dtype: float64
```




```
--- Column: w47 ---  
count      5.000000  
mean       22.800000  
std        16.543881  
min         6.000000  
25%         7.000000  
50%        22.000000  
75%        38.000000  
max        41.000000  
Name: w47, dtype: float64
```


```
--- Column: w48 ---  
count      4.000000  
mean       20.000000  
std        14.628739  
min         8.000000  
25%         8.750000  
50%        16.500000  
75%        27.750000  
max        39.000000  
Name: w48, dtype: float64
```




```
--- Column: w49 ---  
count      4.000000  
mean       22.000000  
std        16.062378  
min         9.000000  
25%         9.000000  
50%        18.500000  
75%        31.500000  
max        42.000000  
Name: w49, dtype: float64
```




```
--- Column: w50 ---
count      4.000000
mean       26.000000
std        18.529256
min        10.000000
25%        11.500000
50%        22.500000
75%        37.000000
max        49.000000
Name: w50, dtype: float64
```




```
--- Column: w51 ---
count      4.000000
mean       29.250000
std        18.997807
min        12.000000
25%        13.500000
50%        28.000000
75%        43.750000
max        49.000000
Name: w51, dtype: float64
```




```
--- Column: w52 ---
count      4.000000
mean       30.250000
std        17.211914
min        15.000000
25%        15.750000
50%        29.000000
75%        43.500000
max        48.000000
Name: w52, dtype: float64
```




```
--- Column: w53 ---
count      4.000000
mean       33.250000
std        17.095321
min        17.000000
25%        19.250000
50%        33.500000
75%        47.500000
max        49.000000
Name: w53, dtype: float64
```




```
--- Column: w54 ---  
count      2.000000  
mean       19.500000  
std        3.535534  
min        17.000000  
25%        18.250000  
50%        19.500000  
75%        20.750000  
max        22.000000  
Name: w54, dtype: float64
```




```
--- Column: w55 ---
count      2.000000
mean       21.500000
std        0.707107
min        21.000000
25%        21.250000
50%        21.500000
75%        21.750000
max        22.000000
Name: w55, dtype: float64
```


```
--- Column: w56 ---  
count      2.000000  
mean       25.500000  
std        0.707107  
min        25.000000  
25%        25.250000  
50%        25.500000  
75%        25.750000  
max        26.000000  
Name: w56, dtype: float64
```




```
--- Column: w57 ---
count      2.00000
mean       27.50000
std        2.12132
min        26.00000
25%        26.75000
50%        27.50000
75%        28.25000
max        29.00000
Name: w57, dtype: float64
```




```
--- Column: w58 ---  
count      2.000000  
mean       31.500000  
std        0.707107  
min        31.000000  
25%        31.250000  
50%        31.500000  
75%        31.750000  
max        32.000000  
Name: w58, dtype: float64
```




```
--- Column: w59 ---  
count      2.000000  
mean       35.500000  
std        4.949747  
min        32.000000  
25%        33.750000  
50%        35.500000  
75%        37.250000  
max        39.000000  
Name: w59, dtype: float64
```




```
--- Column: w60 ---  
count      2.000000  
mean       38.000000  
std        1.414214  
min        37.000000  
25%        37.500000  
50%        38.000000  
75%        38.500000  
max        39.000000  
Name: w60, dtype: float64
```




```
--- Column: w61 ---  
count      2.000000  
mean       42.500000  
std        0.707107  
min        42.000000  
25%        42.250000  
50%        42.500000  
75%        42.750000  
max        43.000000  
Name: w61, dtype: float64
```




```
--- Column: w62 ---  
count      2.000000  
mean       44.500000  
std        3.535534  
min        42.000000  
25%        43.250000  
50%        44.500000  
75%        45.750000  
max        47.000000  
Name: w62, dtype: float64
```




```
--- Column: w63 ---
count      2.000000
mean       47.500000
std        3.535534
min        45.000000
25%        46.250000
50%        47.500000
75%        48.750000
max        50.000000
Name: w63, dtype: float64
```

```
--- Column: w64 ---  
count      2.0  
mean       50.0  
std        0.0  
min        50.0  
25%        50.0  
50%        50.0  
75%        50.0  
max        50.0  
Name: w64, dtype: float64
```



```
--- Column: w65 ---
count      1.0
mean       49.0
std        NaN
min        49.0
25%        49.0
50%        49.0
75%        49.0
max        49.0
Name: w65, dtype: float64
```



c) The dataset is in wide format, use an appropriate function to switch to a long format, and

```
[26]: mlt_df = billboard_df.melt(
      id_vars=["artist", "track", "time", "date_e"],
      value_vars=billboard_df.columns[billboard_df.columns.str.startswith('w')],
      var_name='W',
      value_name='Value'
    )
mlt_df
```

```
[26]:
```

	artist	track	time	\
0	Destiny's Child	Independent Women Part I	3:38	
1	Santana	Maria, Maria	4:18	
2	Savage Garden	I Knew I Loved You	4:07	
3	Madonna	Music	3:45	
4	Aguilera, Christina	Come On Over Baby (All I Want Is You)	3:38	
...	
20600	Ghostface Killah	Cherchez LaGhost	3:04	
20601	Smith, Will	Freakin' It	3:58	
20602	Zombie Nation	Kernkraft 400	3:30	
20603	Eastsidaz, The	Got Beef	3:58	

20604 Fragma Toca's Miracle 3:22

	date_e	W	Value
0	2000-09-23	w1	78.0
1	2000-02-12	w1	15.0
2	1999-10-23	w1	71.0
3	2000-08-12	w1	41.0
4	2000-08-05	w1	57.0
...
20600	2000-08-05	w65	NaN
20601	2000-02-12	w65	NaN
20602	2000-09-02	w65	NaN
20603	2000-07-01	w65	NaN
20604	2000-10-28	w65	NaN

[20605 rows x 6 columns]

d) Write a code that shows mlt_df's every 1200 rows.

```
[27]: mlt_df[0:20605:1200]
```

```
[27]:
```

	artist	track	time	date_e	\
0	Destiny's Child	Independent Women Part I	3:38	2000-09-23	
1200	Limp Bizkit	N 2 Gether Now	3:55	1999-12-04	
2400	Urban, Keith	Your Everything	4:10	2000-07-15	
3600	Mary Mary	Shackles (Praise You)	3:12	2000-03-25	
4800	Pink	There U Go	3:23	2000-03-04	
6000	Levert, Gerald	Baby U Are	4:10	2000-08-19	
7200	Jay-Z	Do It Again (Put Ya Hands Up)	3:47	2000-01-15	
8400	Tippin, Aaron	Kiss This	2:53	2000-08-26	
9600	Rimes, LeAnn	Big Deal	3:03	1999-10-16	
10800	98;	Give Me Just One Night (Una Noche)	3:24	2000-08-19	
12000	Estefan, Gloria	No Me Dejes De Querer	3:25	2000-06-10	
13200	Diffie, Joe	It's Always Somethin'	2:55	2000-08-12	
14400	Aaliyah	I Don't Wanna	4:15	2000-01-29	
15600	Barenaked Ladies	Pinch Me	3:46	2000-09-09	
16800	Fragma	Toca's Miracle	3:22	2000-10-28	
18000	Drama	Left, Right, Left	3:37	2000-02-12	
19200	Lopez, Jennifer	Feelin' Good	4:28	2000-02-19	
20400	Counting Crows	Hanginaround	4:07	1999-11-06	

	W	Value
0	w1	78.0
1200	w4	78.0
2400	w8	59.0
3600	w12	39.0
4800	w16	14.0
6000	w19	NaN

7200	w23	NaN
8400	w27	NaN
9600	w31	NaN
10800	w35	NaN
12000	w38	NaN
13200	w42	NaN
14400	w46	NaN
15600	w50	NaN
16800	w53	NaN
18000	w57	NaN
19200	w61	NaN
20400	w65	NaN

e) Run the following code first and answer the question. Could you have done this by using Boolean masking?

```
[28]: mlt_df.query('artist == "Spears, Britney"')
```

```
[28]:
```

	artist	track	time	date_e	\
51	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	
63	Spears, Britney	From The Bottom Of My Broken Heart	4:30	2000-01-29	
93	Spears, Britney	Lucky	3:23	2000-08-12	
368	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	
380	Spears, Britney	From The Bottom Of My Broken Heart	4:30	2000-01-29	
...	
20034	Spears, Britney	From The Bottom Of My Broken Heart	4:30	2000-01-29	
20064	Spears, Britney	Lucky	3:23	2000-08-12	
20339	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	
20351	Spears, Britney	From The Bottom Of My Broken Heart	4:30	2000-01-29	
20381	Spears, Britney	Lucky	3:23	2000-08-12	

	W	Value
51	w1	67.0
63	w1	76.0
93	w1	61.0
368	w2	38.0
380	w2	59.0
...
20034	w64	NaN
20064	w64	NaN
20339	w65	NaN
20351	w65	NaN
20381	w65	NaN

[195 rows x 6 columns]

Could have done this by using Boolean masking too?

Answer:

```
[30]: BM = mlt_df.artist == "Spears, Britney"
mlt_df[BM]
```

```
[30]:
```

	artist	track	time	date_e	\
51	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	
63	Spears, Britney	From The Bottom Of My Broken Heart	4:30	2000-01-29	
93	Spears, Britney	Lucky	3:23	2000-08-12	
368	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	
380	Spears, Britney	From The Bottom Of My Broken Heart	4:30	2000-01-29	
...	
20034	Spears, Britney	From The Bottom Of My Broken Heart	4:30	2000-01-29	
20064	Spears, Britney	Lucky	3:23	2000-08-12	
20339	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	
20351	Spears, Britney	From The Bottom Of My Broken Heart	4:30	2000-01-29	
20381	Spears, Britney	Lucky	3:23	2000-08-12	

	W	Value
51	w1	67.0
63	w1	76.0
93	w1	61.0
368	w2	38.0
380	w2	59.0
...
20034	w64	NaN
20064	w64	NaN
20339	w65	NaN
20351	w65	NaN
20381	w65	NaN

[195 rows x 6 columns]

f) Use either the approach in e or the Boolean mask, to extract all the unique songs that Br

```
[33]: mlt_df[BM].track.unique()
```

```
[33]: array(['Oops!.. I Did It Again', 'From The Bottom Of My Broken Heart',
        'Lucky'], dtype=object)
```

g) In mlt_df show all of the weeks that the song "Oops!.. I Did It Again" was on the top-100

```
[44]: BM2 = mlt_df.track == "Oops!.. I Did It Again"
BM3 = mlt_df.Value <= 100
print(mlt_df[BM & BM2 & BM3])
mlt_df[BM & BM2 & BM3].W.unique()
```

	artist	track	time	date_e	W	Value
51	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w1	67.0
368	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w2	38.0
685	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w3	26.0

1002	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w4	19.0
1319	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w5	15.0
1636	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w6	13.0
1953	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w7	12.0
2270	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w8	9.0
2587	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w9	9.0
2904	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w10	10.0
3221	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w11	14.0
3538	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w12	15.0
3855	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w13	17.0
4172	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w14	24.0
4489	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w15	35.0
4806	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w16	49.0
5123	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w17	77.0
5440	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w18	85.0
5757	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w19	90.0
6074	Spears, Britney	Oops!.. I Did It Again	3:30	2000-04-22	w20	96.0

```
[44]: array(['w1', 'w2', 'w3', 'w4', 'w5', 'w6', 'w7', 'w8', 'w9', 'w10', 'w11',
          'w12', 'w13', 'w14', 'w15', 'w16', 'w17', 'w18', 'w19', 'w20'],
          dtype=object)
```

Excercise 4 We will use LaqnData.csv for this exercise. Each row of this dataset shows an hourly measurement recording of one of the five following air pollutants: NO, NO2, NOX, PM10, and PM2.5. The data was collected in a location in London for the entirety of year 2017. Read the data using Pandas and perform the following tasks.

```
[48]: import pandas as pd
import matplotlib.pyplot as plt

air_df = pd.read_csv('LaqnData.csv')
air_df
```

```
[48]:
```

	Site	Species	ReadingDateTime	Value	Units	Provisional or Ratified
0	CT3	NO	01/01/2017 00:00	3.5	ug m-3	R
1	CT3	NO	01/01/2017 01:00	3.6	ug m-3	R
2	CT3	NO	01/01/2017 02:00	2.2	ug m-3	R
3	CT3	NO	01/01/2017 03:00	2.1	ug m-3	R
4	CT3	NO	01/01/2017 04:00	3.3	ug m-3	R
...
43795	CT3	PM2.5	31/12/2017 19:00	-2.0	ug m-3	R
43796	CT3	PM2.5	31/12/2017 20:00	6.0	ug m-3	R
43797	CT3	PM2.5	31/12/2017 21:00	5.0	ug m-3	R
43798	CT3	PM2.5	31/12/2017 22:00	5.0	ug m-3	R
43799	CT3	PM2.5	31/12/2017 23:00	6.0	ug m-3	R

```
[43800 rows x 6 columns]
```

a) The dataset has 6 columns. Three of them, naming 'Site', 'Units', 'Provisional or Ratified' :

```
[49]: air_df.drop(columns=['Site', 'Units', 'Provisional or Ratified'], inplace=True)
air_df
```

```
[49]:
```

	Species	ReadingDateTime	Value
0	NO	01/01/2017 00:00	3.5
1	NO	01/01/2017 01:00	3.6
2	NO	01/01/2017 02:00	2.2
3	NO	01/01/2017 03:00	2.1
4	NO	01/01/2017 04:00	3.3
...
43795	PM2.5	31/12/2017 19:00	-2.0
43796	PM2.5	31/12/2017 20:00	6.0
43797	PM2.5	31/12/2017 21:00	5.0
43798	PM2.5	31/12/2017 22:00	5.0
43799	PM2.5	31/12/2017 23:00	6.0

[43800 rows x 3 columns]

b) The dataset is in a long format. Apply the appropriate function to switch it to the wide format :

```
[51]: pvt_df = air_df.pivot(
        index="ReadingDateTime",
        columns="Species",
        values="Value"
    )
pvt_df
```

```
[51]:
```

Species	NO	NO2	NOX	PM10	PM2.5
ReadingDateTime					
01/01/2017 00:00	3.5	30.8	36.2	35.7	NaN
01/01/2017 01:00	3.6	31.5	37.0	28.5	NaN
01/01/2017 02:00	2.2	27.3	30.7	22.7	NaN
01/01/2017 03:00	2.1	23.5	26.8	20.5	NaN
01/01/2017 04:00	3.3	28.0	33.0	22.1	NaN
...
31/12/2017 19:00	0.7	17.5	18.5	16.3	-2.0
31/12/2017 20:00	0.7	17.5	18.6	14.5	6.0
31/12/2017 21:00	0.7	14.1	15.1	8.6	5.0
31/12/2017 22:00	1.1	22.0	23.6	12.5	5.0
31/12/2017 23:00	0.9	19.4	20.7	10.4	6.0


[8760 rows x 5 columns]


c) Draw and study the histogram and boxplots for columns of pvt_df.


```
[55]: import matplotlib.pyplot as plt

for col in pvt_df:
    print(f"\n--- Column: {col} ---")
    print(pvt_df[col].describe())
    pvt_df[col].plot.hist(title=col)
    plt.show()
    pvt_df[col].plot.box(title=col)
    plt.show()
```

```
--- Column: NO ---
count      8660.000000
mean        15.045115
std         26.678565
min         -2.000000
25%          3.100000
50%          7.100000
75%         15.500000
max         401.799990
Name: NO, dtype: float64
```







```


--- Column: NO2 ---
count      8660.000000
mean        38.010185
std         18.580841
min          2.800000
25%         23.800000
50%         35.300000
75%         49.900000
max        120.200000
Name: NO2, dtype: float64


```





```
--- Column: NOX ---  
count      8660.000000  
mean        61.078661  
std         54.584805  
min          1.000000  
25%         30.300000  
50%         46.900000  
75%         73.800000  
max        734.299990  
Name: NOX, dtype: float64
```







```


--- Column: PM10 ---
count      8657.000000
mean        22.551704
std         15.344755
min         -5.600000
25%         13.900000
50%         19.500000
75%         27.300000
max         633.099980
Name: PM10, dtype: float64


```





```
--- Column: PM2.5 ---
count    5933.000000
mean      14.999831
std       13.558588
min       -8.000000
25%        7.000000
50%       11.000000
75%       18.000000
max       128.000000
Name: PM2.5, dtype: float64
```



Exercise 5 We will continue working with LaqnData.csv.

a) Run the following code, see its output, and then study the code to answer what each line of

```
[56]: import pandas as pd
air_df = pd.read_csv('LaqnData.csv')
air_df.drop(columns=['Site', 'Units', 'Provisional or Ratified'], inplace=True)
datetime_df = air_df.ReadingDateTime.str.split(' ', expand=True)
datetime_df.columns = ['Date', 'Time']
date_df = datetime_df.Date.str.split('/', expand=True)
date_df.columns = ['Day', 'Month', 'Year']
air_df = air_df.join(date_df).join(datetime_df.Time).
    ↳drop(columns=['ReadingDateTime', 'Year'])
air_df
```

```
[56]:
```

	Species	Value	Day	Month	Time
0	NO	3.5	01	01	00:00
1	NO	3.6	01	01	01:00
2	NO	2.2	01	01	02:00
3	NO	2.1	01	01	03:00
4	NO	3.3	01	01	04:00

```

...      ...      ...      ..      ...      ...
43795    PM2.5    -2.0    31      12    19:00
43796    PM2.5     6.0    31      12    20:00
43797    PM2.5     5.0    31      12    21:00
43798    PM2.5     5.0    31      12    22:00
43799    PM2.5     6.0    31      12    23:00

```

[43800 rows x 5 columns]

Answer:

- Imports the Pandas library and assigns it the alias `pd` for easier use.
- Reads the CSV file `LaqnData.csv` into a Pandas DataFrame called `air_df`.
- Drops unnecessary columns ('Site', 'Units', 'Provisional or Ratified') from `air_df` permanently (`inplace=True`).
- Takes the `ReadingDateTime` column (e.g., "01/01/2020 00:00") and splits it into two columns using space (' ') as the delimiter, `expand=True` ensures the split parts are returned as separate columns in a new DataFrame (`datetime_df`).
- Renames the columns of `datetime_df` to 'Date' and 'Time'.
- Takes the `Date` column (e.g., "01/01/2020") and splits it into three columns using / as the delimiter. The result is stored in a new DataFrame (`date_df`).
- Renames the columns of `date_df` to 'Day', 'Month', and 'Year'.
- Joins the new columns (Day, Month, Time) back to the original `air_df`. Drops the original `ReadingDateTime` and `Year` columns (since they are redundant after splitting).

b) Run the following code, see its output, and then study the code to answer what does this 1.

```
[57]: air_df = air_df.set_index(['Month', 'Day', 'Time', 'Species'])
air_df
```

```
[57]:
```

				Value
Month	Day	Time	Species	
01	01	00:00	NO	3.5
		01:00	NO	3.6
		02:00	NO	2.2
		03:00	NO	2.1
		04:00	NO	3.3
...				...
12	31	19:00	PM2.5	-2.0
		20:00	PM2.5	6.0
		21:00	PM2.5	5.0
		22:00	PM2.5	5.0
		23:00	PM2.5	6.0

[43800 rows x 1 columns]

Answer:

Converts `air_df` into a `MultiIndex DataFrame` by setting the columns ['Month', 'Day', 'Time', 'Species'] as the new index levels.

c) Run the following code, see its output, and then study the code to answer what does this 1.

```
[58]: air_df.unstack()
```

```
[58]:
```

			Value				
Species			NO	NO2	NOX	PM10	PM2.5
Month	Day	Time					
01	01	00:00	3.5	30.8	36.2	35.7	NaN
		01:00	3.6	31.5	37.0	28.5	NaN
		02:00	2.2	27.3	30.7	22.7	NaN
		03:00	2.1	23.5	26.8	20.5	NaN
		04:00	3.3	28.0	33.0	22.1	NaN
...		
12	31	19:00	0.7	17.5	18.5	16.3	-2.0
		20:00	0.7	17.5	18.6	14.5	6.0
		21:00	0.7	14.1	15.1	8.6	5.0
		22:00	1.1	22.0	23.6	12.5	5.0
		23:00	0.9	19.4	20.7	10.4	6.0

[8760 rows x 5 columns]

Answer:

Reshapes a MultiIndex DataFrame by “pivoting” one or more levels of the row index into columns, creating a wide-format table.

- d) Compare the output of the code above with pvt_df from Exercise 4. Are they basically the same?
- e) Explain what are the differences and similarities between the pair .melt()/.pivot() and the pair .stack()/.unstack().
- f) If you were to choose one counterpart for .melt() between .stack()/.unstack() which one would you choose?

Answer:

- d) Yes
- e) .melt()/.pivot() is to reshape between wide and long formats. .stack()/.unstack() is to reshape between hierarchical indexes and columns. .melt()/.pivot() input dataframes with regular columns. .stack()/.unstack() input dataframes with multiindex rows/columns.
- f) Choose .stack() as the counterpart to .melt().