

# Excercises Ch2

April 20, 2025

## 1 Hands-On Data Preprocessing in Python

Learn how to effectively prepare data for successful data analytics

AUTHOR: Dr. Roy Jafari

### 1.0.1 Chapter 1: Review of the core modules NumPy, Pandas, and Matplotlib

```
[2]: import pandas as pd
```

**Excercise 1** Use adult.csv and Boolean Masking to answer the following questions.

```
[3]: adult_df = pd.read_csv('adult.csv')
adult_df.head()
```

```
[3]:
```

	age	workclass	fnlwgt	education	education-num \	
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	race	sex \	
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

	capitalGain	capitalLoss	hoursPerWeek	nativeCountry	income	
0	2174	0	40	United-States	<=50K	
1	0	0	13	United-States	<=50K	
2	0	0	40	United-States	<=50K	
3	0	0	40	United-States	<=50K	
4	0	0	40	Cuba	<=50K	

- Calculate the mean and median of education-num for every race in the data.
- Draw one histogram of education-num that includes the data for each race in the data.
- Draw a comparative boxplot that compares the education-num for each race.

d) Create a subplot that puts the visual from b on top of the one on c.

```
[11]: # a) Calculate the mean and median of education-num for every race in the data.
      races = adult_df['race'].unique()

      for race in races:
          mask = adult_df['race'] == race
          mean = adult_df.loc[mask, 'education-num'].mean()
          median = adult_df.loc[mask, 'education-num'].median()
          print(f"Race: {race}\n Mean: {mean: 1f}, Median: {median: .0f}\n")
```

Race: White

Mean: 10.135246, Median: 10

Race: Black

Mean: 9.486236, Median: 9

Race: Asian-Pac-Islander

Mean: 10.960539, Median: 10

Race: Amer-Indian-Eskimo

Mean: 9.311897, Median: 9

Race: Other


Mean: 8.841328, Median: 9

```
[26]: # b) Draw one histogram of education-num that includes the data for each race
      ↪ in the data.
```

```
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
for race in adult_df['race'].unique():
    plt.hist(
        adult_df[adult_df['race'] == race]['education-num'],
        label=race
    )

plt.title('Education-Num Distribution by Race')
plt.xlabel('Education-Num (Years)')
plt.ylabel('Frequency')
plt.legend()
plt.show()
```



[39]: # c) Draw a comparative boxplot that compares the education-num for each race.


```
import matplotlib.pyplot as plt

# Prepare data: List of arrays (one per race)
data = [
    adult_df[adult_df['race'] == race]['education-num'].values
    for race in adult_df['race'].unique()
]
labels = adult_df['race'].unique()

# Create boxplot
plt.figure(figsize=(12, 10))
plt.boxplot(data, tick_labels=labels)

# Customize
plt.title('Education-Num Distribution by Race')
plt.xlabel('Race')
plt.ylabel('Education-Num (Years)')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```



[48]: # d) Create a subplot that puts the visual from b on top of the one on c.

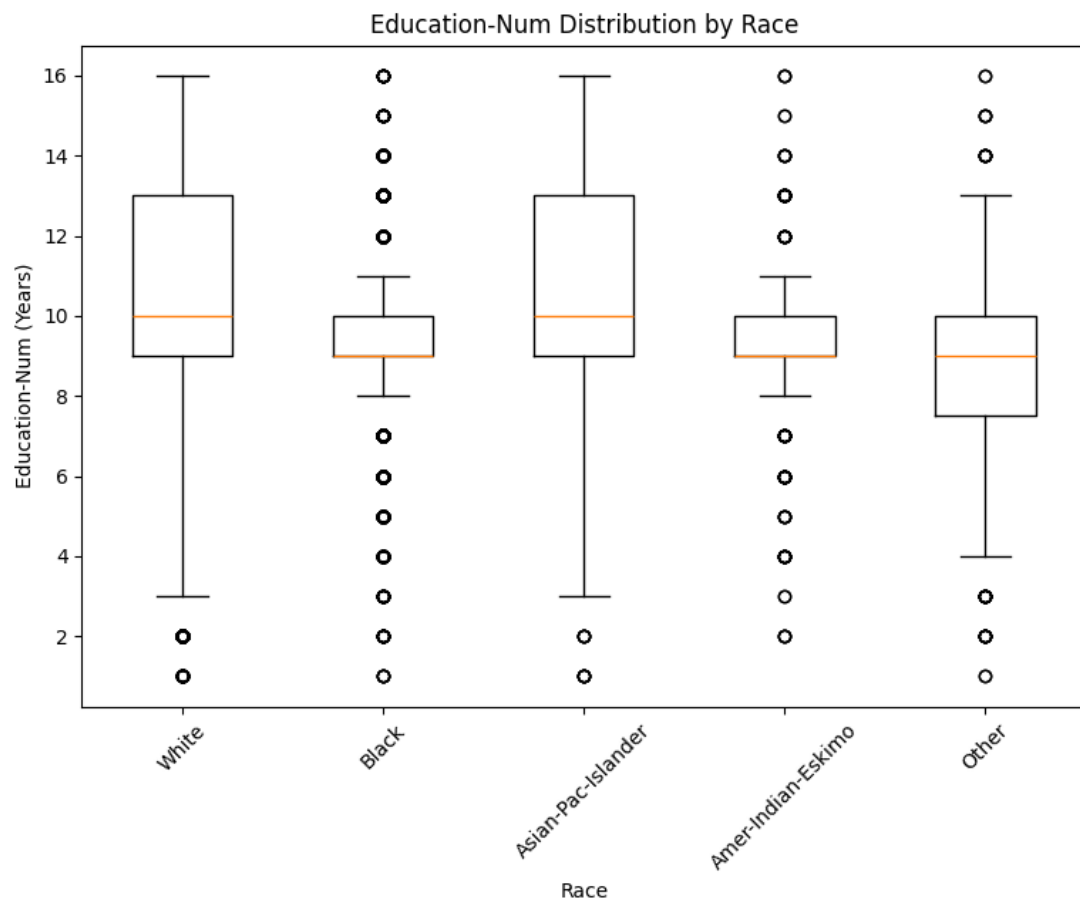
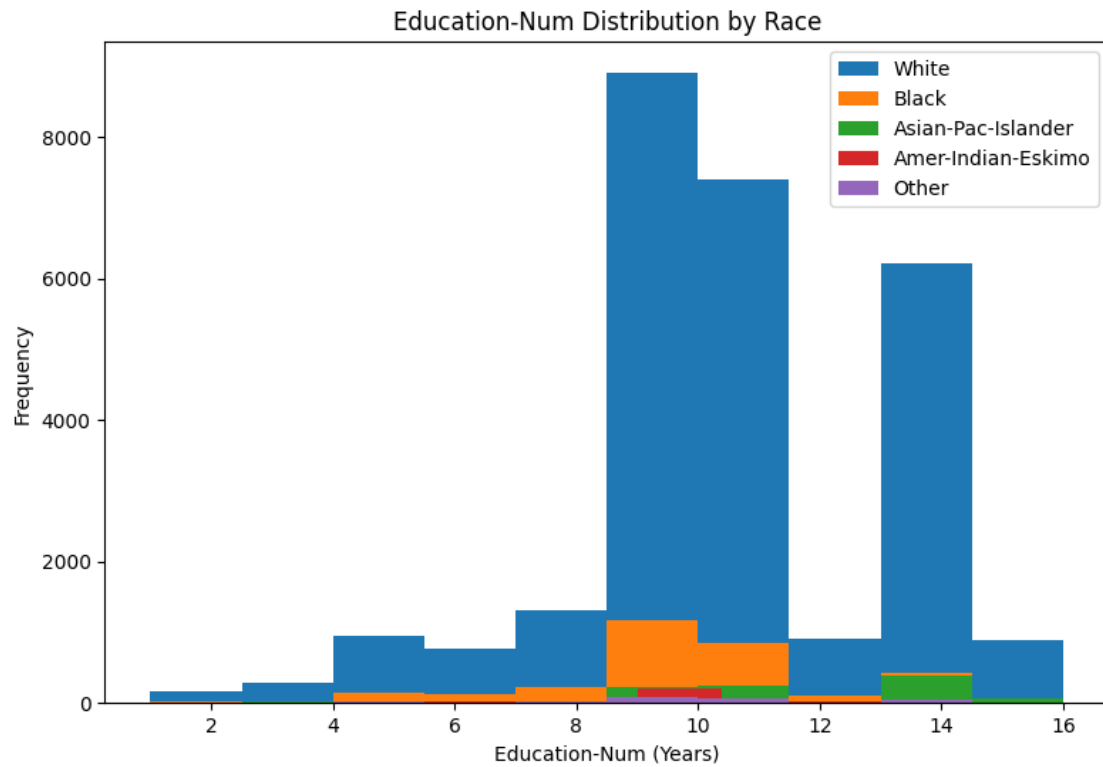
```
plt.figure(figsize=(8, 12))

plt.subplot(2, 1, 1)
for race in adult_df['race'].unique():
    plt.hist(
        adult_df[adult_df['race'] == race]['education-num'],
        label=race
    )
plt.title('Education-Num Distribution by Race')
plt.xlabel('Education-Num (Years)')
plt.ylabel('Frequency')
plt.legend()

plt.subplot(2, 1, 2)
data = [
    adult_df[adult_df['race'] == race]['education-num'].values
```

```
    for race in adult_df['race'].unique()
]
labels = adult_df['race'].unique()
plt.boxplot(data, tick_labels=labels)
plt.title('Education-Num Distribution by Race')
plt.xlabel('Race')
plt.ylabel('Education-Num (Years)')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```



## Excercise 2

- a) Repeat the analysis on Exercise 1. a), but this time use groupby function.
- b) compare the runtime of using BM vs. groupby. (hint: you can import the module time and use

```
[54]: # a) Repeat the analysis on Exercise 1. a), but this time use groupby function.
```

```
print(adult_df.groupby('race')['education-num'].agg(['mean', 'median']))
```

	mean	median
race		
Amer-Indian-Eskimo	9.311897	9.0
Asian-Pac-Islander	10.960539	10.0
Black	9.486236	9.0
Other	8.841328	9.0
White	10.135246	10.0

```
[56]: # b) compare the runtime of using BM vs. groupby. (hint: you can import the
      ↪ module time and use the fuction .time())
```

```
import time
```

```
# Boolean Masking Approach
```

```
def boolean_masking_stats():
```

```
    races = adult_df['race'].unique()
```

```
    for race in races:
```

```
        mask = adult_df['race'] == race
```

```
        mean = adult_df.loc[mask, 'education-num'].mean()
```

```
        median = adult_df.loc[mask, 'education-num'].median()
```

```
# GroupBy Approach
```

```
def groupby_stats():
```

```
    grouped_stats = adult_df.groupby('race')['education-num'].agg(['mean',
```

```
    ↪ 'median'])
```

```
# Measure runtime
```

```
n_iterations = 100
```

```
# Time Boolean Masking
```

```
bm_time = time.time()
```

```
for _ in range(n_iterations):
```

```
    boolean_masking_stats()
```

```
bm_time = (time.time() - bm_time) / n_iterations
```

```
# Time GroupBy
```

```

gb_time = time.time()
for _ in range(n_iterations):
    groupby_stats()
gb_time = (time.time() - gb_time) / n_iterations

print(f"Boolean Masking: {bm_time:.6f} sec per iteration")
print(f"GroupBy: {gb_time:.6f} sec per iteration")
print(f"GroupBy is {bm_time/gb_time:.1f}x faster")

```

Boolean Masking: 0.010407 sec per iteration  
 GroupBy: 0.002198 sec per iteration  
 GroupBy is 4.7x faster

### Exercise 3

If you have not already, solve exercise 4 in the previous chapter. After you created pvt\_df for

```

[57]: import seaborn as sns

air_df = pd.read_csv('LaqnData.csv')
air_df.drop(columns=['Site', 'Units', 'Provisional or Ratified'], inplace=True)


pvt_df = air_df.pivot(
    index="ReadingDateTime",
    columns="Species",
    values="Value"
)

sns.pairplot(pvt_df)

```

[57]: <seaborn.axisgrid.PairGrid at 0x7f0f25a75990>





The code outputs what is known as scatter matrix. This code takes advantage of seaborn module wh

```
[59]: # Get species names and count
species = pvt_df.columns
n_species = len(species)

# Create figure with proper sizing
fig = plt.figure(figsize=(12, 12))
plt.suptitle('Air Quality Species Relationships', y=1.02, fontsize=14)

# Create gridspec for more control
gs = plt.GridSpec(n_species, n_species, wspace=0.1, hspace=0.1)
```

```

# Plot each combination
for i in range(n_species):
    for j in range(n_species):
        ax = plt.subplot(gs[i, j])

        if i == j:
            # Diagonal - show histogram
            ax.hist(pvt_df.iloc[:, i], bins=20, color='skyblue')
            ax.set_xticks([])
            ax.set_yticks([])
        else:
            # Off-diagonal - show scatter plot
            ax.scatter(pvt_df.iloc[:, j], pvt_df.iloc[:, i],
                       s=5, alpha=0.5, color='navy')
            ax.set_xticks([])
            ax.set_yticks([])

        # Add axis labels only on edges
        if i == n_species-1:
            ax.set_xlabel(species[j], fontsize=8)
        if j == 0:
            ax.set_ylabel(species[i], fontsize=8)

plt.tight_layout()
plt.show()

```

/tmp/ipykernel\_253923/3982319793.py:35: UserWarning: This figure includes Axes that are not compatible with tight\_layout, so results might be incorrect.

```
plt.tight_layout()
```

## Air Quality Species Relationships

