

Excercises Ch3

April 21, 2025

0.1 Hands-On Data Preprocessing in Python

Learn how to effectively prepare data for successful data analytics

AUTHOR: Dr. Roy Jafari

0.1.1 Chapter 3: Data

Excercise 1

- 1) From 5 colleagues or classmates ask to provide a definition for the term data.
 - a) Report these definitions and indicate the similarity among them.
 - b) In your own words, define the all-encompassing definition of data put forth in this chapter.
 - c) Indicate the two important aspects of the definition in b.
 - d) Compare the 5 definitions of data from your colleagues with the all-encompassing definitions and indicate
- a) Colleague Definitions & Similarities Colleague 1: “Numbers and facts collected for analysis” Colleague 2: “Information stored in computers” Colleague 3: “Measurements used to make decisions” Colleague 4: “Observations about the world” Colleague 5: “Raw material for creating knowledge”

Similarities:

All associate data with information collection

3/5 mention analysis or decision-making

None explicitly mention representation or structure

- b) All-Encompassing Definition Data: Structured or unstructured representations of objects, events, or their attributes, captured in a reproducible format to enable interpretation and analysis.
- c) Two Key Aspects Representation: Data encodes reality in symbolic forms (numbers, text, images)

Purpose: Collected to support reasoning, analysis, or decision-making

- d) Comparison vs. Colleague Definitions Colleague Definition | Matches All-Encompassing? | Missing Elements “Numbers and facts” | Partial | Representation format, purpose “Computer-stored info” | Partial | Real-world connection, structure “Measurements” | Closest | Represen-

tation diversity “Observations” | Partial | Storage/structured aspect “Raw material” | Partial
| Representation method

Excercise 2 For this exercise, we are going to use covid_impact_on_airport_traffic.csv. Answer the following questions. This dataset is from Kaggle.com, use this link to see its page: <https://www.kaggle.com/terenceshin/covid19s-impact-on-airport-traffic>. The key attribute of this dataset is PercentOfBaseline which shows the ratio of air traffic in the specific day compared to pre-pandemic time (1st Feb to 15th March 2020)

```
[12]: import pandas as pd
covid_df = pd.read_excel("covid_impact_on_airport_traffic.xlsx")
covid_df
```

```
[12]:
```

	AggregationMethod	Date	Version	AirportName	\
0	Daily	2020-04-03	1	Kingsford Smith	
1	Daily	2020-04-13	1	Kingsford Smith	
2	Daily	2020-07-10	1	Kingsford Smith	
3	Daily	2020-09-02	1	Kingsford Smith	
4	Daily	2020-10-31	1	Kingsford Smith	
...	
7242	Daily	2020-06-05	1	Seattle-Tacoma International	
7243	Daily	2020-10-03	1	Seattle-Tacoma International	
7244	Daily	2020-07-16	1	Seattle-Tacoma International	
7245	Daily	2020-07-31	1	Seattle-Tacoma International	
7246	Daily	2020-08-30	1	Seattle-Tacoma International	

	PercentOfBaseline	Centroid	City	\
0	64	POINT(151.180087713813 -33.9459774986125)	Sydney	
1	29	POINT(151.180087713813 -33.9459774986125)	Sydney	
2	54	POINT(151.180087713813 -33.9459774986125)	Sydney	
3	18	POINT(151.180087713813 -33.9459774986125)	Sydney	
4	22	POINT(151.180087713813 -33.9459774986125)	Sydney	
...	
7242	80	POINT(-122.308661576118 47.4505828917119)	SeaTac	
7243	55	POINT(-122.308661576118 47.4505828917119)	SeaTac	
7244	76	POINT(-122.308661576118 47.4505828917119)	SeaTac	
7245	69	POINT(-122.308661576118 47.4505828917119)	SeaTac	
7246	68	POINT(-122.308661576118 47.4505828917119)	SeaTac	

	State	ISO_3166_2	Country	\
0	New South Wales	AU	Australia	
1	New South Wales	AU	Australia	
2	New South Wales	AU	Australia	
3	New South Wales	AU	Australia	
4	New South Wales	AU	Australia	
...	
7242	Washington	US-WA	United States of America (the)	

7243	Washington	US-WA	United States of America (the)
7244	Washington	US-WA	United States of America (the)
7245	Washington	US-WA	United States of America (the)
7246	Washington	US-WA	United States of America (the)

```

                                Geography
0    POLYGON((151.164354085922 -33.9301772341877, 1...
1    POLYGON((151.164354085922 -33.9301772341877, 1...
2    POLYGON((151.164354085922 -33.9301772341877, 1...
3    POLYGON((151.164354085922 -33.9301772341877, 1...
4    POLYGON((151.164354085922 -33.9301772341877, 1...
...
7242 POLYGON((-122.297594547272 47.434474106872, -1...
7243 POLYGON((-122.297594547272 47.434474106872, -1...
7244 POLYGON((-122.297594547272 47.434474106872, -1...
7245 POLYGON((-122.297594547272 47.434474106872, -1...
7246 POLYGON((-122.297594547272 47.434474106872, -1...

```

[7247 rows x 11 columns]

- What is the best definition of the data object for this dataset?
- Are there any attributes in the data that only have one value? Use `.unique()` function to check.
- What type of values do the remaining attributes carry?
- How much statistical information the attribute 'PercentOfBaseline' has?

Answer:

- time-series panel data object

```

[13]: # b)

single_value_cols = [
    col for col in covid_df.columns
    if len(covid_df[col].unique()) == 1
]
print("Single-value columns:", single_value_cols)

```

Single-value columns: ['AggregationMethod', 'Version']

```

[14]: covid_df = covid_df.drop(columns=single_value_cols)

```

```

[15]: # c)
print(covid_df.dtypes)

# Key types:
# Date                interval
# AirportName         nominal
# PercentOfBaseline   ratio
# Centroid            nominal

```

```
# City          nominal
# State         nominal
# ISO_3166_2    nominal
# Country       nominal
# Geography     nominal
```

```
Date          datetime64[ns]
AirportName    object
PercentOfBaseline  int64
Centroid       object
City           object
State          object
ISO_3166_2     object
Country        object
Geography      object
dtype: object
```

```
[16]: # d)
print(covid_df['PercentOfBaseline'].describe())
```

```
count    7247.000000
mean      66.651442
std       22.134433
min        0.000000
25%       53.000000
50%       67.000000
75%       84.000000
max      100.000000
Name: PercentOfBaseline, dtype: float64
```

Exercise 3 For this exercise, we are going to use US_Accidents.csv. Answer the following questions. This dataset is from Kaggle.com, use this link to see its page: <https://www.kaggle.com/sobhanmoosavi/us-accidents>. This dataset shows all the car accidents in the US from February 2016 to Dec 2020.

```
[19]: # import pandas as pd
# pd.read_csv('US_Accidents.csv')
```

- What is the best definition of the data object for this dataset?
- Are there any attributes in the data that only have one value? Use `.unique()` function to check.
- What type of values do the remaining attributes carry?
- How much statistical information the numerical attributes of the dataset carry?
- Compare the statistical information of the numerical attributes and see if any of them are

Answer:

```
[20]: # skip because dataset too large
```

Excercise 4 For this exercise, we are going to use fatal-police-shootings-data.csv. There are a lot of debates, discussions, dialogues, and protests happening in the US surrounding police killings. The Washington Post has been collecting data on all fatal police shootings in the US. The dataset available to the government and the public alike has date, age, gender, race, location, and other situational information of these fatal police shootings. You can read more about this data on <https://www.washingtonpost.com/graphics/investigations/police-shootings-database/>, and you can download the last version of the data from <https://github.com/washingtonpost/data-police-shootings>

```
[21]: import pandas as pd

fatal_df = pd.read_csv('fatal-police-shootings-data.csv')
fatal_df.head()
```

```
[21]:    id      name      date  manner_of_death  armed  age \
0    3    Tim Elliot  2015-01-02          shot      gun  53.0
1    4  Lewis Lee Lembke  2015-01-02          shot      gun  47.0
2    5  John Paul Quintero  2015-01-03  shot and Tasered  unarmed  23.0
3    8   Matthew Hoffman  2015-01-04          shot  toy weapon  32.0
4    9  Michael Rodriguez  2015-01-04          shot   nail gun  39.0
```

```
    gender race      city state  signs_of_mental_illness  threat_level \
0      M    A    Shelton    WA              True      attack
1      M    W    Aloha    OR              False      attack
2      M    H    Wichita    KS              False      other
3      M    W  San Francisco    CA              True      attack
4      M    H    Evans    CO              False      attack
```

```
    flee  body_camera  longitude  latitude  is_geocoding_exact
0  Not fleeing      False   -123.122    47.247              True
1  Not fleeing      False   -122.892    45.487              True
2  Not fleeing      False   -97.281    37.695              True
3  Not fleeing      False  -122.422    37.763              True
4  Not fleeing      False  -104.692    40.384              True
```

- What is the best definition of the data object for this dataset?
 - Are there any attributes in the data that only have one value? Use `.unique()` function to check.
 - What type of values do the remaining attributes carry?
 - How much statistical information the numerical attributes of the dataset carry?
 - Compare the statistical information of the numerical attributes and see if any of them are
- a) This dataset is a structured event-log data object

```
[22]: # b)

single_value_cols = [
    col for col in fatal_df.columns
    if fatal_df[col].nunique() == 1
]
```

```
single_value_cols
```

```
[22]: []
```

```
[23]: # c)
```

```
fatal_df.dtypes
```

```
# id                nominal
# name              nominal
# date              interval
# manner_of_death   nominal
# armed             nominal
# age               ratio
# gender            nominal
# race              nominal
# city              nominal
# state             nominal
# signs_of_mental_illness nominal
# threat_level      ordinal
# flee              ordinal
# body_camera        nominal
# longitude          interval
# latitude           interval
# is_geocoding_exact nominal
```

```
[23]: id                int64
name                object
date                object
manner_of_death     object
armed               object
age                 float64
gender              object
race                object
city                object
state               object
signs_of_mental_illness bool
threat_level        object
flee                object
body_camera         bool
longitude            float64
latitude             float64
is_geocoding_exact  bool
dtype: object
```

```
[24]: # d)
fatal_df.describe()
```

```
[24]:
```

	id	age	longitude	latitude
count	6068.000000	5802.000000	5771.000000	5771.000000
mean	3363.828444	37.160290	-97.238779	36.660617
std	1898.394615	13.039693	16.629397	5.377469
min	3.000000	6.000000	-158.137000	19.498000
25%	1716.750000	27.000000	-112.117000	33.475500
50%	3369.500000	35.000000	-94.386000	36.096000
75%	5019.250000	46.000000	-83.067000	39.986500
max	6629.000000	91.000000	-68.014000	71.301000

```
[25]: # e)

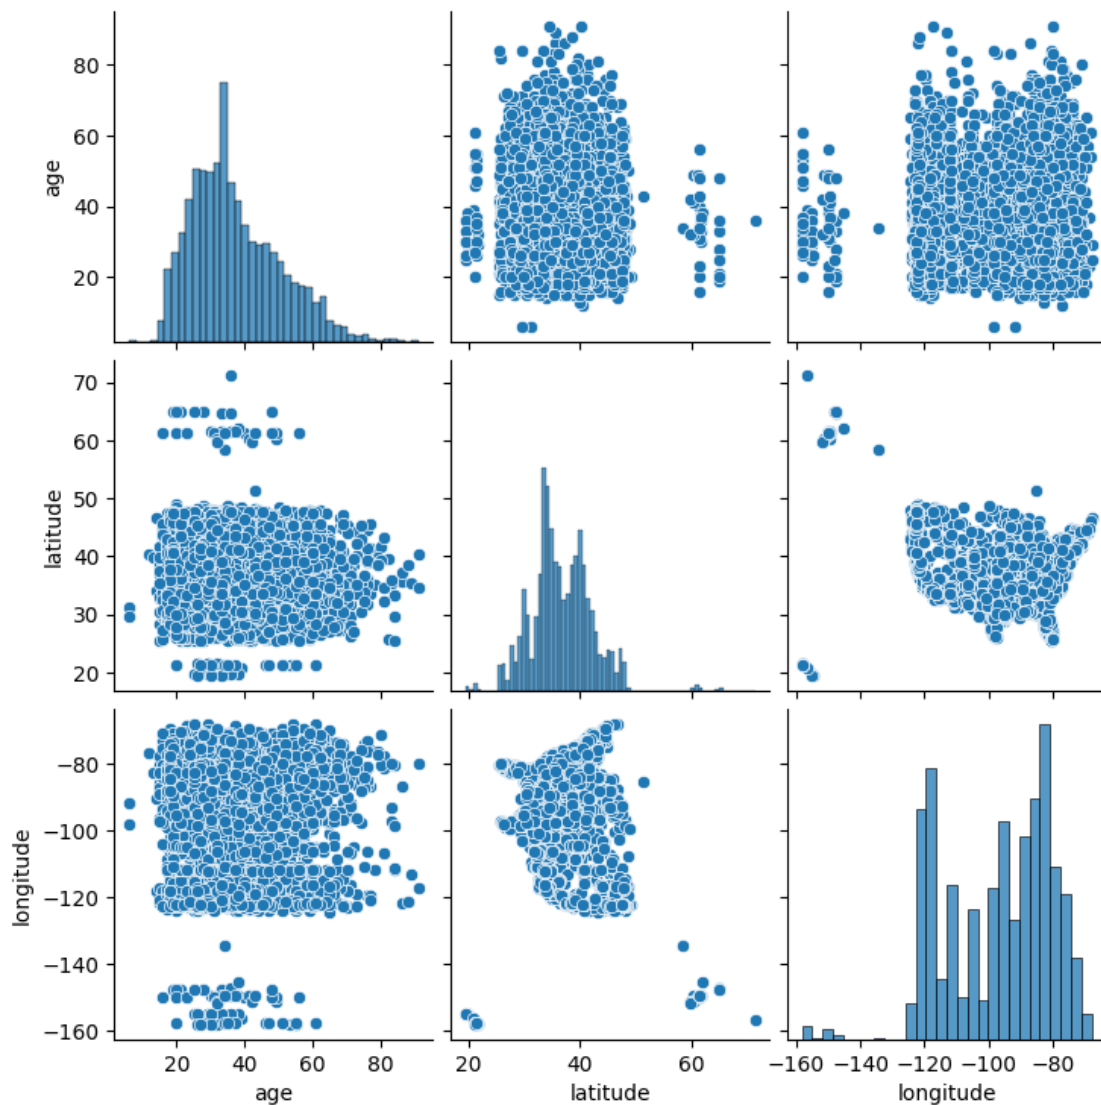
fatal_df[['age', 'latitude', 'longitude']].corr()
```

```
[25]:
```

	age	latitude	longitude
age	1.000000	-0.014671	0.057657
latitude	-0.014671	1.000000	-0.116595
longitude	0.057657	-0.116595	1.000000

```
[26]: import matplotlib.pyplot as plt
import seaborn as sns

sns.pairplot(fatal_df[['age', 'latitude', 'longitude']])
plt.show()
```



Exercise 5 For this exercise, we will be using `electricity_prediction.csv`. The screenshot below shows the 5 rows of this dataset and a linear regression model created to predict electricity consumption based on the weekday and daily average temperature.

```
[27]: import pandas as pd

electricity_df = pd.read_csv('electricity_prediction.csv')
electricity_df.head()
```

```
[27]:
```

	Date	Weekday	Consumption	Average Temperature
0	1/1/2016	4	2581914	80
1	1/2/2016	5	2663011	77
2	1/3/2016	6	2725351	78

3	1/4/2016	0	3092978	80
4	1/5/2016	1	3231827	81

```
[28]: from sklearn.linear_model import LinearRegression

X = electricity_df[['Weekday', 'Average Temperature']]
y = electricity_df['Consumption']

lrm = LinearRegression()
lrm.fit(X, y)

print('intercept ', lrm.intercept_)
print(pd.DataFrame({'Predictor': X.columns, 'coefficient': lrm.coef_}))
```

```
intercept    3074181.495015881
            Predictor coefficient
0              Weekday -55710.145405
1  Average Temperature -3476.377056
```

The regression model that is derived from the data is presented below.

Consumption = 3074181.5- 55710.1 × Weekday-3476.4 ×Average Temperature

What is the fundamental mistake in this analysis? Describe it and provides possible solutions for it.

Answer:

The primary issue is treating the Weekday variable (0-6) as a continuous/numeric predictor in the linear regression mode. Weekdays are categorical (nominal) variables, not continuous. The model assumes a linear relationship between Weekday and Consumption, implying that the difference between Monday (0) and Tuesday (1) is the same as between Saturday (5) and Sunday (6). This is illogical because weekdays have no inherent numerical meaning (e.g., Sunday isn't "greater than" Monday).

Solution is to treat Weekday as Categorical (One-Hot Encoding)

Excercise 6 For this exercise, we will be using adult.csv. we used this dataset extensively in chapter 1. Read the dataset using Padans and call it adult_df.

```
[29]: import pandas as pd

adult_df = pd.read_csv('adult.csv')
```

- What type of values does the attribute education carry?
- Run 'adult_df.education.unique()', study the results, and explain what the code does.
- Based on your understandings, order the output of the code you ran for b).
- Run 'pd.get_dummies(adult_df.education)', study the results, and explain what the code does.
- Run 'adult_df.sort_values(['education-num']).iloc[1:32561:1200]', study the results and explain.
- Compare your answer to c) and what you learned from e). Was the order you came up with in c) correct?
- Education is an ordinal attribute, translating an ordinal attribute from an analytic perspective.

- 'adult_df.education'
 - 'pd.get_dummies(adult_df.education)'
 - 'adult_df['education']'
- h) Either of the choices has some advantages and some disadvantages. Select which programming value representation is better for an ordinal attribute.
- If an ordinal attribute is presented using this programming value representation, no bias is introduced.
 - If an ordinal attribute is presented using this programming value representation, the data is more readable.
 - If an ordinal attribute is presented using this programming value representation, there is no loss of information.

Answer:

a) Ordinal

[30]: # b)

```
adult_df.education.unique()
```

```
[30]: array(['Bachelors', 'HS-grad', '11th', 'Masters', '9th', 'Some-college',
          'Assoc-acdm', 'Assoc-voc', '7th-8th', 'Doctorate', 'Prof-school',
          '5th-6th', '10th', '1st-4th', 'Preschool', '12th'], dtype=object)
```

Returns an array of unique values in the education column.

c)

1. Preschool
2. 1st-4th
3. 5th-6th
4. 7th-8th
5. 9th
6. 10th
7. 11th
8. 12th
9. HS-grad
10. Some-college
11. Assoc-acdm
12. Assoc-voc
13. Bachelors
14. Masters
15. Prof-school
16. Doctorate

```
[31]: # d)
```

```
pd.get_dummies(adult_df.education)
```

```
[31]:
```

	10th	11th	12th	1st-4th	5th-6th	7th-8th	9th	Assoc-acdm	\
0	False	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	False	
3	False	True	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	False	
...	
32556	False	False	False	False	False	False	False	True	
32557	False	False	False	False	False	False	False	False	
32558	False	False	False	False	False	False	False	False	
32559	False	False	False	False	False	False	False	False	
32560	False	False	False	False	False	False	False	False	

	Assoc-voc	Bachelors	Doctorate	HS-grad	Masters	Preschool	\
0	False	True	False	False	False	False	
1	False	True	False	False	False	False	
2	False	False	False	True	False	False	
3	False	False	False	False	False	False	
4	False	True	False	False	False	False	
...	
32556	False	False	False	False	False	False	
32557	False	False	False	True	False	False	
32558	False	False	False	True	False	False	
32559	False	False	False	True	False	False	
32560	False	False	False	True	False	False	

	Prof-school	Some-college
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
...
32556	False	False
32557	False	False
32558	False	False
32559	False	False
32560	False	False

```
[32561 rows x 16 columns]
```

Performs one-hot encoding, converting each education level to a binary column

[32]: # e)

```
adult_df.sort_values(['education-num']).iloc[1:32561:1200]
```

```
[32]:
```

	age	workclass	fnlwgt	education	education-num \
32432	36	Private	208068	Preschool	1
32467	55	Private	199713	9th	5
2449	25	Private	345121	10th	6
14720	17	Private	99462	11th	7
4862	35	Private	46385	HS-grad	9
29380	25	Private	410240	HS-grad	9
27398	50	State-gov	97778	HS-grad	9
25406	26	Private	122206	HS-grad	9
23272	25	Private	109419	HS-grad	9
21374	44	Federal-gov	320071	HS-grad	9
19242	56	Private	134756	HS-grad	9
17302	60	Private	132529	HS-grad	9
15265	58	Private	126991	HS-grad	9
20514	25	State-gov	108542	Some-college	10
22642	18	Private	266489	Some-college	10
25709	24	Private	186213	Some-college	10
29946	38	Private	64879	Some-college	10
31715	22	State-gov	124942	Some-college	10
29794	47	Private	358465	Some-college	10
11908	26	Private	78424	Assoc-voc	11
1317	35	NaN	327120	Assoc-acdm	12
25132	41	Local-gov	177599	Bachelors	13
14210	54	Private	256908	Bachelors	13
7184	42	Private	168103	Bachelors	13
16832	49	Private	224393	Bachelors	13
17112	46	Private	198774	Masters	14
26083	71	Self-emp-inc	38822	Masters	14
7690	34	Private	228873	Doctorate	16

	marital-status	occupation	relationship \
32432	Divorced	Other-service	Not-in-family
32467	Married-civ-spouse	Craft-repair	Husband
2449	Separated	Other-service	Own-child
14720	Never-married	Other-service	Own-child
4862	Married-civ-spouse	Transport-moving	Husband
29380	Never-married	Craft-repair	Own-child
27398	Married-civ-spouse	Protective-serv	Husband
25406	Never-married	Craft-repair	Other-relative
23272	Never-married	Adm-clerical	Own-child
21374	Married-civ-spouse	Craft-repair	Husband
19242	Never-married	Other-service	Not-in-family
17302	Married-civ-spouse	Other-service	Husband

15265	Divorced	Other-service	Unmarried
20514	Married-civ-spouse	Protective-serv	Husband
22642	Never-married	Exec-managerial	Not-in-family
25709	Married-civ-spouse	Tech-support	Husband
29946	Divorced	Sales	Not-in-family
31715	Never-married	Other-service	Own-child
29794	Divorced	Adm-clerical	Not-in-family
11908	Never-married	Sales	Unmarried
1317	Never-married	NaN	Not-in-family
25132	Divorced	Prof-specialty	Unmarried
14210	Married-civ-spouse	Tech-support	Husband
7184	Married-civ-spouse	Exec-managerial	Husband
16832	Married-civ-spouse	Exec-managerial	Husband
17112	Divorced	Exec-managerial	Unmarried
26083	Married-civ-spouse	Exec-managerial	Husband
7690	Married-civ-spouse	Prof-specialty	Husband

	race	sex	capitalGain	capitalLoss	hoursPerWeek	\
32432	Other	Male	0	0	72	
32467	White	Male	0	0	48	
2449	White	Female	0	0	25	
14720	Amer-Indian-Eskimo	Female	0	0	20	
4862	White	Male	5178	0	90	
29380	White	Male	0	0	40	
27398	White	Male	0	0	40	
25406	White	Male	0	0	40	
23272	White	Female	0	0	40	
21374	White	Male	0	0	40	
19242	Black	Female	0	0	40	
17302	White	Male	0	0	40	
15265	Black	Female	0	0	20	
20514	White	Male	0	0	40	
22642	White	Male	0	0	50	
25709	White	Male	0	0	50	
29946	White	Female	0	0	40	
31715	White	Male	0	0	45	
29794	White	Female	0	0	40	
11908	White	Female	0	0	54	
1317	White	Male	0	0	55	
25132	White	Female	0	0	35	
14210	White	Male	0	0	25	
7184	White	Male	0	0	40	
16832	White	Male	0	0	55	
17112	White	Female	0	323	45	
26083	White	Male	99999	0	40	
7690	White	Male	0	0	45	

	nativeCountry	income
32432	Mexico	<=50K
32467	United-States	<=50K
2449	United-States	<=50K
14720	United-States	<=50K
4862	United-States	>50K
29380	United-States	<=50K
27398	United-States	<=50K
25406	United-States	<=50K
23272	United-States	<=50K
21374	United-States	>50K
19242	United-States	<=50K
17302	United-States	<=50K
15265	United-States	<=50K
20514	United-States	<=50K
22642	United-States	<=50K
25709	United-States	>50K
29946	United-States	<=50K
31715	United-States	<=50K
29794	United-States	<=50K
11908	United-States	<=50K
1317	United-States	<=50K
25132	United-States	<=50K
14210	United-States	>50K
7184	United-States	>50K
16832	United-States	>50K
17112	United-States	<=50K
26083	United-States	>50K
7690	United-States	>50K

Sorts the DataFrame by education-num (numeric education ranking).

f)

Hypothesis (c): Ordered education levels subjectively.

Reality (e): The correct order is defined by education-num

```
[34]: adult_df[['education', 'education-num']].drop_duplicates().
      ↪sort_values('education-num')
```

```
[34]:
```

	education	education-num
224	Preschool	1
160	1st-4th	2
56	5th-6th	3
15	7th-8th	4
6	9th	5
77	10th	6
3	11th	7

415	12th	8
2	HS-grad	9
10	Some-college	10
14	Assoc-voc	11
13	Assoc-acdm	12
0	Bachelors	13
5	Masters	14
52	Prof-school	15
20	Doctorate	16

g)

`adult_df.education` String

`pd.get_dummies(adult_df.education)` Boolean (One-hot)

`adult_df['education']` String

h)

Advantages/Disadvantages of Representations

String Representation (`adult_df.education`):

- Statement: “No bias, but algorithms can’t use it.”
- Pros: Preserves original labels.
- Cons: Cannot be used directly in math models.

One-Hot Encoding (`get_dummies`):

- Statement: “Algorithms can use it, but data size increases.”
- Pros: Works with ML algorithms.
- Cons: Creates many columns; loses ordinality.

Integer Representation (`education-num`):

- Statement: “No size concerns, but assumes linear bias.”
- Pros: Compact; respects order.
- Cons: Implies equal intervals between levels (e.g., “HS-grad \rightarrow Bachelors” = “Bachelors \rightarrow Masters”).