

# Chaper 5 Excercises

April 22, 2025

## 0.1 Hands-On Data Preprocessing in Python

Learn how to effectively prepare data for successful data analytics

AUTHOR: Dr. Roy Jafari

### 0.1.1 Chapter 5: Data Visualization

#### Excercises

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from ipywidgets import interact, widgets
```

## 1 Excercise 1

In this exercise, we will be using Universities\_imputed\_reduced.csv. Draw the following described visualizations.

- Use boxplots to compare the student to faculty ratio (stud./fac. ratio) for the two populations.
- Use a histogram to compare the student to faculty ratio (stud./fac. ratio) for the two populations.
- use subplots to put the results of a and b on top of one another to create a visual that compares the two populations.

```
[2]: uni_df = pd.read_csv('Universities_imputed_reduced.csv')
uni_df.head()
```

```
[2]:
```

	College Name	State	Public/Private	num_appli_rec	\
0	Alaska Pacific University	AK	Private	193	
1	University of Alaska at Fairbanks	AK	Public	1852	
2	University of Alaska Southeast	AK	Public	146	
3	University of Alaska at Anchorage	AK	Public	2065	
4	Alabama Agri. & Mech. Univ.	AL	Public	2817	

	num_appl_accepted	num_new_stud_enrolled	in-state tuition	\
0	146	55	7560	
1	1427	928	1742	
2	117	89	1742	
3	1598	1162	1742	
4	1920	984	1700	

	out-of-state tuition	% fac. w/PHD	stud./fac. ratio	Graduation rate
0	7560	76	11.9	15
1	5226	67	10.0	60
2	5226	39	9.5	39
3	5226	48	13.7	60
4	3400	53	14.3	40


a.

```
[3]: private_public_possiblilites = uni_df['Public/Private'].unique()

box_sr = pd.Series('', index=private_public_possiblilites)

for poss in private_public_possiblilites:
    BM = uni_df['Public/Private'] == poss
    box_sr[poss] = uni_df[BM]['stud./fac. ratio']

plt.boxplot(box_sr, vert=False)
plt.yticks([1, 2], private_public_possiblilites)
plt.xlabel('stud./fac. ratio')
plt.ylabel('Public/Private')
plt.show()
```




b.

```
[4]: private_public_possiblilites = uni_df['Public/Private'].unique()

for poss in private_public_possiblilites:
    BM = uni_df['Public/Private'] == poss
    plt.hist(uni_df[BM]['stud./fac. ratio'],
             histtype='step', label=poss)

plt.legend()
plt.xlabel('stud./fac. ratio')
plt.ylabel('Public/Private')
plt.show()
```



c.

```
[5]: private_public_possiblilites = uni_df['Public/Private'].unique()

dataForBox_dic = {}
```

```


for poss in private_public_possiblilites:
    BM = uni_df['Public/Private'] == poss
    dataForBox_dic[poss] = uni_df[BM]['stud./fac. ratio']

plt.subplot(2, 1, 1)
plt.boxplot(dataForBox_dic.values(), vert=False)
plt.yticks([1, 2], private_public_possiblilites)

plt.subplot(2, 1, 2)
for poss in private_public_possiblilites:
    BM = uni_df['Public/Private'] == poss
    plt.hist(uni_df[BM]['stud./fac. ratio'],
             histtype='step', label=poss)
plt.legend()

plt.tight_layout()
plt.show()

```




## 2 Exercise 2

In this exercise, we will continue using `Universities_imputed_reduced.csv`. Draw the following described visualizations.

- Use a bar chart to compare the private/public ratio of all the states in the dataset. In the bar chart, the x-axis represents the percentage of public and private schools, and the y-axis represents the states.
- Improve the visualizations by sorting the states on the visuals based on the total number of schools.
- Create a stacked bar chart that shows the compare the percentages of public and private schools.

a.


```
[6]: uni_df.groupby(['State', 'Public/Private']).size().unstack().plot.  
      ↪ barh(stacked=True)  
plt.legend(loc=4)  
plt.show()
```



b.

```
[7]: state_counts = uni_df.groupby(['State', 'Public/Private']).size().unstack()  
state_counts['Total'] = state_counts.sum(axis=1)  
  
state_counts.sort_values(by='Total', ascending=False).drop(columns='Total').  
      ↪ plot.barh(stacked=True)  
plt.legend(loc=4)
```


```
plt.show()
```



c.

```
[8]: state_counts = uni_df.groupby(['State', 'Public/Private']).size().unstack()
state_counts['Total'] = state_counts.sum(axis=1)
state_percentages = state_counts.div(state_counts['Total'], axis=0) * 100
state_percentages = state_percentages.drop(columns='Total')
state_percentages = state_percentages.sort_values(by='Private', ascending=False)

state_percentages.plot.barh(stacked=True)
plt.legend(loc=4)
plt.show()
```



### 3 Exercice 3

For this example, we will be using WH Report\_preprocessed.csv. Draw the following described visualizations.

- Create a visual that compares the relationship between all the happiness indices.
- Use the visual you created in a) to report the happiness indices with strong relationships
- Confirm the relationship you found and described by calculating their correlation coefficient

```
[9]: report_df = pd.read_csv('WH Report_preprocessed.csv')
report_df.head()
```

```
[9]:
```

	Name	Continent	year	population	Life_Ladder	Log_GDP_per_capita	\
0	Afghanistan	Asia	2010	29185507.0	4.758	7.647	
1	Afghanistan	Asia	2011	30117413.0	3.832	7.620	
2	Afghanistan	Asia	2012	31161376.0	3.783	7.705	
3	Afghanistan	Asia	2013	32269589.0	3.572	7.725	
4	Afghanistan	Asia	2014	33370794.0	3.131	7.718	

	Social_support	Healthy_life_expectancy_at_birth	\
0	0.539	51.60	
1	0.521	51.92	

2	0.521	52.24
3	0.484	52.56
4	0.526	52.88


	Freedom_to_make_life_choices	Generosity	Perceptions_of_corruption \
0	0.600	0.121	0.707
1	0.496	0.162	0.731
2	0.531	0.236	0.776
3	0.578	0.061	0.823
4	0.509	0.104	0.871

	Positive_affect	Negative_affect
0	0.618	0.275
1	0.611	0.267
2	0.710	0.268
3	0.621	0.273
4	0.532	0.375

a.

```
[10]: happiness_indices = [
        'Life_Ladder',
        'Log_GDP_per_capita',
        'Social_support',
        'Healthy_life_expectancy_at_birth',
        'Freedom_to_make_life_choices',
        'Generosity',
        'Perceptions_of_corruption',
        'Positive_affect',
        'Negative_affect'
    ]
sns.pairplot(report_df[happiness_indices])
plt.show()
```





b.

```
[11]: # base on a
```

c.

```
[12]: corr_matrix = report_df[happiness_indices].corr()
corr_matrix
```

```
[12]:
```

	Life_Ladder	Log_GDP_per_capita	\
Life_Ladder	1.000000	0.798912	
Log_GDP_per_capita	0.798912	1.000000	
Social_support	0.723686	0.718969	
Healthy_life_expectancy_at_birth	0.758287	0.857981	

Freedom_to_make_life_choices	0.518618	0.357799
Generosity	0.198072	0.010562
Perceptions_of_corruption	-0.465268	-0.368602
Positive_affect	0.518226	0.296845
Negative_affect	-0.302013	-0.261958

	Social_support \
Life_Ladder	0.723686
Log_GDP_per_capita	0.718969
Social_support	1.000000
Healthy_life_expectancy_at_birth	0.629507
Freedom_to_make_life_choices	0.421854
Generosity	0.099973
Perceptions_of_corruption	-0.258575
Positive_affect	0.429687
Negative_affect	-0.425569

	Healthy_life_expectancy_at_birth \
Life_Ladder	0.758287
Log_GDP_per_capita	0.857981
Social_support	0.629507
Healthy_life_expectancy_at_birth	1.000000
Freedom_to_make_life_choices	0.393043
Generosity	0.018837
Perceptions_of_corruption	-0.353142
Positive_affect	0.339499
Negative_affect	-0.209444

	Freedom_to_make_life_choices	Generosity \
Life_Ladder	0.518618	0.198072
Log_GDP_per_capita	0.357799	0.010562
Social_support	0.421854	0.099973
Healthy_life_expectancy_at_birth	0.393043	0.018837
Freedom_to_make_life_choices	1.000000	0.325176
Generosity	0.325176	1.000000
Perceptions_of_corruption	-0.504291	-0.296068
Positive_affect	0.635665	0.359233
Negative_affect	-0.313267	-0.121400

	Perceptions_of_corruption	Positive_affect \
Life_Ladder	-0.465268	0.518226
Log_GDP_per_capita	-0.368602	0.296845
Social_support	-0.258575	0.429687
Healthy_life_expectancy_at_birth	-0.353142	0.339499
Freedom_to_make_life_choices	-0.504291	0.635665
Generosity	-0.296068	0.359233
Perceptions_of_corruption	1.000000	-0.320755

Positive_affect	-0.320755	1.000000
Negative_affect	0.345491	-0.372535

	Negative_affect
Life_Ladder	-0.302013
Log_GDP_per_capita	-0.261958
Social_support	-0.425569
Healthy_life_expectancy_at_birth	-0.209444
Freedom_to_make_life_choices	-0.313267
Generosity	-0.121400
Perceptions_of_corruption	0.345491
Positive_affect	-0.372535
Negative_affect	1.000000

### Strong & Notable Relationships Among Happiness Indices

1. Life Ladder & Log GDP per Capita Correlation coefficient: +0.80

Interpretation: A very strong positive relationship. Countries with higher income per person tend to report greater happiness. Economic stability clearly supports life satisfaction.

2. Life Ladder & Healthy Life Expectancy Correlation coefficient: +0.76

Interpretation: People in countries with longer, healthier lives tend to be happier. This emphasizes the impact of health systems and longevity on well-being.

3. Life Ladder & Social Support Correlation coefficient: +0.72

Interpretation: Strong social connections contribute significantly to happiness. Societies with tight-knit communities and supportive networks foster greater life satisfaction.

4. Log GDP per Capita & Healthy Life Expectancy Correlation coefficient: +0.86

Interpretation: The strongest correlation in your data. Countries with high economic wealth also have high life expectancy — likely due to better healthcare, nutrition, and living conditions.

5. Life Ladder & Freedom to Make Life Choices Correlation coefficient: +0.52

Interpretation: People who feel more in control of their lives tend to be happier. This moderate relationship highlights the importance of autonomy and civil liberties.

6. Life Ladder & Positive Affect Correlation coefficient: +0.52

Interpretation: Happier people tend to experience more positive emotions (joy, laughter, etc.). Though not a perfect match, the emotional tone of life plays a big role.

7. Life Ladder & Perceptions of Corruption Correlation coefficient: -0.47

Interpretation: A moderate negative relationship. Higher levels of perceived corruption are associated with lower happiness, likely due to mistrust and systemic dissatisfaction.

### Weaker or Less Clear Relationships

1. Generosity & Life Ladder: +0.20

Suggests generosity doesn't directly track with happiness, or that it varies a lot by culture/context.

2. Negative Affect & Life Ladder:  $-0.30$

A mild negative relationship: more negative emotions slightly predict lower happiness, as expected.

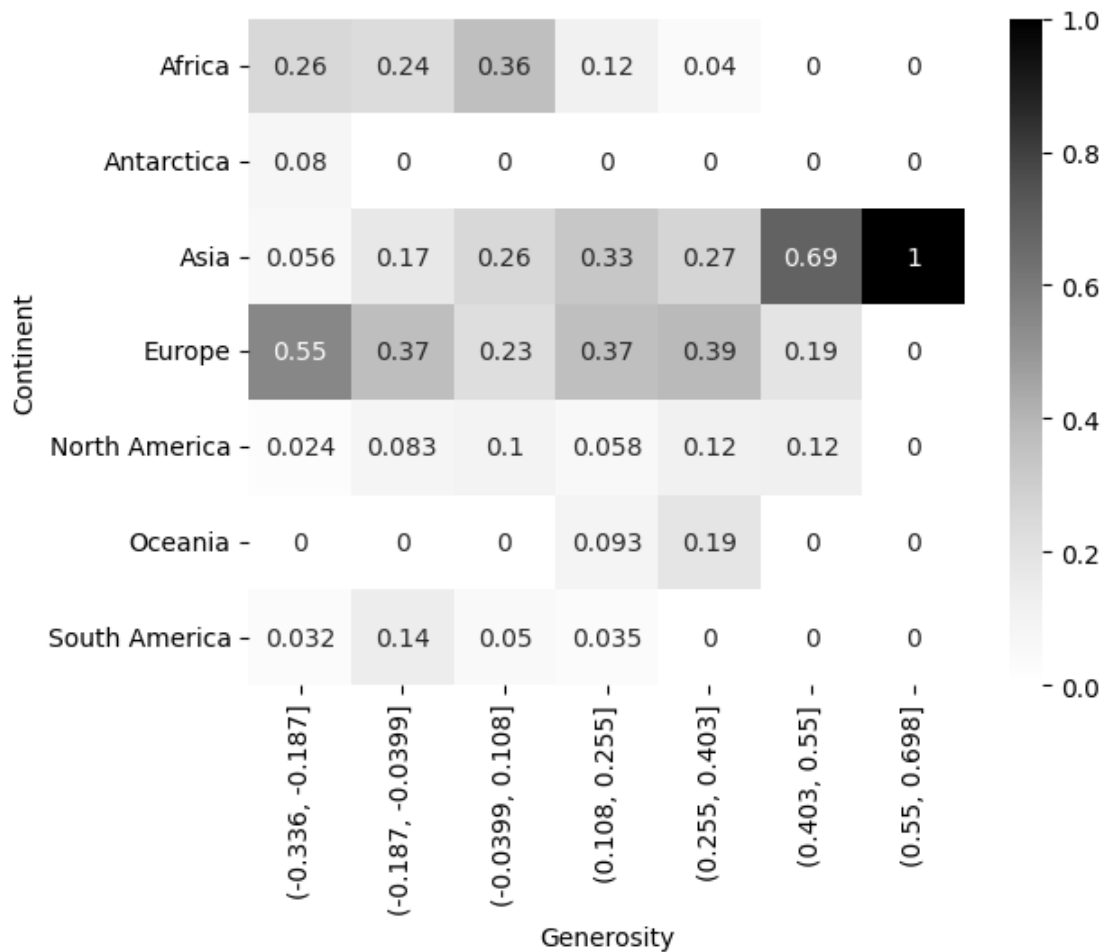
## 4 Exercise 4

For this exercise, we will continue using WH Report\_preprocessed.csv. Draw the following described visualizations.

- Draw a visual that examine the relationship between the two attributes Continent and Generosity.
- Based on the visual, is there a relationship between the two attributes? Explain why.

a.

```
[13]: generosity_discretized = pd.cut(report_df['Generosity'], bins=7)
contingency_tbl = pd.crosstab(report_df['Continent'], generosity_discretized)
probability_tbl = contingency_tbl / contingency_tbl.sum()
sns.heatmap(probability_tbl, annot=True, center=0.5, cmap="Greys")
plt.show()
```



b.

Yes, there is a relationship between Continent and Generosity, as seen in the heatmap and table. Some continents have a higher likelihood of countries falling into specific generosity ranges than others — which indicates a non-random distribution

Continent	Pattern in Generosity Distribution
<b>Africa</b>	Majority falls in lower bins (-0.336 to 0.108) – <b>high concentration in low generosity ranges.</b>
<b>Asia</b>	<b>Spread across the entire spectrum, only region with values in the highest bin (0.55, 0.698].</b> Indicates diverse and high generosity levels.
<b>Europe</b>	Strong presence in low to moderate bins, but <b>absent from the top generosity bin.</b>
<b>North America</b>	Appears more in <b>lower-middle generosity ranges</b> , no countries in top bins.
<b>South America</b>	Heavily weighted toward <b>lower generosity</b> , barely any representation in higher bins.
<b>Oceania</b>	<b>Very limited data</b> , but slight mid-range generosity.
<b>Antarctica</b>	Likely an outlier or placeholder (ignore in analysis).

## 5 Exercise 5

For this exercise, we will be using `whickham.csv`. Draw the following described visualizations.

- What is the numerical attribute in this dataset? Draw two different plots that summarize the data.
- What are the categorical attributes in this dataset? Draw a plot per attribute that summarizes the data.
- Draw a visual that examine the relationship between `outcome` and `smoker`. Do you notice anything surprising?
- To demystify the surprising relationship you observed on c) run the following code, and store the result in `person_df`.

```
person_df = pd.read_csv('whickham.csv') person_df['age_discretized'] =  
pd.cut(person_df.age, bins = 4, labels=False) person_df.groupby(['age_discretized', 'smoker']).  
plt.show()
```

- Using the visual that was created under d) explain the surprising observation under c).
- How many dimensions the visual that was created under d) has? How did we manage to add dimensions?


```
[14]: person_df = pd.read_csv('whickham.csv')  
person_df.head()
```

```
[14]:  outcome  smoker  age  
0   Alive    Yes   23  
1   Alive    Yes   18  
2    Dead    Yes   71  
3   Alive    No   67  
4   Alive    No   64
```

a.


```
[15]: numerical_att = 'age'  
person_df[numerical_att].plot.hist()
```

```
[15]: <Axes: ylabel='Frequency'>
```



```
[16]: person_df[numerical_att].plot.box(vert=False)
```

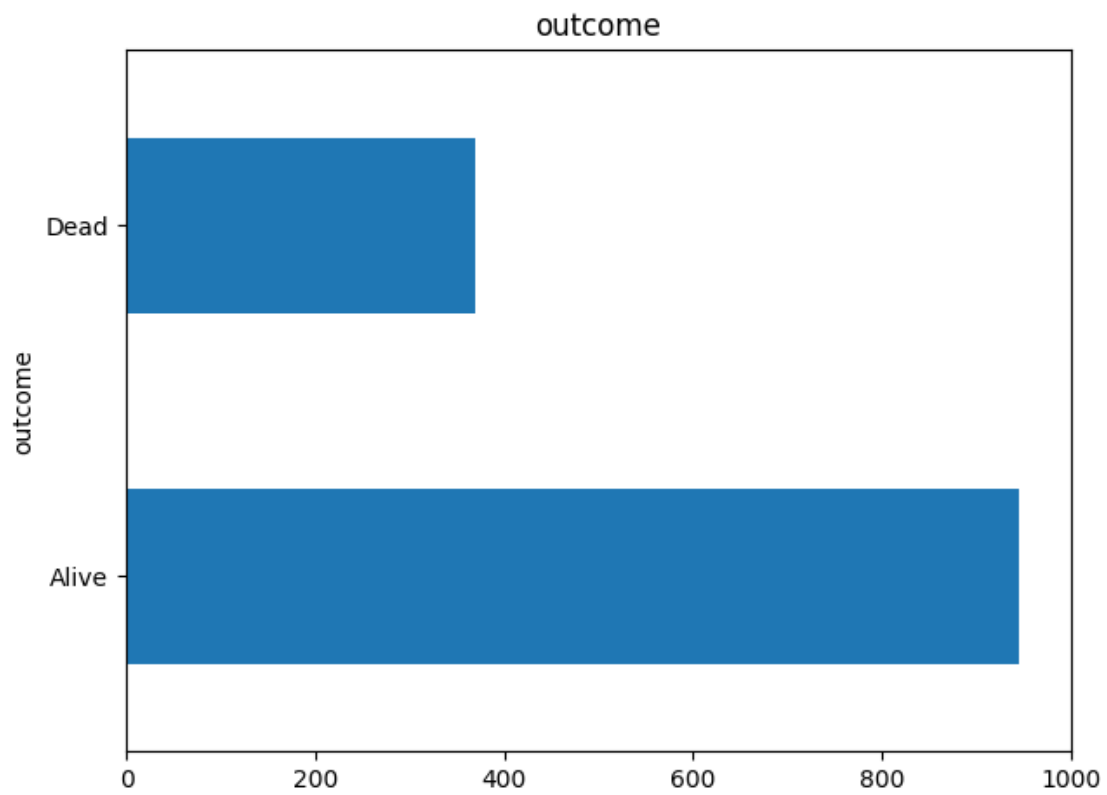
```
[16]: <Axes: >
```




b.

```
[17]: categorical_attributes = ['outcome', 'smoker']

for att in categorical_attributes:
    person_df[att].value_counts().plot.barh()
    plt.title(att)
    plt.tight_layout()
    plt.xlim(0, 1000)
    plt.show()
```

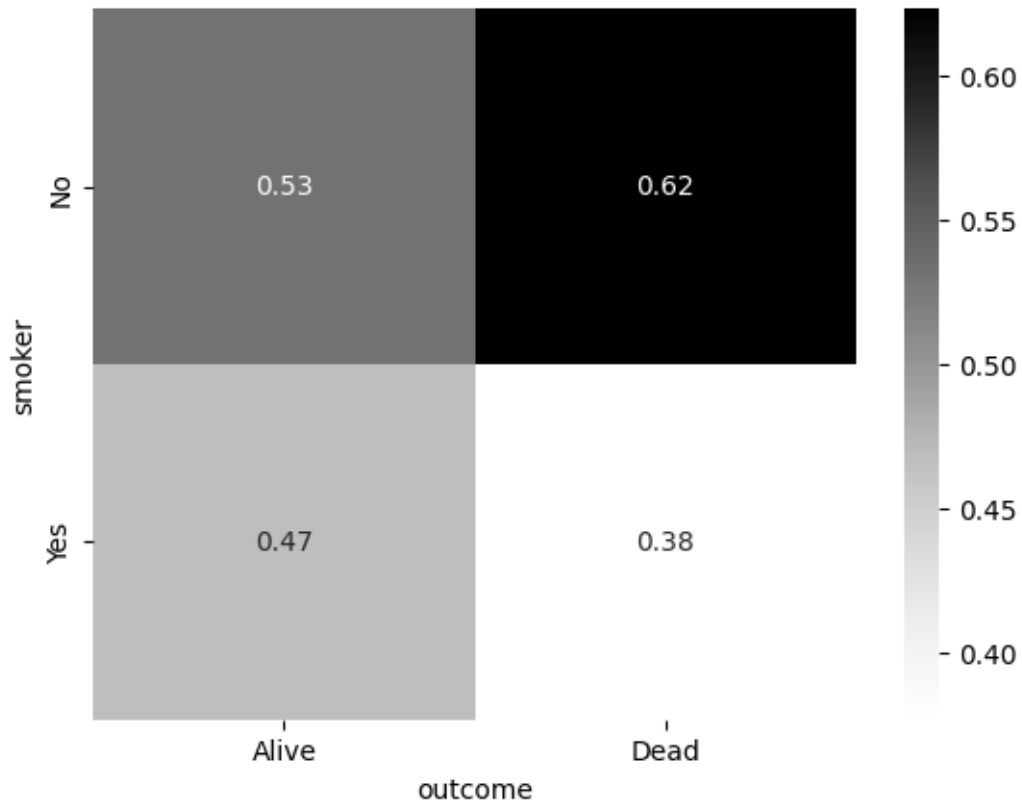






c.

```
[18]: contingency_tbl = pd.crosstab(person_df['smoker'], person_df['outcome'])  
probability_tbl = contingency_tbl / contingency_tbl.sum()  
sns.heatmap(probability_tbl, annot=True, center=0.5, cmap="Greys")  
plt.show()
```




More non-smokers died than smokers (contrary to expectations).

But this is misleading because it doesn't account for age distribution (smokers might be younger).

d.

```
[19]: person_df = pd.read_csv('whickham.csv')
person_df['age_discretized'] = pd.cut(person_df.age, bins = 4, labels=False)
person_df.groupby(['age_discretized', 'smoker']).outcome.value_counts().
    ↪unstack().unstack().plot.bar(stacked=True)
plt.show()
```



e.

The surprising observation (smokers appearing to survive more) is due to:

1. Age Confounding:

- Smokers are younger (mean age = 45) vs. non-smokers (mean age = 50).
- Younger people naturally survive longer, regardless of smoking.

2. Simpson's Paradox:

- Aggregating data hides subgroup trends.
- When stratified by age, smoking reduces survival in every age group.

Example:

Age Group 0 (youngest): Smokers have 90% survival vs. non-smokers' 95%.

Age Group 3 (oldest): Smokers have 20% survival vs. non-smokers' 30%.

f.

Number of Dimensions: 3

1. Primary Dimension (x-axis): age\_discretized (4 age groups).
2. Secondary Dimension (color stacks): smoker (Yes/No).

3. Tertiary Dimension (stack height): outcome (Alive/Dead).

How Dimensions Were Added:

- `groupby()`: Split data by `age_discretized` and `smoker` (2 dimensions).
- `value_counts()` + `unstack()`: Pivoted outcome into columns (3rd dimension).
- `plot.bar(stacked=True)`: Visualized stacks for Alive/Dead within each smoker-age group.

## 6 Exercise 6

For this exercise, we will be using `WH Report_preprocessed.csv`.

- Use this dataset to create a 5-dimensional scatterplot to show the interactions between the variables.
- Interact with and study the visual you created under a) and report your observations.

```
[20]: report_df = pd.read_csv('WH Report_preprocessed.csv')
      report_df.head()
```

```
[20]:
```

	Name	Continent	year	population	Life_Ladder	Log_GDP_per_capita	\
0	Afghanistan	Asia	2010	29185507.0	4.758	7.647	
1	Afghanistan	Asia	2011	30117413.0	3.832	7.620	
2	Afghanistan	Asia	2012	31161376.0	3.783	7.705	
3	Afghanistan	Asia	2013	32269589.0	3.572	7.725	
4	Afghanistan	Asia	2014	33370794.0	3.131	7.718	

	Social_support	Healthy_life_expectancy_at_birth	\
0	0.539	51.60	
1	0.521	51.92	
2	0.521	52.24	
3	0.484	52.56	
4	0.526	52.88	

	Freedom_to_make_life_choices	Generosity	Perceptions_of_corruption	\
0	0.600	0.121	0.707	
1	0.496	0.162	0.731	
2	0.531	0.236	0.776	
3	0.578	0.061	0.823	
4	0.509	0.104	0.871	

	Positive_affect	Negative_affect
0	0.618	0.275
1	0.611	0.267
2	0.710	0.268
3	0.621	0.273
4	0.532	0.375

a.

```

[39]: import pandas as pd
import matplotlib.pyplot as plt
from ipywidgets import interact, IntSlider

# Define continent colors
continent_poss = report_df['Continent'].unique()
colors_dic = {
    'Asia': 'blue',
    'Europe': 'green',
    'Africa': 'red',
    'South America': 'cyan',
    'Oceania': 'magenta',
    'North America': 'yellow',
    'Antarctica': 'black'
}

# Sort by population for better visualization
report_df = report_df.sort_values('population', ascending=False)

def plot_year(year):
    plt.figure(figsize=(12, 8))

    for continent in continent_poss:
        # Filter data for year and continent
        mask = (report_df['year'] == year) & (report_df['Continent'] ==
↪continent)
        df_filtered = report_df[mask]

        # Skip if no data for this continent-year combination
        if len(df_filtered) == 0:
            continue

        # Create scatter plot with multiple dimensions
        plt.scatter(
            x=df_filtered['Healthy_life_expectancy_at_birth'], # x-axis
            y=df_filtered['Life_Ladder'], # y-axis
            s=df_filtered['population']/500000, # marker size
↪(scaled)
            c=df_filtered['Social_support'], # marker color
            cmap='viridis', # color map
            vmin=0, vmax=1, # color scale
↪limits
            alpha=0.7, # transparency
            edgecolors='w', # white borders
            linewidths=0.5, # border width
            label=continent
        )

```

```

# Formatting
plt.title(f'World Happiness Report - {year}')
plt.xlabel('Healthy Life Expectancy at Birth')
plt.ylabel('Life Ladder (Happiness Score)')
plt.xlim(40, 85)
plt.ylim(2, 8)
plt.grid(alpha=0.2)

# Add colorbar for Social_support
plt.colorbar(label='Social Support')

# Add legend for continents
plt.legend(title='Continent', markerscale=0.7, loc=2)

plt.show()

# Create interactive widget
interact(
    plot_year,
    year=IntSlider(min=report_df['year'].min(),
                    max=report_df['year'].max(),
                    step=1,
                    value=report_df['year'].median())
)

```

```

interactive(children=(IntSlider(value=2014, description='year', max=2019,
    min=2010), Output()), _dom_classes=(...

```

```
[39]: <function __main__.plot_year(year)>
```

b.

```
[35]: # based on a
```

## 7 Exercise 7

For this exercise, we will continue using WH Report\_preprocessed.csv.

- Create a visual that shows the trend of change for the attribute Generosity for all the countries.
- Add three more line plots to the previous visual using the color blue and a thicker line (3px).

Figure 5. 23. Line plot comparing Generosity across all countries in 2010 and 2019 with emphasis on the United States, India, and China

- Report your observations from the visual. Make sure to employ all of the line plots (grey and blue).

```

[71]: import pandas as pd
import matplotlib.pyplot as plt

# Load data
country_df = pd.read_csv('WH_Report_preprocessed.csv')


# Aggregate Generosity by Continent and Year
byCountryYear_df = country_df.groupby(['Name', 'year'])['Generosity'].mean()

# Plot setup
plt.figure(figsize=(5, 7))

# Plot lines connecting 2010-2019 for each continent
for i, continent in enumerate(byCountryYear_df.index.get_level_values(0).
    ↪unique()):
    plt.plot([2010, 2019],
             byCountryYear_df.loc[continent, [2010, 2019]],
             color='grey', # Using grey for all lines
             )

# Formatting
plt.xticks([2010, 2019])
plt.title('Generosity Trends by Country (2010 vs 2019)')
plt.ylabel('Average Generosity Score')
plt.grid(alpha=0.2)

```



b.

```
[72]: import pandas as pd
import matplotlib.pyplot as plt

# Load data
country_df = pd.read_csv('WH_Report_preprocessed.csv')

# Aggregate Generosity by Country and Year
byCountryYear_df = country_df.groupby(['Name', 'year'])['Generosity'].mean()

# Plot setup
```



```

plt.figure(figsize=(5, 7))

# Plot grey lines for all countries
for country in byCountryYear_df.index.get_level_values(0).unique():
    plt.plot([2010, 2019],
              byCountryYear_df.loc[country, [2010, 2019]],
              color='grey',
              linewidth=0.8)

highlight_markers = {
    'United States': '*', # Star
    'China': 's', # Square
    'India': 'o' # Circle
}

# Highlight specific countries in blue
for country, marker in highlight_markers.items():
    plt.plot([2010, 2019],
              byCountryYear_df.loc[country, [2010, 2019]],
              color='blue',
              marker=marker,
              linewidth=1.8,
              label=country)


# Formatting
plt.xticks([2010, 2019])
plt.title('Generosity Trends (2010 vs 2019)\nGrey: All Countries | Blue: ↪Highlighted Countries')
plt.ylabel('Average Generosity Score')
plt.grid(alpha=0.2)

# Add legend only for highlighted countries
plt.legend(loc=9)

plt.tight_layout()
plt.show()

```

Generosity Trends (2010 vs 2019)  
Grey: All Countries | Blue: Highlighted Countries



c.

[57]: # refer to b