

Chapter 5

April 21, 2025

1 Hands-On Data Preprocessing in Python

Learn how to effectively prepare data for successful data analytics

AUTHOR: Dr. Roy Jafari

2 Chapter 5: Data Visualization

2.1 Summarizing a population



2.1.1 Example of summarizing numerical attributes



```
[4]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```


```
[5]: adult_df = pd.read_csv('adult.csv')


numerical_attributes = ['age', 'fnlwgt', 'education-num', 'capitalGain',
                        'capitalLoss', 'hoursPerWeek']


for att in numerical_attributes:
    plt.subplot(2,1,1)
    adult_df[att].plot.hist()
    plt.subplot(2,1,2)
    adult_df[att].plot.box(vert=False)
plt.tight_layout()
plt.savefig('{} .png'.format(att), dpi=600)
plt.show()
```










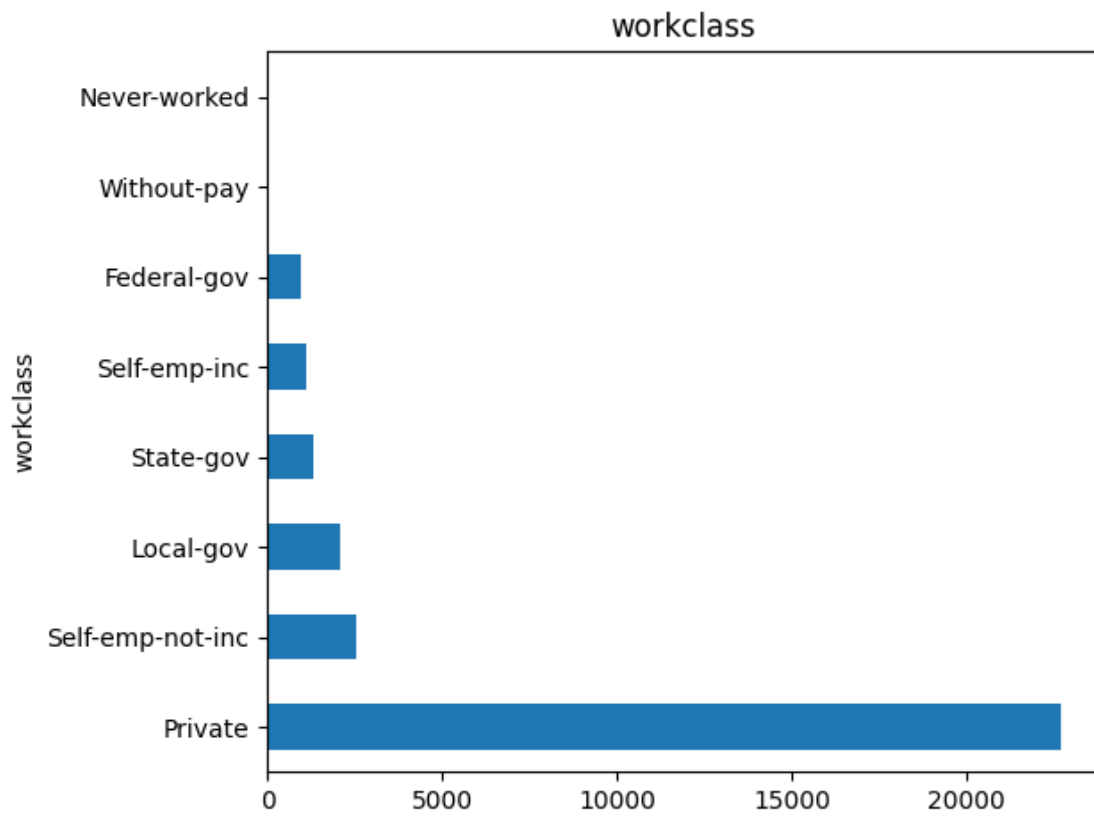


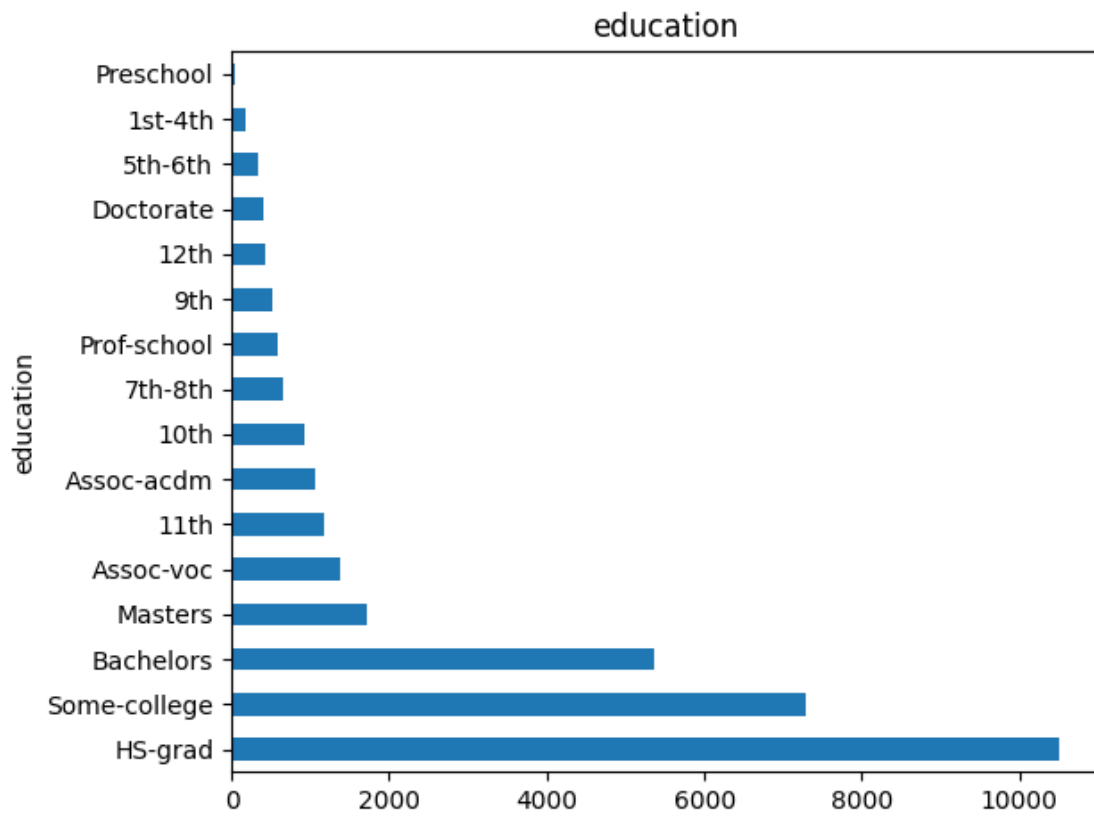



2.1.2 Example of summarizing categorical attributes


```
[6]: categorical_attributes = ['workclass', 'education', 'marital-status',
                              'occupation', 'relationship', 'race',
                              'sex', 'nativeCountry', 'income']

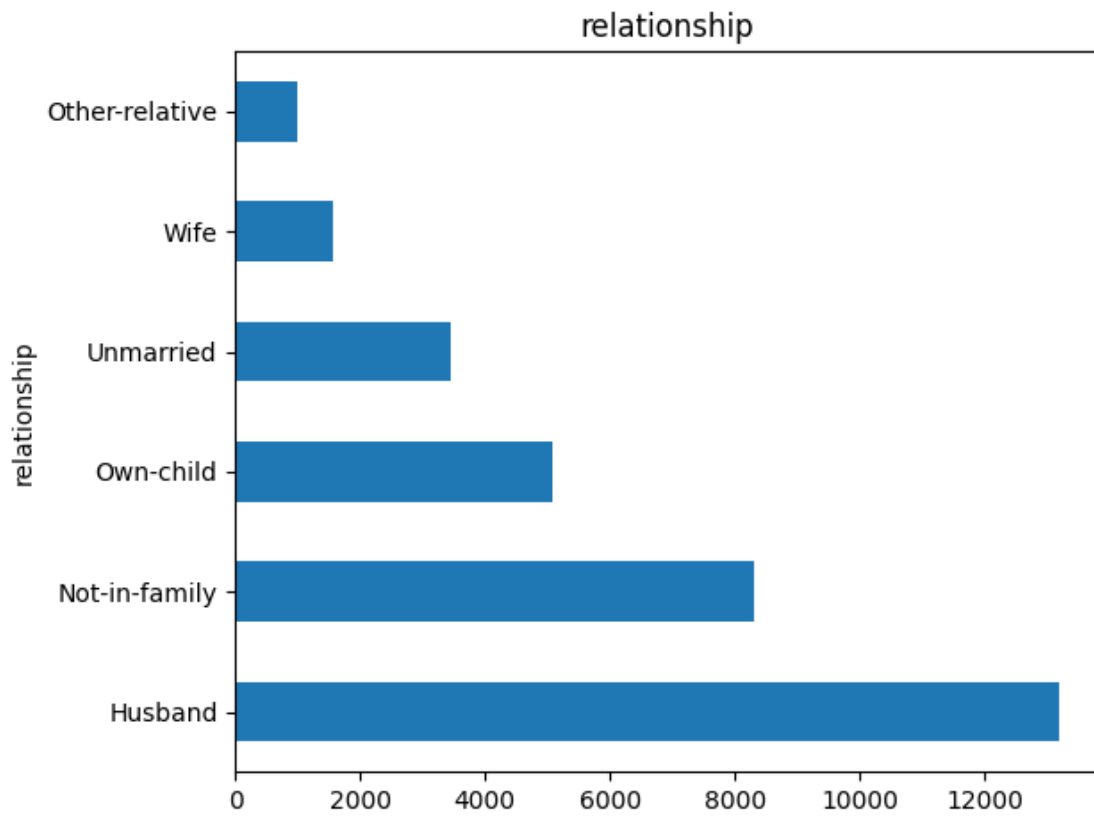
for att in categorical_attributes:
    adult_df[att].value_counts().plot.barh()
    plt.title(att)
    plt.tight_layout()
    plt.savefig('{} .png'.format(att), dpi=600)
    plt.show()
```

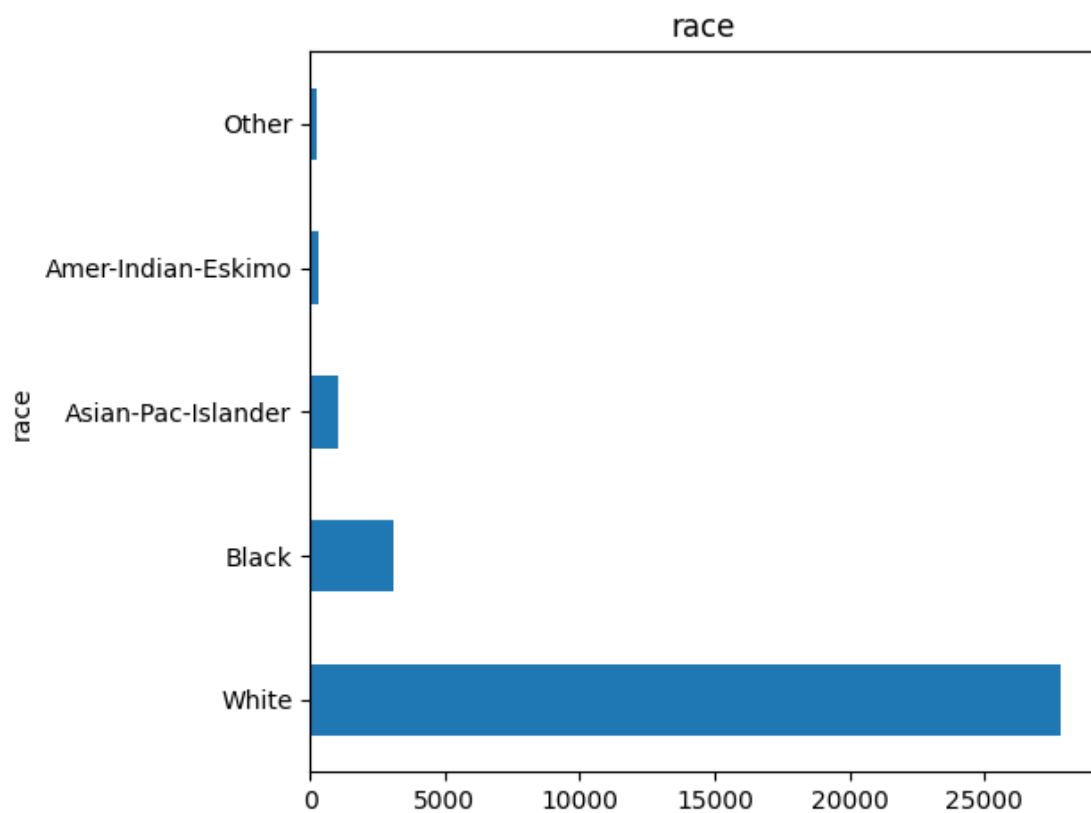


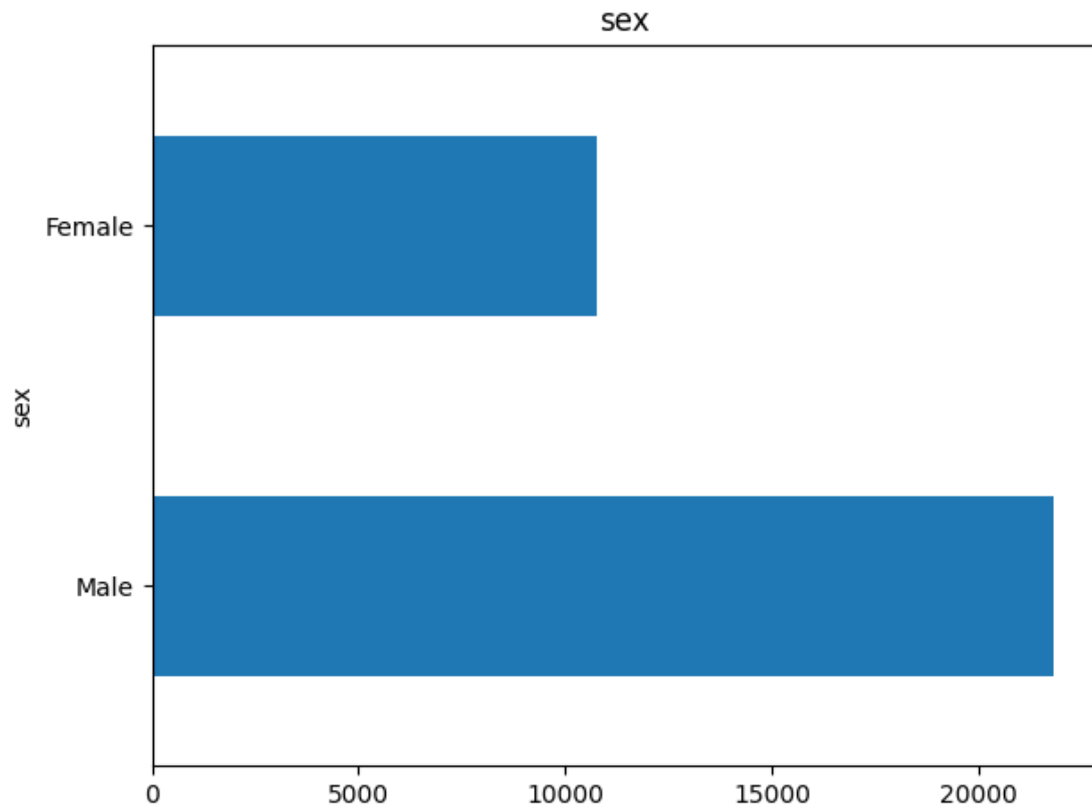


















2.2 Comparing populations


2.2.1 Example of comparing populations using boxplots

```
[7]: income_possibilities = adult_df.income.unique()

box_sr = pd.Series('', index = income_possibilities)

for poss in income_possibilities:
    BM = adult_df.income == poss
    box_sr[poss] = adult_df[BM]['education-num']

plt.boxplot(box_sr, vert=False)
plt.yticks([1,2], income_possibilities)
plt.show()
```





```
[8]: income_possibilities = adult_df.income.unique()

dataForBox_dic= {}

for poss in income_possibilities:
    BM = adult_df.income == poss
    dataForBox_dic[poss] = adult_df[BM]['education-num']


plt.boxplot(dataForBox_dic.values(),vert=False)
plt.yticks([1,2],income_possibilities)
plt.show()
```



2.2.2 Example of comparing populations using histograms

```
[9]: income_possibilities = adult_df.income.unique()

for poss in income_possibilities:
    BM = adult_df.income == poss
    plt.hist(adult_df[BM]['education-num'],
             histtype='step', label=poss)
plt.legend()
plt.show()
```



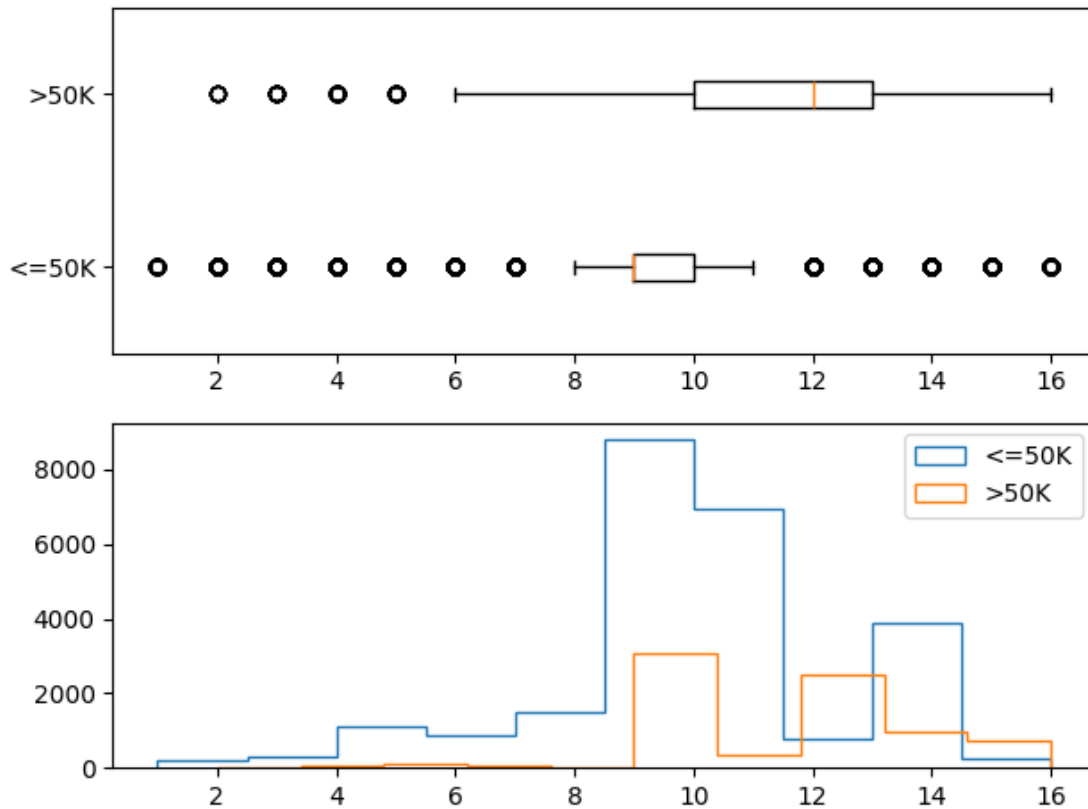
```
[10]: income_possibilities = adult_df.income.unique()

dataForBox_dic= {}

for poss in income_possibilities:
    BM = adult_df.income == poss
    dataForBox_dic[poss] = adult_df[BM]['education-num']

plt.subplot(2,1,1)
plt.boxplot(dataForBox_dic.values(),vert=False)
plt.yticks([1,2],income_possibilities)

plt.subplot(2,1,2)
for poss in income_possibilities:
    BM = adult_df.income == poss
    plt.hist(adult_df[BM]['education-num'],
             histtype='step',label=poss)
plt.legend()
plt.tight_layout()
plt.show()
```

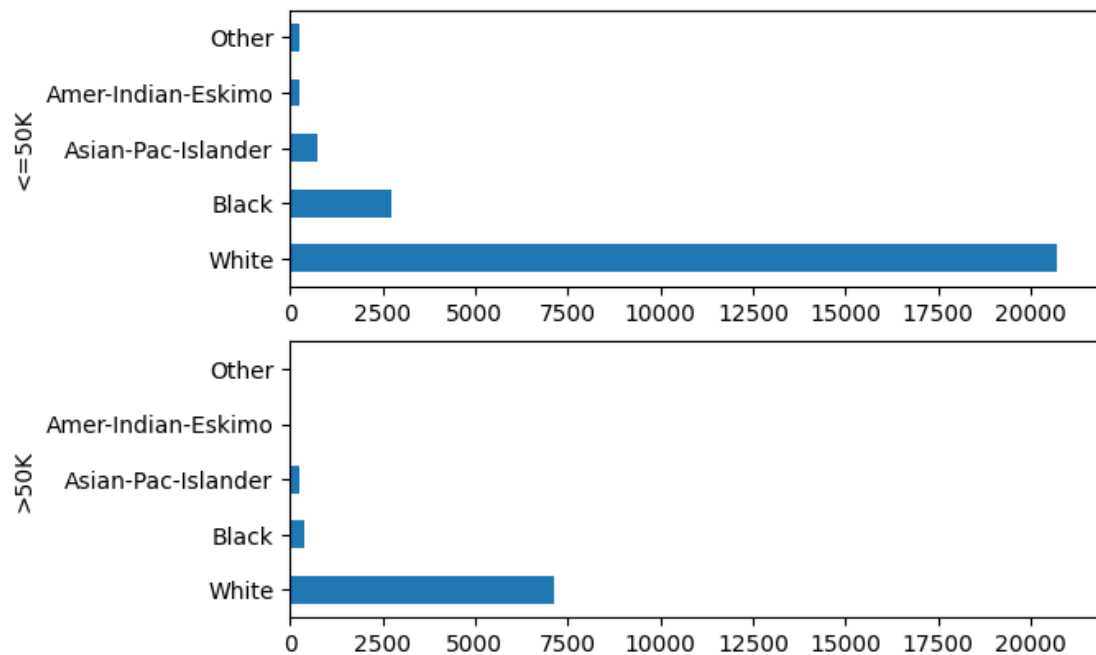


2.2.3 Example of comparing populations using bar charts

The first way of solving


```
[11]: income_possibilities = adult_df.income.unique()

for i, poss in enumerate(income_possibilities):
    plt.subplot(2,1,i+1)
    BM = adult_df.income == poss
    adult_df[BM].race.value_counts().plot.barh()
    plt.xlim([0,22000])
    plt.ylabel(poss)
```




The second way of solving

```
[12]: adult_df.groupby(['income', 'race']).size().plot.barh()
plt.show()
```




The third way of solving

```
[13]: adult_df.groupby(['income', 'race']).size().unstack().plot.barh()  
plt.show()
```



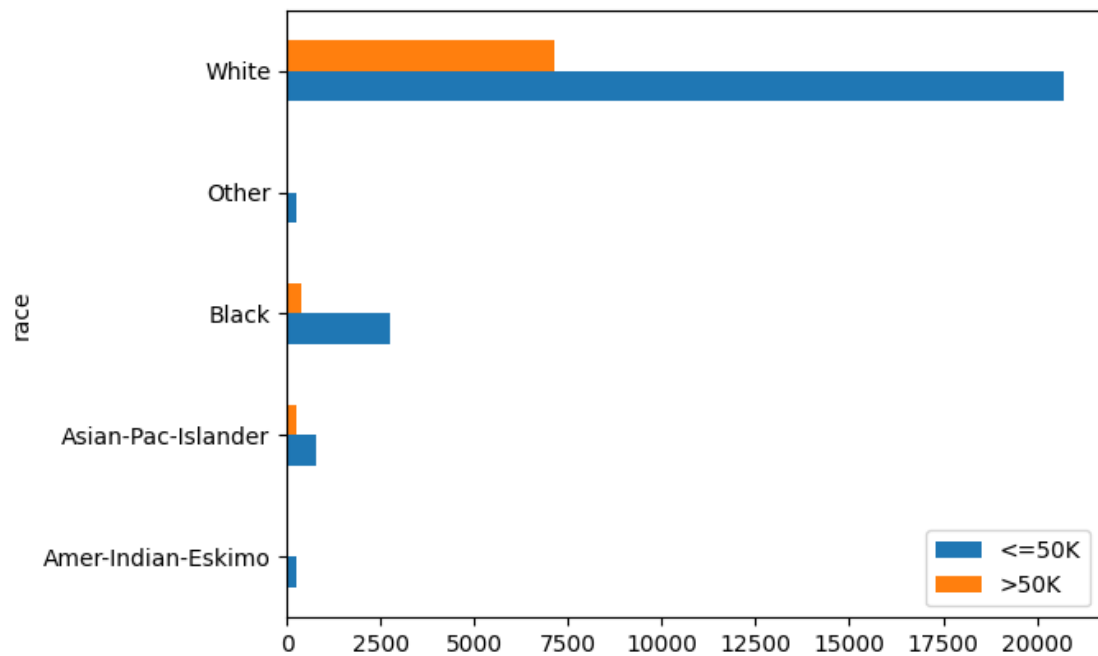
The fourth way of solving

```
[14]: adult_df.groupby(['race', 'income']).size().plot.barh()  
plt.show()
```




The fifth way of solving

```
[15]: adult_df.groupby(['race', 'income']).size().unstack().plot.barh()
plt.legend(loc=4)
plt.show()
```



The sixth way of solving

```
[16]: adult_df.groupby(['race', 'income']).size().unstack().plot.barh(stacked=True)
plt.legend(loc=4)
plt.show()
```

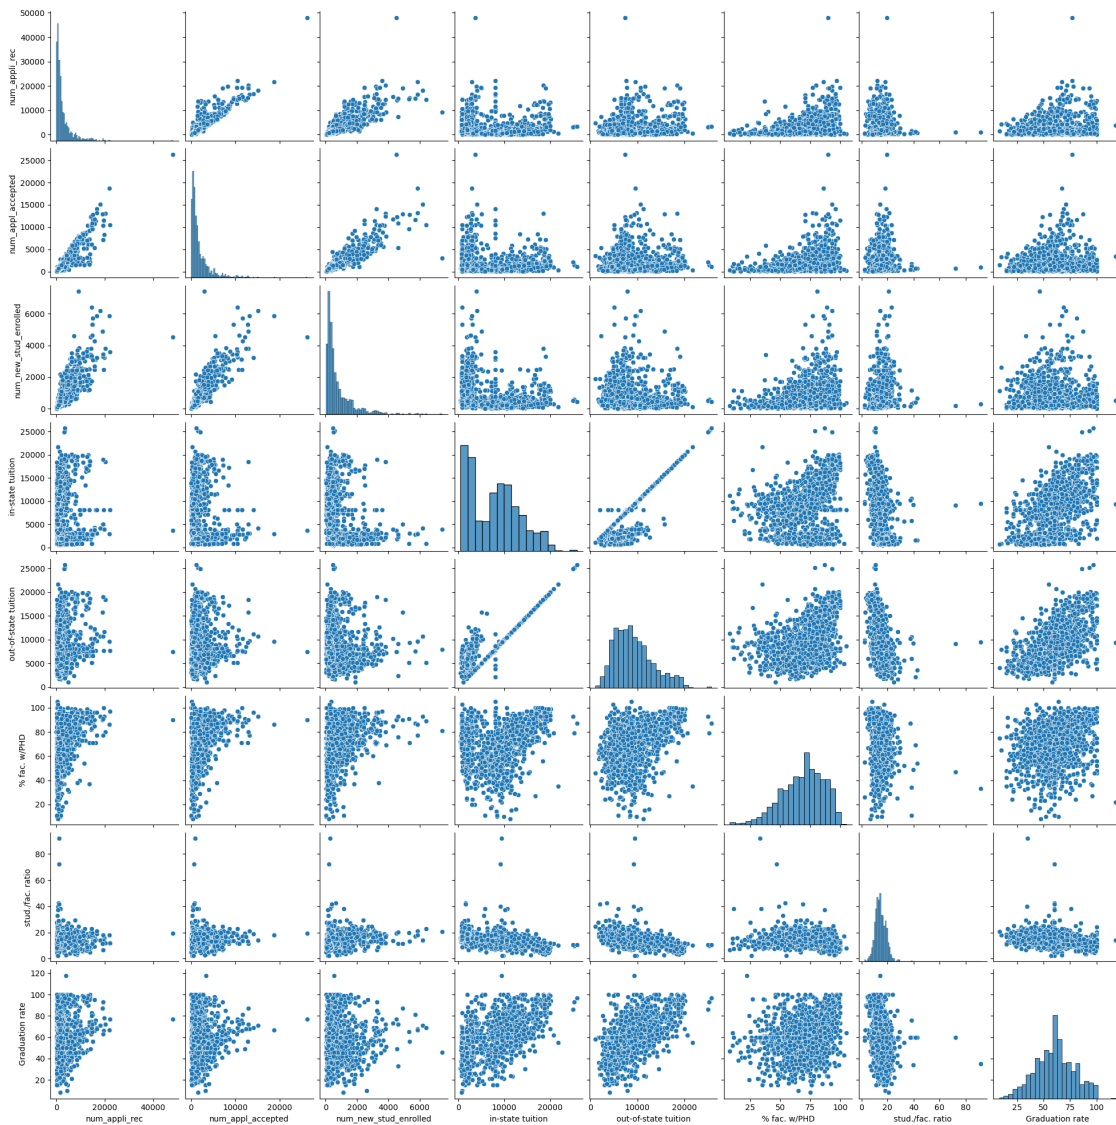


2.3 Investigating the relationship between two attributes

2.3.1 Visualizing the relationship between two numerical attributes

Example of using scatterplots to investigate between the numerical attributes

```
[17]: import seaborn as sns
uni_df = pd.read_csv('Universities_imputed_reduced.csv')
sns.pairplot(uni_df)
plt.show()
```


2.3.2 Visualizing the relationship between two categorical attributes

Example of using contingency table to examine the relationship between two categorical (binary) attributes

```
[18]: contingency_tbl = pd.crosstab(adult_df.income, adult_df.sex)
contingency_tbl
```

```
[18]: sex      Female    Male
income
<=50K      9592    15128
>50K       1179     6662
```

```
[19]: probablity_tbl = contingency_tbl/ contingency_tbl.sum()
sns.heatmap(probablity_tbl, annot=True, center=0.5 ,cmap="Greys")
plt.show()
```



Example of using contingency table to relationship between two categorical (non-binary) attributes


```
[20]: contingency_tbl = pd.crosstab(adult_df.occupation,adult_df.race)
probablity_tbl = contingency_tbl/ contingency_tbl.sum()
sns.heatmap(probablity_tbl, annot=True, center=0.5 ,cmap="Greys")
plt.show()
```



2.3.3 Visualizing the relationship between a numerical attribute and a categorical attribute

Visualizing the relationship between a numerical attribute and a categorical attribute

```
[21]: age_discretized = pd.cut(adult_df.age, bins = 5)
contingency_tbl = pd.crosstab(age_discretized,adult_df.race)
probability_tbl = contingency_tbl/ contingency_tbl.sum()
sns.heatmap(probability_tbl, annot=True, center=0.5 ,cmap="Greys")
plt.show()
```



Another example of examining the relationship between a categorical attribute and a numerical attribute

```
[22]: pd.DataFrame(adult_df.groupby(['education', 'education-num']).size()).
      > drop(columns=[0]).reset_index().sort_values('education-num').
      > reset_index(drop=True).transpose()
```

```
[22]:
```

	0	1	2	3	4	5	6	7	\
education	Preschool	1st-4th	5th-6th	7th-8th	9th	10th	11th	12th	
education-num	1	2	3	4	5	6	7	8	

	8	9	10	11	12	\
education	HS-grad	Some-college	Assoc-voc	Assoc-acdm	Bachelors	
education-num	9	10	11	12	13	


	13	14	15
education	Masters	Prof-school	Doctorate

education-num	14	15	16
---------------	----	----	----

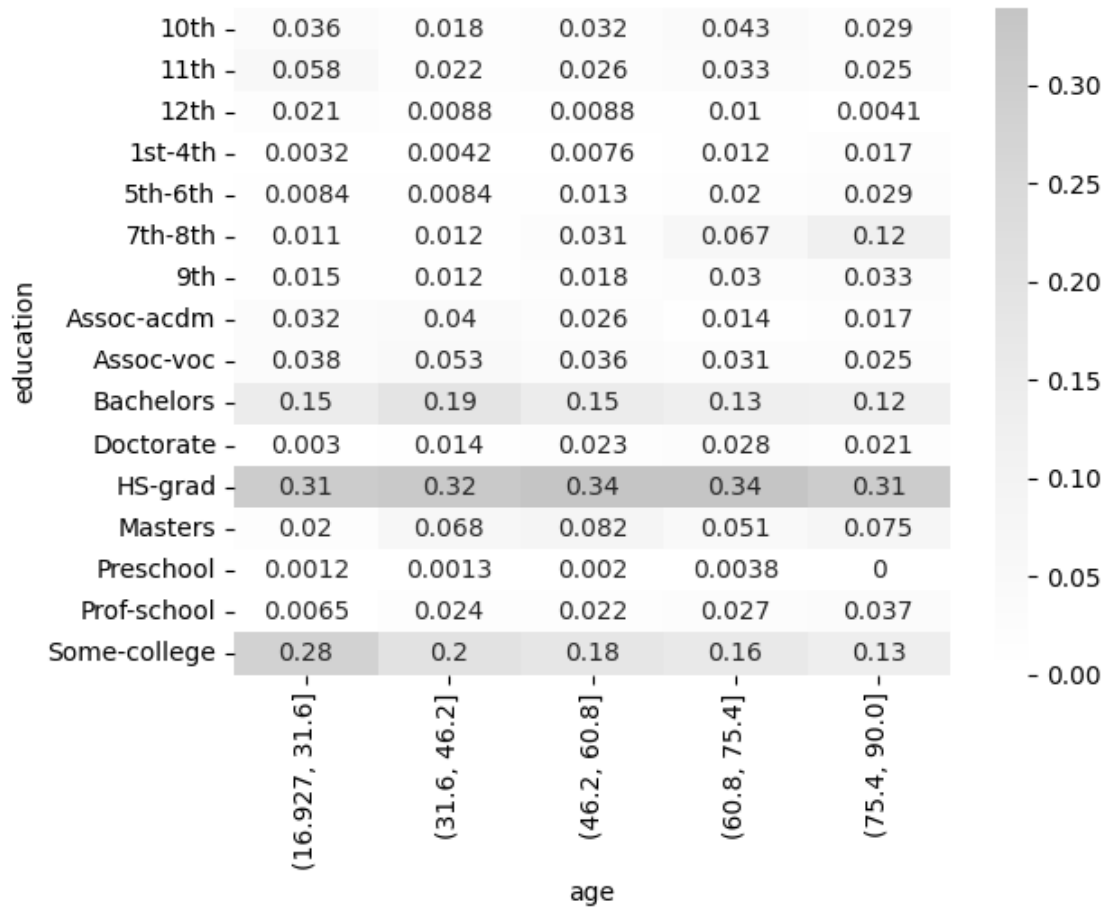
```
[23]: adult_df.groupby(['education', 'education-num']).size()
```

```
[23]: education    education-num
10th             6             933
11th             7            1175
12th             8             433
1st-4th          2             168
5th-6th          3             333
7th-8th          4             646
9th              5             514
Assoc-acdm       12            1067
Assoc-voc        11            1382
Bachelors        13            5355
Doctorate         16             413
HS-grad           9           10501
Masters           14            1723
Preschool         1              51
Prof-school       15             576
Some-college      10            7291
dtype: int64
```

```
[24]: adult_df.plot.scatter(x='age', y='education-num')
plt.show()
```



```
[25]: age_discretized = pd.cut(adult_df['age'], bins = 5)
contingency_tbl = pd.crosstab(adult_df.education, age_discretized)
probability_tbl = contingency_tbl / contingency_tbl.sum()
sns.heatmap(probability_tbl, annot=True, center=0.5, cmap="Greys")
plt.show()
```



2.4 Adding visual dimensions

2.4.1 Example of a 5-dimensional scatterplot

```
[26]: country_df = pd.read_csv('WH Report_preprocessed.csv')
plt.figure(figsize=(15,8))

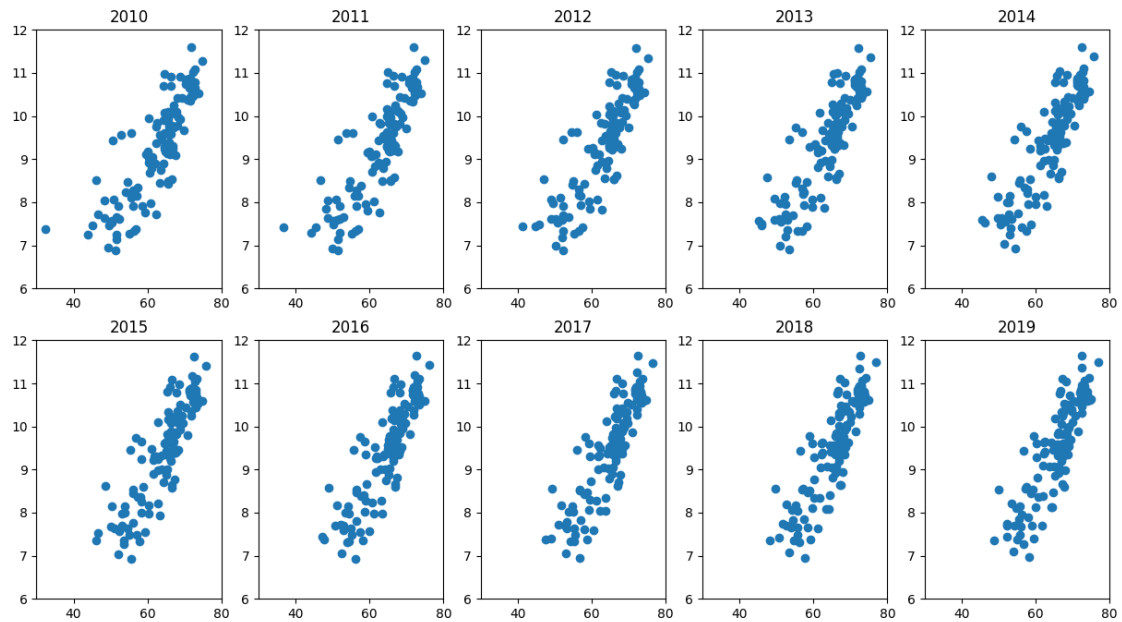
year_poss = country_df.year.unique()

for i, yr in enumerate(year_poss):
    BM = country_df.year == yr
    X= country_df[BM].Healthy_life_expectancy_at_birth
    Y= country_df[BM].Log_GDP_per_capita

    plt.subplot(2,5,i+1)
    plt.scatter(X,Y)
    plt.title(yr)
    plt.xlim([30,80])
```

```
plt.ylim([6,12])

plt.show()
plt.tight_layout()
```




<Figure size 640x480 with 0 Axes>

```
[27]: def plotyear(year):
        BM = country_df.year == year
        X= country_df[BM].Healthy_life_expectancy_at_birth
        Y= country_df[BM].Log_GDP_per_capita

        plt.scatter(X,Y)
        plt.xlabel('Healthy_life_expectancy_at_birth')
        plt.ylabel('Log_GDP_per_capita')
        plt.xlim([30,80])
        plt.ylim([6,12])
        plt.show()
```

```
[28]: plotyear(2011)
```

```
[30]: from ipywidgets import interact, widgets

interact(plotyear, year=widgets.IntSlider(min=2010, max=2019, step=1, value=2010))

interactive(children=(IntSlider(value=2010, description='year', max=2019,
    min=2010), Output()), _dom_classes=(...

[30]: <function __main__.plotyear(year)>
```

The fourth dimension

```
[31]: Continent_poss = country_df.Continent.unique()
colors_dic={'Asia':'b', 'Europe':'g', 'Africa':'r', 'South America':'c',
'Oceania':'m', 'North America':'y', 'Antarctica':'k'}

def plotyear(year):
    for cotinent in Continent_poss:
        BM1 = (country_df.year == year)
        BM2 = (country_df.Continent ==cotinent)
        BM = BM1 & BM2
        X = country_df[BM].Healthy_life_expectancy_at_birth
        Y = country_df[BM].Log_GDP_per_capita
```

```

plt.scatter(X,Y,c=colors_dic[cotinent], marker='o',
            linewidths=0.5,edgecolors='w',label=cotinent)

plt.xlabel('Healthy_life_expectancy_at_birth')
plt.ylabel('Log_GDP_per_capita')
plt.xlim([30,80])
plt.ylim([6,12])
plt.legend(ncol=1)
plt.show()

interact(plotyear,year=widgets.IntSlider(min=2010,max=2019,step=1,value=2010))

```

```

interactive(children=(IntSlider(value=2010, description='year', max=2019,
    min=2010), Output()), _dom_classes=(...

```

```
[31]: <function __main__.plotyear(year)>
```

The fifth dimension

```

[33]: Continent_poss = country_df.Continent.unique()
colors_dic={'Asia':'b', 'Europe':'g', 'Africa':'r', 'South America':'c',
            'Oceania':'m', 'North America':'y', 'Antarctica':'k'}
country_df.sort_values(['population'],inplace = True, ascending=False)

def plotyear(year):
    for cotinent in Continent_poss:
        BM1 = (country_df.year == year)
        BM2 = (country_df.Continent ==cotinent)
        BM = BM1 & BM2
        size = country_df[BM].population/200000
        X = country_df[BM].Healthy_life_expectancy_at_birth
        Y= country_df[BM].Log_GDP_per_capita
        plt.scatter(X,Y,s=size,c=colors_dic[cotinent], marker='o',
                    linewidths=0.5,edgecolors='w',label=cotinent)

    plt.xlabel('Healthy_life_expectancy_at_birth')
    plt.ylabel('Log_GDP_per_capita')
    plt.xlim([30,80])
    plt.ylim([6,12])
    plt.legend(markerscale=0.5)
    plt.show()

interact(plotyear,year=widgets.IntSlider(min=2010,max=2019,step=1,value=2010))

```

```

interactive(children=(IntSlider(value=2010, description='year', max=2019,
    min=2010), Output()), _dom_classes=(...

```

```
[33]: <function __main__.plotyear(year)>
```

2.5 Showing and comparing Trends

```
[34]: amazon_df = pd.read_csv('Amazon Stock.csv')
apple_df = pd.read_csv('Apple Stock.csv')
show_table = amazon_df.iloc[5031:5041][['Date', 'Close']]
show_table.columns = ['Date', 'Amazon']
show_table = show_table.join(apple_df.iloc[5031:5041]['Close'])
show_table.columns = ['Date', 'Amazon', 'Apple']
show_table = show_table.transpose()
show_table.columns = show_table.loc['Date']
show_table.drop(index=['Date'])
```

```
[34]: Date      1/2/2020      1/3/2020      1/6/2020      1/7/2020      1/8/2020  \
Amazon  1898.01001  1874.969971  1902.880005  1906.859985  1891.969971
Apple   74.333511   73.61084    74.197395   73.848442   75.036385

Date      1/9/2020      1/10/2020     1/13/2020     1/14/2020     1/15/2020
Amazon  1901.050049  1883.160034  1891.300049  1869.439941  1862.02002
Apple   76.630219   76.803459   78.444321   77.385063   77.053429
```

2.5.1 Example of visualizing and comparing trends

```
[35]: country_df = pd.read_csv('WH Report_preprocessed.csv')
continent_poss = country_df.Continent.unique()
byContinentYear_df = country_df.groupby(['Continent', 'year']).
    ↳ Perceptions_of_corruption.mean()
Markers_options = ['o', '^', 'P', '8', 's', 'p', '*']

for i, c in enumerate(continent_poss):
    plt.plot([2010, 2019], byContinentYear_df.loc[c, [2010, 2019]],
             label=c, marker=Markers_options[i])
plt.xticks([2010, 2019])
plt.legend(bbox_to_anchor=(1.05, 1))
plt.title('Aggregated values per each continent in 2010 and 2019')
plt.ylabel('Perceptions_of_corruption')
plt.show()
```

