

Chapter 7

April 24, 2025

0.1 Hands-On Data Preprocessing in Python

Learn how to effectively prepare data for successful data analytics

AUTHOR: Dr. Roy Jafari

1 Chapter 7: Classification

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

1.1 Classification models

1.1.1 Example of designing a classification model

1.1.2 Classification Algorithms

1.2 K-Nearest Neighbors (KNN)

1.2.1 Example of using KNN for classification

```
[2]: applicant_df = pd.read_csv('CustomerLoan.csv')
applicant_df.drop(columns = ['Name'],inplace=True)
applicant_df
```

```
[2]:
```

	income	score	default
0	78479	800	NO
1	95483	801	NO
2	101641	815	NO
3	104234	790	NO
4	108726	795	NO
5	112845	750	NO
6	114114	799	NO
7	114799	801	NO
8	119147	805	NO
9	119976	790	NO
10	84519	740	Yes
11	86504	753	Yes
12	89292	750	Yes

13	93941	706	Yes
14	97262	777	Yes
15	102658	680	Yes
16	103760	740	Yes
17	104451	730	Yes
18	107388	789	Yes
19	107400	690	Yes
20	98487	785	NaN

```
[3]: newApplicant = applicant_df.iloc[20]
newApplicant
```

```
[3]: income      98487
score          785
default        NaN
Name: 20, dtype: object
```

```
[4]: applicant_df = pd.read_csv('CustomerLoan.csv')
applicant_df.drop(index = [20],inplace=True)
fig, ax = plt.subplots()

subset = applicant_df.loc[applicant_df['default']=='Yes']
ax.scatter(subset.income, subset.score, marker='o', label='Default-YES',
           color='C1')

subset = applicant_df.loc[applicant_df['default']=='NO']
ax.scatter(subset.income, subset.score, marker='D', label='Default-NO',
           color='C0')

ax.scatter(newApplicant.income, newApplicant.score, marker='*', label='New
Applicant', color='black', s=150)


plt.xlabel('income') # set x-axis label
plt.ylabel('score') # set y-axis label

for _, row in applicant_df.iterrows():
    ax.annotate(row.Name, (row.income -700, row.score-10))

handles, labels = ax.get_legend_handles_labels()

ax.legend(handles, labels, bbox_to_anchor=(1.05, 1))

plt.show()
```



```
[5]: applicant_df = pd.read_csv('CustomerLoan.csv')
applicant_df['income_Normalized'] = (applicant_df.income - applicant_df.income.
    ↪min())/(applicant_df.income.max() - applicant_df.income.min())
applicant_df['score_Normalized'] = (applicant_df.score - applicant_df.score.
    ↪min())/(applicant_df.score.max() - applicant_df.score.min())
```

```
[6]: applicant_df.drop(columns = ['Name'])
```

```
[6]:
```

	income	score	default	income_Normalized	score_Normalized
0	78479	800	NO	0.000000	0.888889
1	95483	801	NO	0.409765	0.896296
2	101641	815	NO	0.558161	1.000000
3	104234	790	NO	0.620647	0.814815
4	108726	795	NO	0.728896	0.851852
5	112845	750	NO	0.828156	0.518519
6	114114	799	NO	0.858737	0.881481
7	114799	801	NO	0.875244	0.896296
8	119147	805	NO	0.980023	0.925926
9	119976	790	NO	1.000000	0.814815
10	84519	740	Yes	0.145553	0.444444
11	86504	753	Yes	0.193387	0.540741
12	89292	750	Yes	0.260573	0.518519
13	93941	706	Yes	0.372605	0.192593
14	97262	777	Yes	0.452635	0.718519
15	102658	680	Yes	0.582669	0.000000
16	103760	740	Yes	0.609225	0.444444

17	104451	730	Yes	0.625877	0.370370
18	107388	789	Yes	0.696653	0.807407
19	107400	690	Yes	0.696942	0.074074
20	98487	785	NaN	0.482155	0.777778

```
[7]: from sklearn.neighbors import KNeighborsClassifier

predictors = ['income_Normalized', 'score_Normalized']
target = 'default'

Xs = applicant_df[predictors].drop(index=[20])
y = applicant_df[target].drop(index=[20])

knn = KNeighborsClassifier(n_neighbors=4)
knn.fit(Xs, y)

newApplicant = pd.DataFrame({'income_Normalized':
                             applicant_df.iloc[20].income_Normalized,
                             'score_Normalized':
                             applicant_df.iloc[20].score_Normalized},
                             index = [20])

predict_y = knn.predict(newApplicant)
print(predict_y)
```

['NO']

1.3 Decision Trees

1.3.1 Example of using Decision Trees for classification

```
[10]: from sklearn.tree import DecisionTreeClassifier

predictors = ['income', 'score']
target = 'default'

Xs = applicant_df[predictors].drop(index=[20])
y = applicant_df[target].drop(index=[20])

classTree = DecisionTreeClassifier()
classTree.fit(Xs, y)

newApplicant = pd.DataFrame({'income':
                             applicant_df.iloc[20].income,
                             'score':
                             applicant_df.iloc[20].score},
                             index = [20])

predict_y = classTree.predict(newApplicant)
```

```
print(predict_y)
```

```
['Yes']
```

```
[11]: from sklearn.tree import plot_tree
plot_tree(classTree,
          feature_names=predictors,
          class_names=y.unique(),
          filled=True,
          impurity=False)
plt.show()
```

